

— C++ PROGRAMS TO ACCOMPANY —

PROGRAMMING LOGIC AND DESIGN

Eighth Edition

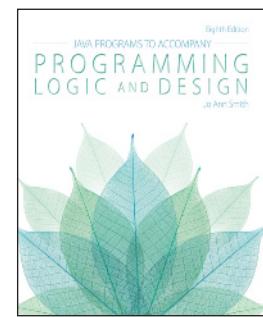
Learn how to transform program logic and design concepts into working programs with this outstanding supplemental handbook. Specifically designed to be paired with the latest edition of Joyce Farrell's highly successful and widely used textbook, *Programming Logic and Design*, this innovative guide, developed by experienced industry practitioner Jo Ann Smith, combines the power of C++ with the popular, language-independent, logical approach of Farrell's text. Together, the two books provide an excellent opportunity for those who want to learn the fundamentals of programming in tandem with the basic concepts of a leading programming language. The guide combines clear explanations of concepts and syntax with pseudocode, complete programming examples, numerous visuals, and real-world, business-related C++ code examples.

Features of the Text

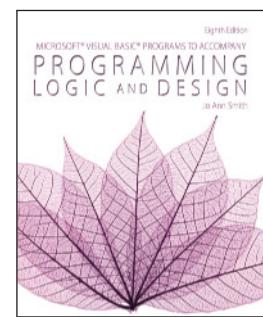
- Fully revised and updated to align with the eighth edition of Joyce Farrell's *Programming Logic and Design* text, both the Comprehensive and Introductory editions.
- Engaging Lab Exercises let students apply newly learned logic and design concepts in C++.
- Complete Programming Examples demonstrate applications from start to finish, and focus on practical business situations without requiring special mathematics or accounting knowledge.

About the Author

An experienced industry practitioner — she has directed a computer training company in the Chicago area and worked for nine years at AT&T Bell Laboratories — Jo Ann Smith is currently a computer consultant in higher education. Ms. Smith previously served as Assistant Professor of Computer Information Systems at Harper College in Palatine, Illinois, and taught at the College of DuPage in Glen Ellyn, Illinois, and at the University of St. Francis in Joliet, Illinois.



Java™ Programs
to Accompany
*Programming Logic
and Design*,
Eighth Edition
Jo Ann Smith
978-1-285-86740-3



Microsoft® Visual
Basic® Programs
to Accompany
*Programming Logic
and Design*,
Eighth Edition
Jo Ann Smith
978-1-285-86739-7



To learn more about Cengage Learning, visit www.cengage.com

Purchase any of our products at your local college store or at our preferred online store www.cengagebrain.com



Smith

C++

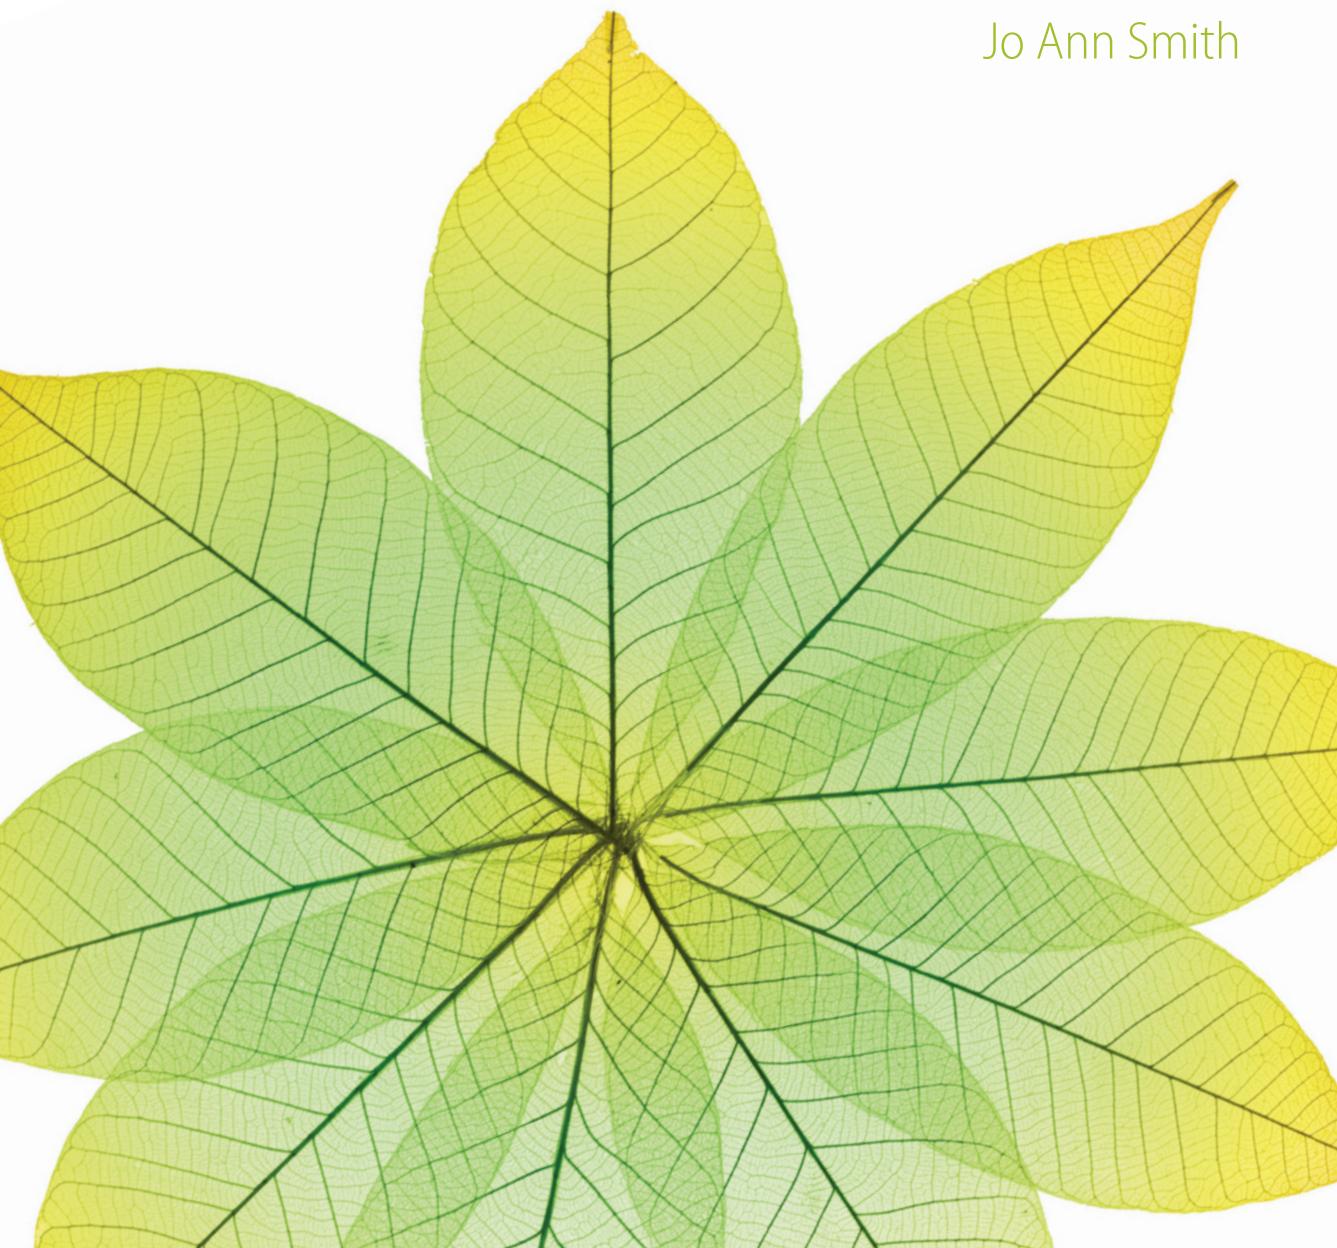
Eighth Edition

Cengage
Learning

— C++ PROGRAMS TO ACCOMPANY —

PROGRAMMING LOGIC AND DESIGN

Eighth Edition



EIGHTH EDITION

C++ PROGRAMS TO ACCOMPANY PROGRAMMING LOGIC AND DESIGN

BY JO ANN SMITH



Australia • Brazil • Japan • Korea • Mexico • Singapore • Spain • United Kingdom • United States



**C++ Programs to Accompany
Programming Logic and Design
Eighth Edition**
Jo Ann Smith

Senior Product Manager: Jim Gish

Senior Content Developer: Alyssa Pratt

Content Project Manager:
Jennifer Feltri-George

Product Assistant: Gillian Daniels

Art Director: Cheryl Pearl, GEX

Proofreader: Lisa Weidenfeld

Copyeditor: Mark Goodin

Indexer: Sharon Hilgenberg

Compositor: Integra Software Services

Cover Designer: GEX Publishing Services

Image Credit: © Kasia/Shutterstock.com

© 2015 Cengage Learning

ALL RIGHTS RESERVED. No part of this work covered by the copyright herein may be reproduced, transmitted, stored or used in any form or by any means graphic, electronic, or mechanical, including but not limited to photocopying, recording, scanning, digitizing, taping, Web distribution, information networks, or information storage and retrieval systems, except as permitted under Section 107 or 108 of the 1976 United States Copyright Act, without the prior written permission of the publisher.

For product information and technology assistance, contact us at
Cengage Learning Customer & Sales Support, cengage.com/support

For permission to use material from this text or product,
submit all requests online at www.cengage.com/permissions

Further permissions questions can be e-mailed to
permissionrequest@cengage.com

Library of Congress Control Number: 2013953075

ISBN-13: 978-1-285-86741-0

Cengage Learning
20 Channel Center Street
Boston, MA 02210
USA

Cengage Learning is a leading provider of customized learning solutions with office locations around the globe, including Singapore, the United Kingdom, Australia, Mexico, Brazil, and Japan. Locate your local office at: www.international.cengage.com/region

Cengage Learning products are represented in Canada by Nelson Education, Ltd.

Visit our corporate Web site at cengage.com.

The programs in this book are for instructional purposes only.

They have been tested with care, but are not guaranteed for any particular intent beyond educational purposes. The author and the publisher do not offer any warranties or representations, nor do they accept any liabilities with respect to the programs.

Brief Contents

iii

Preface	x
Read This Before You Begin	xiii
CHAPTER 1	An Introduction to C++ and the C++ Programming Environment
	1
CHAPTER 2	Variables, Constants, Operators, and Writing Programs Using Sequential Statements
	11
CHAPTER 3	Writing Structured C++ Programs
	29
CHAPTER 4	Writing Programs that Make Decisions
	43
CHAPTER 5	Writing Programs Using Loops
	69
CHAPTER 6	Using Arrays in C++ Programs
	93
CHAPTER 7	File Handling and Applications
	109
CHAPTER 8	Advanced Array Techniques
	125
CHAPTER 9	Advanced Modularization Techniques
	141
CHAPTER 10	Object-Oriented C++
	175
Index	195

Contents

v

Preface	x
Read This Before You Begin	xiii
CHAPTER 1	
An Introduction to C++ and the C++ Programming Environment	1
The C++ Programming Language	2
An Introduction to Object-Oriented Terminology	2
The Structure of a C++ Program	4
The C++ Development Cycle	6
Writing C++ Source Code	6
Compiling a C++ Program	7
Executing a C++ Program	8
Exercise 1-1: Understanding the C++ Compiler	9
Lab 1-1: Compiling and Executing a C++ Program	9
CHAPTER 2	
Variables, Constants, Operators, and Writing Programs Using Sequential Statements	11
Variables	12
Variable Names	12
C++ Data Types	13
Exercise 2-1: Using C++ Variables, Data Types, and Keywords	14
Declaring and Initializing Variables	14
Exercise 2-2: Declaring and Initializing C++ Variables	16
Lab 2-1: Declaring and Initializing C++ Variables	16
Constants	17
Unnamed Constants	17
Named Constants	17
Exercise 2-3: Declaring and Initializing C++ Constants	17
Lab 2-2: Declaring and Initializing C++ Constants	18
Arithmetic and Assignment Operators	18
Arithmetic Operators	18
Assignment Operators and the Assignment Statement	19
Precedence and Associativity	20
Exercise 2-4: Understanding Operator Precedence and Associativity	22

Lab 2-3: Understanding Operator Precedence and Associativity	22
Sequential Statements, Comments, and Interactive Input Statements	23
Exercise 2-5: Understanding Sequential Statements	26
Lab 2-4: Understanding Sequential Statements	27
CHAPTER 3 Writing Structured C++ Programs	29
Using Flowcharts and Pseudocode to Write a C++ Program	30
Lab 3-1: Using Flowcharts and Pseudocode to Write a C++ Program	34
Writing a Modular Program in C++	35
Lab 3-2: Writing a Modular Program in C++	40
CHAPTER 4 Writing Programs that Make Decisions	43
Boolean Operators	44
Relational Operators	44
Logical Operators	45
Relational and Logical Operator Precedence and Associativity	46
Comparing <code>strings</code>	48
Decision Statements	50
The <code>if</code> Statement	50
Exercise 4-1: Understanding <code>if</code> Statements	52
Lab 4-1: Understanding <code>if</code> Statements	53
The <code>if else</code> Statement	54
Exercise 4-2: Understanding <code>if else</code> Statements	55
Lab 4-2: Understanding <code>if else</code> Statements	56
Nested <code>if</code> Statements	57
Exercise 4-3: Understanding Nested <code>if</code> Statements	58
Lab 4-3: Understanding Nested <code>if</code> Statements	59
The <code>switch</code> Statement	60
Exercise 4-4: Using a <code>switch</code> Statement	61
Lab 4-4: Using a <code>switch</code> Statement	62
Using Decision Statements to Make Multiple Comparisons	63
Using AND Logic	63
Using OR Logic	64
Exercise 4-5: Making Multiple Comparisons in Decision Statements	64
Lab 4-5: Making Multiple Comparisons in Decision Statements	66

CHAPTER 5	Writing Programs Using Loops	69
The Increment (<code>++</code>) and Decrement (<code>--</code>) Operators	70	
Exercise 5-1: Using the Increment and Decrement Operators	71	
Writing a <code>while</code> Loop in C++	72	
Exercise 5-2: Using a <code>while</code> Loop	74	
Using a Counter to Control a Loop	74	
Exercise 5-3: Using a Counter-Controlled <code>while</code> Loop	75	
Lab 5-1: Using a Counter-Controlled <code>while</code> Loop	76	
Using a Sentinel Value to Control a Loop	76	
Exercise 5-4: Using a Sentinel Value to Control a <code>while</code> Loop	79	
Lab 5-2: Using a Sentinel Value to Control a <code>while</code> Loop	79	
Writing a <code>for</code> Loop in C++	80	
Exercise 5-5: Using a <code>for</code> Loop	81	
Lab 5-3: Using a <code>for</code> Loop	82	
Writing a <code>do while</code> Loop in C++	82	
Exercise 5-6: Using a <code>do while</code> Loop	83	
Lab 5-4: Using a <code>do while</code> Loop	83	
Nesting Loops	84	
Exercise 5-7: Nesting Loops	85	
Lab 5-5: Nesting Loops	86	
Accumulating Totals in a Loop	87	
Exercise 5-8: Accumulating Totals in a Loop	88	
Lab 5-6: Accumulating Totals in a Loop	89	
Using a Loop to Validate Input	89	
Exercise 5-9: Validating User Input	91	
Lab 5-7: Validating User Input	91	
CHAPTER 6	Using Arrays in C++ Programs	93
Array Basics	94	
Declaring Arrays	94	
Initializing Arrays	96	
Accessing Array Elements	97	
Staying Within the Bounds of an Array	98	
Using Constants with Arrays	98	
Exercise 6-1: Array Basics	99	
Lab 6-1: Array Basics	99	
Searching an Array for an Exact Match	100	
Exercise 6-2: Searching an Array for an Exact Match	102	
Lab 6-2: Searching an Array for an Exact Match	103	
Parallel Arrays	103	
Exercise 6-3: Parallel Arrays	106	
Lab 6-3: Parallel Arrays	107	

CHAPTER 7	File Handling and Applications	109
File Handling	110	
Using Input and Output Classes	110	
Opening a File for Reading	110	
Reading Data from an Input File	111	
Reading Data Using a Loop and EOF	112	
Opening a File for Writing	112	
Writing Data to an Output File	113	
Exercise 7-1: Opening Files and Performing File Input	114	
Lab 7-1: Opening Files and Performing File Input	115	
Understanding Sequential Files and Control Break Logic	116	
Exercise 7-2: Accumulating Totals in Single-Level Control Break Programs	121	
Lab 7-2: Accumulating Totals in Single-Level Control Break Programs	121	
CHAPTER 8	Advanced Array Techniques	125
Sorting Data	126	
Swapping Data Values	126	
Exercise 8-1: Swapping Values	127	
Lab 8-1: Swapping Values	127	
Using a Bubble Sort	128	
The <code>main()</code> Function	131	
The <code>fillArray()</code> Function	132	
The <code>sortArray()</code> Function	133	
The <code>displayArray()</code> Function	133	
Exercise 8-2: Using a Bubble Sort	134	
Lab 8-2: Using a Bubble Sort	134	
Using Multidimensional Arrays	135	
Exercise 8-3: Using Multidimensional Arrays	138	
Lab 8-3: Using Multidimensional Arrays	139	
CHAPTER 9	Advanced Modularization Techniques	141
Writing Functions with No Parameters	142	
Exercise 9-1: Writing Functions with No Parameters	144	
Lab 9-1: Writing Functions with No Parameters	145	
Writing Functions that Require a Single Parameter	145	
Exercise 9-2: Writing Functions that Require a Single Parameter	148	
Lab 9-2: Writing Functions that Require a Single Parameter	148	
Writing Functions that Require Multiple Parameters	149	

Exercise 9-3: Writing Functions that Require Multiple Parameters	151
Lab 9-3: Writing Functions that Require Multiple Parameters	152
Writing Functions that Return a Value	152
Exercise 9-4: Writing Functions that Return a Value	155
Lab 9-4: Writing Functions that Return a Value	155
Passing an Array and an Array Element to a Function	156
Exercise 9-5: Passing Arrays to Functions	159
Lab 9-5: Passing Arrays to Functions	160
Passing Arguments by Reference and by Address	160
Pass by Reference	161
Pass by Address	162
Exercise 9-6: Pass by Reference and Pass by Address	167
Lab 9-6: Pass by Reference and Pass by Address	168
Overloading Functions	168
Exercise 9-7: Overloading Functions	171
Lab 9-7: Overloading Functions	172
Using C++ Built-in Functions	172
Exercise 9-8: Using C++ Built-in Functions	173
Lab 9-8: Using C++ Built-in Functions	174
CHAPTER 10 Object-Oriented C++	175
A Programmer-Defined Class	176
Creating a Programmer-Defined Class	176
Adding Attributes to a Class	178
Adding Methods to a Class	179
Exercise 10-1: Creating a Class in C++	182
Lab 10-1: Creating a Class in C++	183
Reusing C++ Classes	184
Defining a Derived Class	185
Using a Derived Class in a C++ Program	191
Exercise 10-2: Using Inheritance to Create a Derived Class in C++	193
Lab 10-2: Using Inheritance to Create a Derived Class in C++	193
Index	195

Preface

x

C++ Programs to Accompany Programming Logic and Design, Eighth Edition (also known as *C++ PAL*) is designed to provide students with an opportunity to write C++ programs as part of an Introductory Programming Logic course. It accompanies the student's primary text, *Programming Logic and Design, Eighth Edition*, by Joyce Farrell. This textbook assumes no programming language experience and provides the beginning programmer with a guide to writing structured programs and simple object-oriented programs using introductory elements of the popular C++ programming language. It is not intended as a textbook for a course in C++ programming. The writing is non-technical and emphasizes good programming practices. The examples do not assume mathematical background beyond high school math. Additionally, the examples illustrate one or two major points; they do not contain so many features that students become lost following irrelevant and extraneous details.

The examples in *C++ PAL* are often examples presented in the primary textbook, *Programming Logic and Design, Eighth Edition*. The following table shows the correlation between topics in the two books.

<i>C++ PAL, Eighth Edition</i>	<i>Programming Logic and Design, Eighth Edition</i>
Chapter 1: An Introduction to C++ and the C++ Programming Environment	Chapter 1: An Overview of Computers and Programming
Chapter 2: Variables, Constants, Operators, and Writing Programs Using Sequential Statements	Chapter 2: Elements of High-Quality Programs Chapter 3: Understanding Structure
Chapter 3: Writing Structured C++ Programs	Chapter 2: Elements of High-Quality Programs Chapter 3: Understanding Structure
Chapter 4: Writing Programs that Make Decisions	Chapter 4: Making Decisions
Chapter 5: Writing Programs Using Loops	Chapter 5: Looping
Chapter 6: Using Arrays in C++ Programs	Chapter 6: Arrays
Chapter 7: File Handling and Applications	Chapter 7: File Handling and Applications
Chapter 8: Advanced Array Techniques	Chapter 8: Advanced Data Handling Concepts
Chapter 9: Advanced Modularization Techniques	Chapter 9: Advanced Modularization Techniques
Chapter 10: Object Oriented C++	Chapter 10: Object Oriented Programming Chapter 11: More Object Oriented Programming Concepts

Organization and Coverage

C++ Programs to Accompany Programming Logic and Design, Eighth Edition provides students with a review of the programming concepts they are introduced to in their primary textbook. It also shows them how to use C++ to transform their program logic and design into working programs. Chapter 1 introduces the structure of a C++ program, how to compile and run a C++ console program, and introductory object-oriented concepts. Chapter 2 discusses C++'s data types, variables, constants, arithmetic and assignment operators, and using sequential statements to write a complete C++ program. In Chapter 3, students learn how to transform pseudocode and flowcharts into C++ programs. Chapters 4 and 5 introduce students to writing C++ programs that make decisions and programs that use looping constructs. Students learn to use C++ to develop more sophisticated programs that include using arrays, control breaks, and file input and output in Chapters 6 and 7. In Chapter 8, students learn about sorting data items in an array and using multidimensional arrays. Passing parameters to functions is introduced in Chapter 9. Lastly, in Chapter 10, students use C++ to write programs that include programmer-defined classes.

This book combines text explanation of concepts and syntax along with pseudocode and actual C++ code examples to provide students with the knowledge they need to implement their logic and program designs using the C++ programming language. This book is written in a modular format and provides paper-and-pencil exercises as well as lab exercises after each major topic is introduced. The exercises provide students with experience in reading and writing C++ code as well as modifying and debugging existing code. In the labs, students are asked to complete partially prewritten C++ programs. Using partially prewritten programs allows students to focus on individual concepts rather than an entire program. The labs also allow students to see their programs execute.

C++ PAL, Eighth Edition is unique because:

- It is written and designed to correspond to the topics in the primary textbook, *Programming Language and Design, Eighth Edition*.
- The examples are everyday examples; no special knowledge of mathematics, accounting, or other disciplines is assumed.
- It introduces students to introductory elements of the C++ programming language rather than overwhelming beginning programmers with more detail than they are prepared to use or understand.
- Text explanations are interspersed with pseudocode from the primary book, thus reinforcing the importance of programming logic.
- Complex programs are built through the use of complete examples. Students see how an application is built from start to finish instead of studying only segments of programs.

Features of the Text

Every chapter in this book includes the following features. These features are both conducive to learning in the classroom and enable students to learn the material at their own pace.

- **Objectives:** Each chapter begins with a list of objectives so the student knows the topics that will be presented in the chapter. In addition to providing a quick reference to topics covered, this feature provides a useful study aid.
- **Figures and illustrations:** This book has plenty of visuals, which provide the reader with a more complete learning experience, rather than one that involves simply studying text.
- **Notes:** These brief notes provide additional information—for example, a common error to watch out for.
- **Exercises:** Each section of each chapter includes meaningful paper-and-pencil exercises that allow students to practice the skills and concepts they are learning in the section.
- **Labs:** Each section of each chapter includes meaningful lab work that allows students to write and execute programs that implement their logic and program design.

Acknowledgments

I would like to thank all of the people who helped to make this book possible. Thanks to Alyssa Pratt, Senior Content Developer, and Jim Gish, Senior Product Manager, for their help and encouragement. I am grateful to Jennifer Feltri-George, Content Project Manager, Serge Palladino, Manuscript Quality Assurance, and Anandhavalli Namachivayam of Integra Software Services, for overseeing the production of the printed book. It is a pleasure to work with so many fine people who are dedicated to producing quality instructional materials.

This book is dedicated to my mom. Her courage and determination continue to inspire me. I am grateful for such an excellent role model.

Jo Ann Smith

Read This Before You Begin

xiii

To the User

Data Files

To complete most of the lab exercises, you will need data files that have been created for this book. Your instructor can provide the data files. You also can obtain the files electronically from the publisher at www.CengageBrain.com (search under the ISBN for this book).

You can use a computer in your school lab or your own computer to complete the lab exercises in this book.

Solutions

Solutions to the Exercises and Labs are provided to instructors on the Cengage Learning Web site at sso.cengage.com. The solutions are password protected.

Using Your Own Computer

To use your own computer to complete the material in this textbook, you will need the following:

- Computer with a 1.6 GHz or faster processor
- Operating System:
 - ◆ Windows® 7 SP1 (x86 & x64)
 - ◆ Windows® 8 (x86 & x64)
 - ◆ Windows Server® 2008 R2 SP1 (x64)
 - ◆ Windows Server 2012 (x64)
- 1 GB RAM (1.5 GB if running on a virtual machine)
- 5 GB of available hard-disk space
- 5400 RPM hard drive
- DirectX 9 capable video card running at 1024 x 768 or higher display resolution

This book was written using Microsoft Windows 8 and Quality Assurance tested using Microsoft Windows 7.

Using the C++ Compiler

To use the C++ compiler (cl) you open the Developer Command Prompt for VS2012.

To do this in Windows 8, right-click on a blank area of the **Start** screen, click **All apps**, scroll until you see the Microsoft Visual Studio 2012 area, and then click **Developer Command Prompt for VS2012**. In the next version of Windows 8, you click the down arrow on the **Start** screen, scroll until you see the Microsoft Visual Studio 2012 area, and then click **Developer Command Prompt for VS2012**.

In Windows 7, click the **Start** button, point to **All Programs**, click **Visual Studio 2012**, click **Visual Studio Tools**, and then click **Developer Command Prompt for VS2012**.

To The Instructor

To complete some of the Exercises and Labs in this book, your students must use the data files provided with this book. These files are available from the publisher at www.CengageBrain.com (search under the ISBN for this book). Follow the instructions in the Help file to copy the data files to your server or standalone computer. You can view the Help file using a text editor such as WordPad or Notepad. Once the files are copied, you may instruct your students to copy the files to their own computers or workstations.

Cengage Learning Data Files

You are granted a license to copy the data files to any computer or computer network used by individuals who have purchased this book.

EIGHTH EDITION

C++ PROGRAMS TO ACCOMPANY PROGRAMMING LOGIC AND DESIGN

1

CHAPTER

An Introduction to C++ and the C++ Programming Environment

After studying this chapter, you will be able to:

- ④ Discuss the C++ programming language and its history
- ④ Explain introductory concepts and terminology used in object-oriented programming
- ④ Recognize the structure of a C++ program
- ④ Complete the C++ development cycle, which includes creating a source code file, compiling the source code, and executing a C++ program

You should do the exercises and labs in this chapter only after you have finished Chapter 1 of your book, *Programming Logic and Design, Eighth Edition*, by Joyce Farrell. This chapter introduces the C++ programming language and its history. It explains some introductory object-oriented concepts, and describes the process of compiling and executing a C++ program. You begin writing C++ programs in Chapter 2 of this book.

2

The C++ Programming Language

The C programming language was written in the early 1970s by Dennis Ritchie at AT&T Bell Labs. C is an important programming language because it is both a high-level and a low-level programming language. It is a **high-level language**, which means that it is more English-like and easier for programmers to use than a low-level language. At the same time, it possesses **low-level language** capabilities that allow programmers to directly manipulate the underlying computer hardware.



Due to their power, C and C++ have been used in the programming of special effects for action movies and video games.

The C++ programming language was developed by Bjarne Stroustrup at AT&T Bell Labs in 1979 and inherited the widespread popularity of C. Because many programmers liked using the powerful C programming language, it was an easy step to move on to the new C++ language.

What makes C++ especially useful for today's programmers is that it is an object-oriented programming language. The term **object-oriented** encompasses a number of concepts explained later in this chapter and throughout this book. For now, all you need to know is that an object-oriented programming language is modular in nature, allowing the programmer to build a program from reusable parts of programs called classes, objects, and methods.

An Introduction to Object-Oriented Terminology

You must understand a few object-oriented concepts to be successful at reading and working with C++ programs in this book. Note, however, that you will not learn enough to make you a C++ programmer. You will have to take additional courses in C++ to become a C++ programmer. This book teaches you only the basics.

To fully understand the term *object-oriented*, you need to know a little about procedural programming. Procedural programming is a style of programming that is older than object-oriented programming. **Procedural programs** consist of statements that the computer runs or **executes**. Many of the statements make calls (a request to run or execute) to groups of other statements that are known as procedures, modules, methods, or subroutines. These programs are known as "procedural" because they perform a sequence of procedures. Procedural programming focuses on writing code that takes some data (for example, some sales figures), performs a specific task using the data (for example, adding up the sales figures),

and then produces output (for example, a sales report). When people who use procedural programs (the **users**) decide they want their programs to do something slightly different, a programmer must revise the program code, taking great care not to introduce errors into the logic of the program.

Today, we need computer programs that are more flexible and easy to revise. Object-oriented programming languages, including C++, were introduced to meet this need. In object-oriented programming, the programmer can focus on the data that he or she wants to manipulate, rather than the individual lines of code required to manipulate that data (although those individual lines still must be written eventually). An **object-oriented program** is made up of a collection of interacting objects. An **object** represents something in the real world, such as a car, an employee, a video game character, or an item in an inventory. An object includes (or **encapsulates**) both the data related to the object and the tasks you can perform on that data. The term **behavior** is sometimes used to refer to the tasks you can perform on an object's data. For example, the data for an inventory object might include a list of inventory items, the number of each item in stock, the number of days each item has been in stock, and so on. The behaviors of the inventory object might include calculations that add up the total number of items in stock and calculations that determine the average amount of time each item remains in inventory.



The preceding assumes you are using classes that someone else previously developed. That programmer must write code that manipulates the object's data.

In object-oriented programming, the data items within an object are known collectively as the object's **attributes**. You can think of an attribute as one of the characteristics of an object, such as its shape, its color, or its name. The tasks the object performs on that data are known as the object's **methods**. (You can also think of a method as an object's behavior.) Because methods are built into objects, when you create a C++ program, you do not always have to write line after line of code telling the program exactly how to manipulate the object's data. Instead, you can write a shorter line of code, known as a **call**, that passes a message to the method indicating that you need it to do something.

For example, you can display dialog boxes, scroll bars, and buttons for a user of your program to type in or click on simply by sending a message to an existing object. At other times, you will be responsible for creating your own classes and writing the code for the methods that are part of that class. Whether you use existing, prewritten classes or create your own classes, one of your main jobs as a C++ programmer is to communicate with the various objects in a program (and the methods of those objects) by passing messages. Individual objects in a program can also pass messages to other objects.

When C++ programmers write an object-oriented program, they begin by creating a class. A **class** can be thought of as a template for a group of similar objects. In a class, the programmer specifies the data (attributes) and behaviors (methods) for all objects that belong to that class. An object is sometimes referred to as an **instance** of a class, and the process of creating an object is referred to as **instantiation**.

To understand the terms *class*, *instance*, and *instantiation*, it is helpful to think of them in terms of a real-world example—baking a chocolate cake. The recipe is similar to a class and an actual cake is an object. If you wanted to, you could create many chocolate cakes that are all based on the same recipe. For example, your mother’s birthday cake, your sister’s anniversary cake, and the cake for your neighborhood bake sale all might be based on a single recipe that contains the same data (ingredients) and methods (instructions). In object-oriented programming, you can create as many objects as you need in your program from the same class.

The Structure of a C++ Program

When a programmer learns a new programming language, the first program he or she traditionally writes is a Hello World program—a program that displays the message “Hello World” on the screen. Creating this simple program illustrates that the language is capable of instructing the computer to communicate with the outside world. The C++ version of the Hello World program is shown in Figure 1-1.

```
#include <iostream>
using namespace std;
int main()
{
    cout << "Hello World" << endl;
    return 0;
}
```

Figure 1-1 Hello World program

At this point, you are not expected to understand all the code in Figure 1-1. Just notice that the code begins with the preprocessor directive, `#include <iostream>`. The C++ **preprocessor** is a program that processes your C++ program before the compiler processes it. The `#include` directive tells the compiler to include another file when the program is compiled. This makes it easy for you to use code previously written by you or others in your programs without having to re-create it. You will learn more about the Visual C++ compiler later in this chapter. Following the `#include`, you see `<iostream>`, which is the name of a header file you want to include in this program. The **iostream header file** gives your program access to what it needs to perform input and output in a C++ program. The name of the header file is placed within angle brackets (`< >`). The angle brackets tell the compiler to look for this file in a directory that is specified by the compiler you are using. You will learn more about preprocessor directives throughout this book.



Namespaces are a relatively new addition to C++ and are used primarily in large programs. In this book, the only namespace you will use is the `std` namespace.

The next line (`using namespace std;`) instructs the compiler to use the `std` namespace. You can think of a **namespace** as a container that holds various program elements. The **std** name space contains everything C++ programs need to use the Standard C++ library. The **Standard C++ library** adds functionality to the C++ programming language. For example, this program needs to use the `std` namespace to have access to `cout` and `endl`, which you see on the fifth line in Figure 1-1. Notice that this line ends with a semicolon (`;`). In fact, all C++ statements end with a semicolon. The reason the previous line, `#include <iostream>`, does not end with a semicolon is that it is a preprocessor directive, not a C++ statement.



You can tell `main()` is a function because of the parentheses; all C++ function names are followed by parentheses.

On the third line you see the start of a function named `main()`. A **function** is a group of C++ statements that perform a specified task. This is a special function in a C++ program; the `main()` function is the first function that executes when any program runs. The programs in the first eight chapters of this book will include only a `main()` function. In later chapters you will be able to include additional functions.

The first part of any function is its **header**. In Figure 1-1, the header for the `main()` function begins with the `int` keyword, followed by the function name, which is `main()`. A **keyword** is a special word that is reserved by C++ to have a special meaning. To understand the keyword `int` you need to know that functions often send values back to a calling function (for example, the result of a calculation), which can then be used elsewhere in the program. Another way to say this is that functions sometimes return a value. In Figure 1-1, the keyword `int` indicates that the `main()` function returns an integer. You will learn more about functions returning values in Chapter 9 of this book.

The opening curly brace (`{`) on the fourth line of Figure 1-1 marks the beginning of the body of the `main()` function and the closing curly brace (`}`) on the last line of Figure 1-1 marks the end of the `main()` function. All the code within this pair of curly braces executes when the `main()` function executes. In Figure 1-1, there are two lines of code between the curly braces. The first is:

```
cout << "Hello World" << endl;
```



Do not confuse the last character in `endl`; it is a lowercase letter `l`, not the numeral `1`.

This is the line that causes the words `Hello World` to appear on the user's screen. This line consists of multiple parts. The first part, `cout`, is the name of an object that represents the user's screen. Next, you see `<<`, which is called the **insertion** or **output** operator. You use `cout` and `<<` to output what follows, which in this example is the string constant `"Hello World"`. (The quotation marks will not appear on the screen, but they are necessary to make the program work.) After `"Hello World"` you see another `<<` which causes `endl` to be displayed (after `Hello World`) on the user's screen. For now, think of `endl` as a

newline that causes the cursor to be positioned on the line after `Hello World`. You will learn more about `endl` in Chapter 9 of this book. Note that the semicolon that ends the `cout << "Hello World" << endl;` statement is required because it tells the compiler that this is the end of the statement.

6

The next line of code is `return 0;`. This statement instructs the compiler to return the value 0 from the `main()` function. Remember, when you saw the header for the `main()` function on the third line of this program, you used the keyword `int` to specify that the `main()` function returns an integer; 0 is the returned integer. Conventionally, when a program returns a 0, it means “everything went well and the program is finished.”

Next, you learn about the C++ development cycle so later in this chapter you can compile the Hello World program and execute it.

The C++ Development Cycle

When you finish designing a program and writing the C++ code that implements your design, you must compile and execute your program. This three-step process of writing code, compiling code, and executing code is called the C++ development cycle. It is illustrated in Figure 1-2.

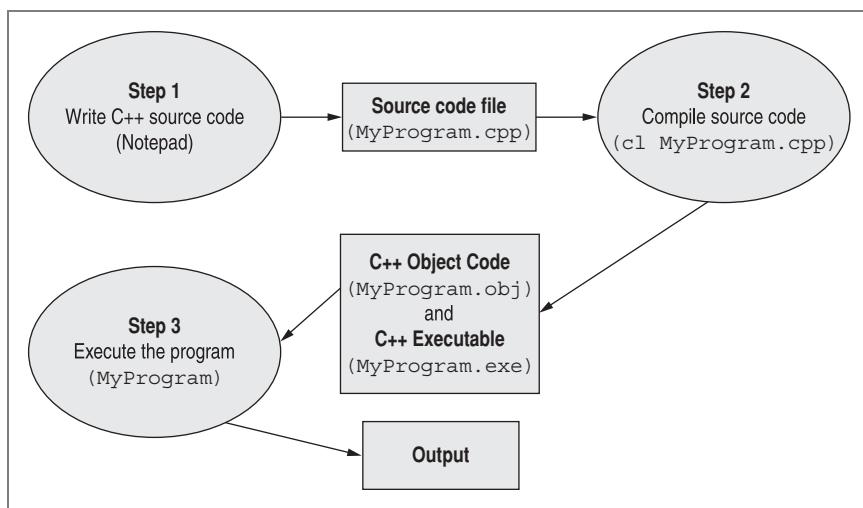


Figure 1-2 C++ development cycle

Writing C++ Source Code

As you learned in the previous section, you write a C++ program by creating a function named `main()` that contains C++ statements. But what do you use to write the program, and where do you save it?

To write a C++ program, you can use any text editor, but the steps in this book assume you are using Windows Notepad. To start Notepad in Windows 7, click the **Start** button, point to **All Programs**, click **Accessories**, and then click **Notepad**. To start Notepad in Windows 8, right-click on a blank area of the **Start** screen, click **All apps** scroll to the end, and then click **Notepad**. In the next version of Windows 8, you click the down arrow on the **Start** screen, scroll to the end, and then click **Notepad**. Once Notepad starts, you simply type in your C++ source code. **Source code** is the term used for the statements that make up a C++ program. For example, the code shown earlier in Figure 1-1 is source code.

When you save the file that contains the source code, it is important to give the file a meaningful name, and then add the extension `.cpp`. For the Hello World program, an appropriate name for the source code file is `HelloWorld.cpp`. Of course, it is also important to remember the location of the folder in which you save your source code file.

Compiling a C++ Program

The Visual C++ compiler is named `c1`. It is a program that is responsible for a two-step process that takes your source code and transforms it into object code and then links the object code to create executable code.



The Visual C++ compiler, `c1`, is actually a compiler and a linker.

Object code is code in computer-readable form that is linked with libraries to create an executable file. **Executable code** is the *1s* and *0s* a computer needs to execute a program. The C++ compiler automatically saves the object code in a file. This file has the same name as the source code file, but it has an `.obj` extension rather than a `.cpp` extension.

The following steps show how to compile a source code file. These steps assume you have already created and saved the `HelloWorld.cpp` source code file.

1. Open a Command Prompt window. To do this in Windows 7, click the **Start** button, point to **All Programs**, click **Accessories**, click **Visual Studio 2012 Express**, and then click **Developer Command Prompt for VS2012**. In Windows 8, right-click on a blank area of the **Start** screen, click **All apps**, scroll until you see the Microsoft Visual Studio 2012 area, and then click **Developer Command Prompt for VS2012**. In the next version of Windows 8, you click the down arrow on the **Start** screen, click **All apps**, scroll until you see the Microsoft Visual Studio 2012 area, and then click **Developer Command Prompt for VS2012**. The cursor blinks to the right of the current file path.
2. To compile your source code file, you first have to change to the file path containing your source code file. To do this, type `cd drive letter:\path` where **drive letter** is the drive containing your file, and **path** is the path to the folder containing your file. For example, to gain access to a file stored in a folder named “Testing”, which is in turn stored in a folder named “My Program”, which is stored on the C drive, you would type `cd C:\My Program\Testing`. After you type the command, press **Enter**. The cursor now blinks next to the file path for the folder containing your source code file.

3. Type the following command, which uses the C++ compiler, `c1`, to compile the program:

c1 HelloWorld.cpp

If there are no syntax errors in your source code, a file named `HelloWorld.obj` and a file named `HelloWorld.exe` are created. You do not see anything special happen. However, the files you just created contain the object code (`HelloWorld.obj`) and executable code (`HelloWorld.exe`) for the Hello World program. If there are syntax errors, you will see error messages on the screen; in that case, you need to go back to Notepad to fix the errors, save the source code file again, and recompile until there are no syntax errors remaining. **Syntax errors** are messages from the compiler that tell you what your errors are and where they are located in your source code file. For example, omitting a semicolon at the end of the statement `cout << "Hello World"` `<< endl` results in a syntax error.



At this point in your programming career, do not expect to understand the contents of files with an `.obj` or `.exe` extension if you open one using a text editor such as Notepad.

4. After the program is compiled, you can use the `dir` command to display a directory listing to see the files named `HelloWorld.obj` and `HelloWorld.exe`. To execute the `dir` command, you type `dir` at the command prompt. For example, if your source code file is located at `C:\My Program\Testing`, the command prompt and `dir` command should be `C:\My Program\Testing> dir`. The `HelloWorld.obj` and `HelloWorld.exe` files should be in the same directory as the source code file `HelloWorld.cpp`.

Executing a C++ Program

To execute the Hello World program, do the following:

1. Open a Command Prompt window. To do this, refer to Step 1 under Compiling a C++ program. Change to the file path containing your executable code file, if necessary, and then enter the following command:

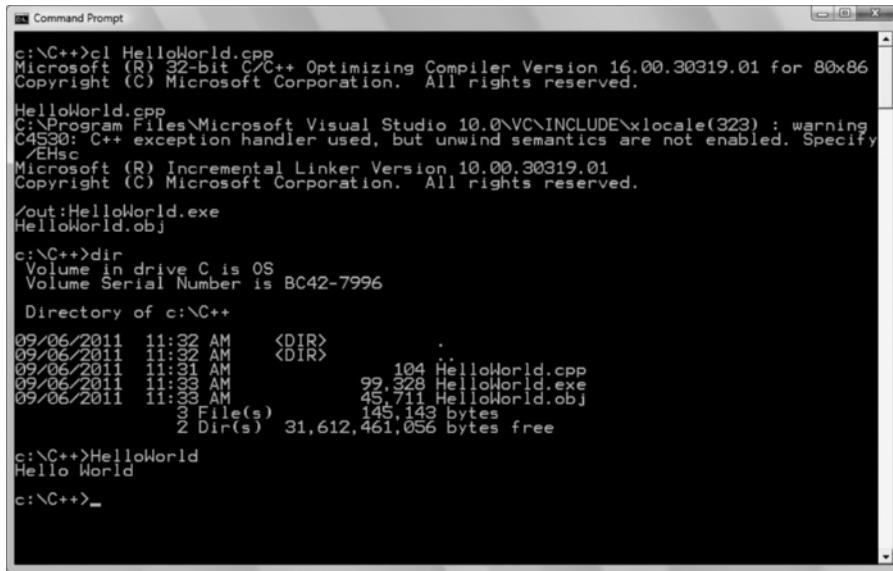
HelloWorld



You must be in the same directory that contains the `.exe` file when you execute the program.

2. When the program executes, the words `Hello World` appear in the Command Prompt window.

Figure 1-3 illustrates the steps involved in compiling `HelloWorld.cpp` using the `cl` compiler, and executing the `dir` command to verify that the files `HelloWorld.obj` and `HelloWorld.exe` were created, as well as the output generated by executing the Hello World program.



The screenshot shows a Windows Command Prompt window titled "Command Prompt". The window contains the following text:

```
c:\C++>cl HelloWorld.cpp
Microsoft (R) 32-bit C/C++ Optimizing Compiler Version 16.00.30319.01 for 80x86
Copyright (C) Microsoft Corporation. All rights reserved.

HelloWorld.cpp
C:\Program Files\Microsoft Visual Studio 10.0\VC\INCLUDE\xlocale(323) : warning
C4530: C++ exception handler used, but unwind semantics are not enabled. Specify
/EHsc
Microsoft (R) Incremental Linker Version 10.00.30319.01
Copyright (C) Microsoft Corporation. All rights reserved.

/out:HelloWorld.exe
HelloWorld.obj

c:\C++>dir
Volume in drive C is OS
Volume Serial Number is BC42-7996

Directory of c:\C++
09/06/2011  11:32 AM    <DIR>
09/06/2011  11:32 AM           104 HelloWorld.cpp
09/06/2011  11:31 AM           99,328 HelloWorld.exe
09/06/2011  11:33 AM           45,711 HelloWorld.obj
               3 File(s)   145,143 bytes
               2 Dir(s)  31,612,461,056 bytes free

c:\C++>HelloWorld
Hello World
c:\C++>
```

Figure 1-3 Compiling and executing the Hello World program

Exercise 1-1: Understanding the C++ Compiler

In this exercise, assume you have written a C++ program and stored your source code in a file named `MyCPlusPlusProgram.cpp`. Answer the following questions:

1. What command would you use to compile the source code?

2. What command would you use to execute the program?

Lab 1-1: Compiling and Executing a C++ Program

In this lab, you compile and execute a prewritten C++ program, and then answer some questions about the program.

1. Open the source code file named `Programming.cpp` using Notepad or the text editor of your choice.
2. Save this source code file in a directory of your choice, and then change to that directory.

- 10
-
3. Compile the source code file. There should be no syntax errors. Record the command you used to compile the source code file.

 4. Execute the program. Record the command you used to execute the program, and record the output of this program.

 5. Modify the program so it displays "I'm learning how to program in C++". Save the file as `C++Programming.cpp`. Compile and execute.
 6. Modify the `C++Programming` program so it prints two lines of output. Add a second output statement that displays "That's Awesome!" Save the modified file as `Awesome.cpp`. Compile and execute the program.

2

CHAPTER

Variables, Constants, Operators, and Writing Programs Using Sequential Statements

After studying this chapter, you will be able to:

- ◎ Name variables and use appropriate data types
- ◎ Declare and initialize variables
- ◎ Understand and use unnamed and named constants
- ◎ Use arithmetic operators in expressions
- ◎ Use assignment operators in assignment statements
- ◎ Write C++ comments
- ◎ Write programs using sequential statements and interactive input statements

In this chapter, you learn about writing programs that use variables, constants, and arithmetic operators. You also learn about programs that receive interactive input from a user of your programs. You begin by reviewing variables and constants and learning how to use them in a C++ program. You should do the exercises and labs in this chapter only after you have finished Chapter 2 and Chapter 3 of your book, *Programming Logic and Design, Eighth Edition*, by Joyce Farrell.

12

Variables

As you know, a **variable** is a named location in the computer's memory whose contents can vary (thus the term *variable*). You use a variable in a program when you need to store values. The values stored in variables often change as a program executes.

In C++, you must declare variables before you can use them in a program. Declaring a variable is a two-part process: first, you give the variable a name, and then you specify its data type. You will learn about data types shortly. But first, you will focus on the rules for naming variables in C++.

Variable Names

Variable names in C++ can consist of letters, numerical digits, and the underscore character, but they cannot begin with a digit. Also, you cannot use a C++ keyword for a variable name. As you learned in Chapter 1 of this book, a keyword is a word with a special meaning in C++. The following are all examples of legal variable names in C++: `my_var`, `num6`, `intValue`, and `firstName`. Table 2-1 lists some examples of invalid variable names and explains why each is invalid.



A variable is sometimes referred to as an identifier.

Name of Variable	Explanation
<code>3 wrong</code>	Invalid because it begins with a digit and includes a space
<code>\$don't</code>	Invalid because it contains a single quotation mark and begins with a dollar sign
<code>int</code>	Invalid because it is a C++ keyword
<code>first name</code>	Invalid because it contains a space

Table 2-1 Invalid variable names

When naming variables, keep in mind that C++ is case sensitive—in other words, C++ knows the difference between uppercase and lowercase characters. That means `value`, `Value`, and `ValuE` are three different variable names in C++.

In C++, variable names can be as long as you want. A good rule is to give variables meaningful names that are long enough to describe how the variable is used, but they should not be so long that you make your program hard to read or cause yourself unnecessary typing. For example, a variable named `firstName` will clearly be used to store someone's first name. The variable name `freshmanStudentFirstName` is descriptive but inconveniently long; the variable name `fn` is too short and not meaningful.

One of the naming conventions used by C++ programmers is called **camel case**. This means:

- Variable names are made up of multiple words, with no spaces between them.
- The first character in the variable name is lowercase.
- The first character of each word after the first word is a capitalized character.

Examples of C++ variable names in camel case include `firstName`, `myAge`, and `salePrice`. You do not include spaces between the words in a variable name.

C++ Data Types

In addition to specifying a name for a variable, you also need to specify a particular data type for that variable. A variable's **data type** dictates the amount of memory that is allocated for the variable, the type of data you can store in the variable, and the types of operations that can be performed on the variable. There are many different kinds of data types, but this book focuses on the most basic kind of data types, known as **primitive data types**. There are five primitive data types in C++: `int`, `float`, `double`, `char`, and `bool`. Some of these data types (such as `int`, `double`, and `float`) are used for variables that store numeric values, and they are referred to as numeric data types. The others have specialized purposes. For example, the `bool` data type is used to store a value of either `true` or `false`, and the `char` data type is used to store a single character.



You used the `int` data type in Chapter 1 as the return type for the `main()` function in the Hello World program.

In the programs you write for this book, you will not use all of the primitive data types found in C++. Instead, you will focus on two of the numeric data types (`int` and `double`). The `int` data type is used for values that are whole numbers. For example, you could use a variable with the data type `int` to store someone's age (for example, 25) or the number of students in a class (for example, 35). A variable of the `int` data type consists of 32 bits (4 bytes) of space in memory.



The actual size of the built-in data types may be different on different computers, but the sizes noted in this book indicate the usual sizes on a 32-bit or 64-bit computer.

You use the data type `double` to store a floating-point value (that is, a fractional value), such as the price of an item (2.95) or a measurement (2.5 feet). A variable of the `double` data type occupies 64 bits (8 bytes) of space in memory. You will learn about using other data types as you continue to learn more about C++ in subsequent courses.



In *Programming Logic and Design*, the data type `num` is used to refer to all numeric data types. That book does not make a distinction between `int` and `double` because the pseudocode used in the book is not specific to any one programming language. However, in C++ this distinction is always maintained.

14

The `int` and `double` data types are adequate for all the numeric variables you will use in this book. But what about when you need to store a group of characters (such as a person's name) in a variable? In programming, you refer to a group of one or more characters as a **string**. An example of a string is the last name *Wallace* or a product type such as a *desk*. There is no primitive data type in C++ for storing strings; instead, they are stored in an object known as a **string** object. In addition to working with the `int` and `double` data types in this book, you will also work with **strings**.

Exercise 2-1: Using C++ Variables, Data Types, and Keywords

In this exercise, you use what you have learned about naming variables, data types, and keywords to answer Questions 1–2.

1. Is each of the following a legal C++ variable name? (Answer yes or no.)

myIDNumber	_____	this_is_a_var	_____	AMOUNT	_____
yourIDNumber	_____	amount	_____	\$amount	_____
int	_____	10amount	_____	doubleAmount	_____
July31	_____	one amount	_____	Amount	_____

2. What data type (`int`, `double`, or `string`) is appropriate for storing each of the following values?

A person's height (in inches) _____

The amount of interest on a loan, such as 1 percent _____

The amount of your car payment _____

Your grandmother's first name _____

The number of siblings you have _____

Declaring and Initializing Variables

Now that you understand the rules for naming a variable, and you understand the concept of a data type, you are ready to learn how to declare a variable. In C++, you must declare all variables before you can use them in a program. When you **declare** a variable, you tell the compiler that you are going to use the variable. In the process of declaring a variable, you must specify the variable's name and its data type. Declaring a variable tells the compiler that it needs to reserve a memory location for the variable. A line of code that declares a variable is known as a **variable declaration**. The C++ syntax for a variable declaration is as follows:

```
dataType variableName;
```

For example, the declaration statement `int counter;` declares a variable named `counter` of the `int` data type. The compiler reserves the amount of memory space allotted to an `int` variable (32 bits, or 4 bytes) for the variable named `counter`. The compiler then assigns the new variable a specific memory address. In Figure 2-1, the memory address for the variable named `counter` is 1000, although you would not typically know the memory address of the variables included in your C++ programs.

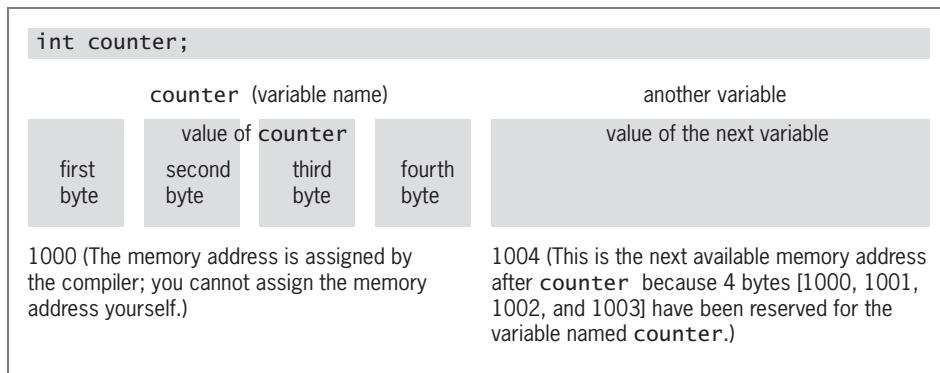


Figure 2-1 Declaration of variable and memory allocation



In C++, variables are not automatically initialized with a value. They contain undetermined values unless you explicitly provide a value.

You can also initialize a C++ variable when you declare it. When you **initialize** a C++ variable, you give it an initial value. For example, you can assign an initial value of 8 to the `counter` variable when you declare it, as shown in the following code:

```
int counter = 8;
```

You can also declare and initialize variables of data type `double` and `string` variables (objects) as shown in the following code:

```
double salary;
double cost = 12.95;
string firstName;
string homeAddress = "123 Main Street";
```

You can declare more than one variable in one statement as long as they have the same data type. For example, the following statement declares two variables, named `counter` and `value`. Both variables are of the `int` data type.

```
int counter, value;
```

Exercise 2-2: Declaring and Initializing C++ Variables

In this exercise, you use what you have learned about declaring and initializing variables to answer Questions 1–2.

16

1. Write a C++ variable declaration for each of the following. Use `int`, `double`, or `string`, and choose meaningful variable names.

Declare a variable to store a student ID number (1 to 1000). _____

Declare a variable to store the number of inches in a foot. _____

Declare a variable to store a temperature (0.0 to 130.0). _____

Declare a variable to store the name of the state you live in. _____

2. Declare and initialize variables to represent the following values. Use `int`, `double`, or `string`, and choose meaningful variable names.

Your hat size (7.5) _____

The number of days in a leap year _____

The name of your cat, Mindre _____

Your cell phone number (999-999-9999) _____

Lab 2-1: Declaring and Initializing C++ Variables

In this lab, you declare and initialize variables in a C++ program provided with the data files for this book. The program, which is saved in a file named `NewAge.cpp`, calculates your age in the year 2050.

1. Open the source code file named `NewAge.cpp` using Notepad or the text editor of your choice.
 2. Declare an integer variable named `myNewAge`.
 3. Declare and initialize an integer variable named `myCurrentAge`. Initialize this variable with your current age.
 4. Declare and initialize an integer variable named `currentYear`. Initialize this variable with the value of the current year. Use four digits for the year.
 5. Save this source code file in a directory of your choice, and then make that directory your working directory.
 6. Compile the source code file `NewAge.cpp`.
 7. Execute the program. Record the output of this program.
-
-

Constants

As you know, a **constant** is a value that never changes. In C++, you can use both unnamed constants as well as named constants in a program. You will learn about named constants shortly. But first, this section focuses on unnamed constants in C++.

17

Unnamed Constants

Computers are able to deal with two basic types of data: text and numeric. When you use a specific numeric value, such as 35, in a program, you write it using the numerals, without quotation marks. A specific numeric value is called a **numeric constant** because it does not change; a 35 always has the value 35. When you use a specific text value, or string of characters, such as “William”, you enclose the **string constant** in double quotation marks. Both of the preceding examples, 35 and “William”, are examples of **unnamed constants** because they do not have specified names as variables do.

Named Constants

In addition to variables, C++ allows you to create named constants. A **named constant** is similar to a variable, except it can be assigned a value only once. You use a named constant when you want to assign a name to a value that will never be changed when a program executes.

To declare a named constant in C++, you use the keyword `const` followed by the data type, followed by the name of the constant. Named constants must be initialized when they are declared, and their contents may not be changed during the execution of the program. For example, the following statement declares an `int` constant named `MAX_STUDENTS` and initializes `MAX_STUDENTS` with the value 35.

```
const int MAX_STUDENTS = 35;
```



By convention, in C++ the names of constants are written in all uppercase letters, and multiple words are separated by the underscore character. This makes it easier for you to spot named constants in a long block of code.

Exercise 2-3: Declaring and Initializing C++ Constants

In this exercise, you use what you have learned about declaring and initializing constants to answer the following question.

1. Declare and initialize constants to represent the following values. Use `int`, `double`, or `string`, and choose meaningful names.

The price of a pizza is \$14.95. _____

The number of days in September is 30. _____

The name of your cat, Yogi. _____

The height of an NBA basketball hoop is 10 feet. _____

Lab 2-2: Declaring and Initializing C++ Constants

In this lab, you declare and initialize constants in a C++ program provided with the data files for this book. The program, which is saved in a file named `NewAge2.cpp`, calculates your age in the year 2050.

18

1. Open the source code file named `NewAge2.cpp` using Notepad or the text editor of your choice.
 2. Declare a constant named `YEAR`, and initialize `YEAR` with the value 2050.
 3. Edit the statement `myNewAge = myCurrentAge + (2050 - currentYear)` so it uses the constant named `YEAR`.
 4. Edit the statement `cout << "I will be " << myNewAge << " in 2050." << endl;` so it uses the constant named `YEAR`.
 5. Save this source code file as `NewAge2.cpp` in a directory of your choice, and then make that directory your working directory.
 6. Compile the source code file `NewAge2.cpp`
 7. Execute the program. Record the output of this program.
-
-

Arithmetic and Assignment Operators

After you declare a variable, you can use it in various tasks. For example, you can use variables in simple arithmetic calculations, such as adding, subtracting, and multiplying. You can also perform other kinds of operations with variables, such as comparing one variable to another to determine which is greater.

To write C++ code that manipulates variables in this way, you need to be familiar with operators. An **operator** is a symbol that tells the computer to perform a mathematical or logical operation. C++ has a large assortment of operators. The discussion begins with a group of operators known as the arithmetic operators.

Arithmetic Operators

Arithmetic operators are the symbols used to perform arithmetic calculations. You are probably already very familiar with the arithmetic operators for addition (+) and subtraction (-). Table 2-2 lists and explains the arithmetic operators found in C++.

Operator Name and Symbol	Example	Comment
Addition +	num1 + num2	
Subtraction -	num1 - num2	
Multiplication *	num1 * num2	
Division /	15/2	Integer division; result is 7; fraction is truncated
	15.0 / 2.0	Floating point division; result is 7.5
	15.0 / 2	Floating point division because one of the operands is a floating point number; result is 7.5
Modulus %	hours % 24	Performs division and finds the remainder; result is 1 if the value of hours is 25
Unary plus +	+num1	Maintains the value of the expression; if the value of num1 is 3, then +num1 is 3
Unary minus -	-(num1 - num2)	If value of (num1 - num2) is 10, then -(num1 - num2) is -10.

Table 2-2 C++ arithmetic operators

You can combine arithmetic operators and variables to create **expressions**. The computer evaluates each expression, and the result is a value. To give you an idea of how this works, assume that the value of num1 is 3 and num2 is 20, and that both are of data type `int`. With this information in mind, study the examples of expressions and their values in Table 2-3.

Expression	Value	Explanation
num1 + num2	23	$3 + 20 = 23$
num1 - num2	-17	$3 - 20 = -17$
num2 % num1	2	$20 / 3 = 6$ remainder 2
num1 * num2	60	$3 * 20 = 60$
num2 / num1	6	$20 / 3 = 6$ (remainder is truncated)
-num1	-3	Value of num1 is 3, therefore -num1 is -3

Table 2-3 Expressions and values

Assignment Operators and the Assignment Statement

Another type of operator is an **assignment operator**. You use an assignment operator to assign a value to a variable. A statement that assigns a value to a variable is known as an **assignment statement**. In C++, there are several types of assignment operators. The one you

will use most often is the = assignment operator, which simply assigns a value to a variable. Table 2-4 lists and explains some of the assignment operators found in C++.

Operator Name and Symbol	Example	Comment
Assignment =	count = 5;	Places the value on the right side into the memory location named on the left side
Initialization =	int count = 5;	Places the value on the right side into the memory location named on the left side when the variable is declared
Assignment += -= *= /= %= +=	num += 20; num -= 20; num *= 20; num /= 20; num %= 20;	Equivalent to num = num + 20; Equivalent to num = num - 20; Equivalent to num = num * 20; Equivalent to num = num / 20; Equivalent to num = num % 20;

Table 2-4 C++ assignment operators

When an assignment statement executes, the computer evaluates the expression on the right side of the assignment operator and then assigns the result to the memory location associated with the variable named on the left side of the assignment operator. An example of an assignment statement is shown in the following code. Notice that the statement ends with a semicolon. In C++, assignment statements always end with a semicolon.

```
answer = num1 * num2;
```

This assignment statement causes the computer to evaluate the expression `num1 * num2`. After evaluating the expression, the computer stores the results in the memory location associated with `answer`. If the value stored in the variable named `num1` is 3, and the value stored in the variable named `num2` is 20, then the value 60 is assigned to the variable named `answer`.

Here is another example:

```
answer += num1;
```

This statement is equivalent to the following statement:

```
answer = answer + num1;
```

If the value of `answer` is currently 10 and the value of `num1` is 3, then the expression on the right side of the assignment statement `answer + num1;` evaluates to 13, and the computer assigns the value 13 to `answer`.

Precedence and Associativity

Once you start to write code that includes operators, you need to be aware of the order in which a series of operations is performed. In other words, you need to be aware of the

precedence of operations in your code. Each operator is assigned a certain level of precedence. For example, multiplication has a higher level of precedence than addition. So in the expression $3 * 7 + 2$, the $3 * 7$ is multiplied first; only after the multiplication is completed would the 2 be added.

But what happens when two operators have the same precedence? The rules of **associativity** determine the order in which operations are evaluated in an expression containing two or more operators with the same precedence. For example, in the expression $3 + 7 - 2$, the addition and subtraction operators have the same precedence. As shown in Table 2-5, the addition and subtraction operators have left-to-right associativity, which causes the expression to be evaluated from left to right ($3 + 7$ added first; then 2 is subtracted). Table 2-5 shows the precedence and associativity of the operators discussed in this chapter.

Operator Name	Operator Symbol(s)	Order of Precedence	Associativity
Parentheses	()	First	Left to right
Unary	- +	Second	Right to left
Multiplication, division, and modulus	* / %	Third	Left to right
Addition and subtraction	+ -	Fourth	Left to right
Assignment	= += -= *= /= %=	Fifth	Right to left

Table 2-5 Order of precedence and associativity

As you can see in Table 2-5, the parentheses operator, (), has the highest precedence. You use this operator to change the order in which operations are performed. Note the following example:

```
average = test1 + test2 / 2;
```

The task of this statement is to find the average of two test scores. The way this statement is currently written, the compiler will divide the value in the `test2` variable by 2, and then add it to the value in the `test1` variable. So, for example, if the value of `test1` is 90 and the value of `test2` is 88, then the value assigned to `average` will be 134, which is obviously not the correct average of these two test scores. By using the parentheses operator in this example, you can force the addition to occur before the division. The correct statement looks like this:

```
average = (test1 + test2) / 2;
```

In this example, the value of `test1`, 90, is added to the value of `test2`, 88, and then the sum is divided by 2. The value assigned to `average`, 89, is the correct result.

Exercise 2-4: Understanding Operator Precedence and Associativity

In this exercise, you use what you have learned about operator precedence and associativity in C++. Study the following code, and then answer Questions 1–2.

22

```
// This program demonstrates the precedence and
// associativity of operators.
#include <iostream>
using namespace std;
int main()
{
    int number1 = 20;
    int number2 = 5;
    int number3 = 17;
    int answer1, answer2, answer3;
    int answer4, answer5, answer6;
    answer1 = number1 * number2 + number3;
    cout << "Answer 1: " << answer1 << endl;
    answer2 = number1 * (number2 + number3);
    cout << "Answer 2: " << answer2 << endl;
    answer3 = number1 + number2 - number3;
    cout << "Answer 3: " << answer3 << endl;
    answer4 = number1 + (number2 - number3);
    cout << "Answer 4: " << answer4 << endl;
    answer5 = number1 + number2 * number3;
    cout << "Answer 5: " << answer5 << endl;
    answer6 = number3 / number2;
    cout << "Answer 6: " << answer6 << endl;
    return 0;
}
```

1. What are the values of `answer1`, `answer2`, `answer3`, `answer4`, `answer5`, and `answer6`?

2. Explain how precedence and associativity affect the result.

Lab 2-3: Understanding Operator Precedence and Associativity

In this lab, you complete a partially written C++ program that is provided in the data files for this book. The program, which was written for a furniture company, prints the name of the furniture item, its retail price, its wholesale price, the profit made on the piece of furniture, a sale price, and the profit made when the sale price is used.

1. Open the file named `Furniture.cpp` using Notepad or the text editor of your choice.
2. The file includes variable declarations and output statements. Read them carefully before you proceed to the next step.
3. Design the logic and write the C++ code that will use assignment statements to first calculate the profit, then calculate the sale price, and finally calculate the profit when the sale price is used. Profit is defined as the retail price minus the wholesale price. The sale price is 25 percent deducted from the retail price. The sale profit is defined as the sale price minus the wholesale price. Perform the appropriate calculations as part of your assignment statements.
4. Save the source code file in a directory of your choice, and then make that directory your working directory.
5. Compile the program.
6. Execute the program. Your output should be as follows:

```
Item Name: TV Stand
Retail Price: $325
Wholesale Price: $200
Profit: $125
Sale Price: $243.75
Sale Profit: $43.75
```



In Chapter 9 of this book, you will learn how to control the number of places after the decimal point when you want to create output with floating-point values.

Next, you will see how to put together all you have learned in this chapter to write a C++ program that uses sequential statements, comments, and interactive input statements.

Sequential Statements, Comments, and Interactive Input Statements

The term **sequential statements** (or **sequence**) refers to a series of statements that must be performed in sequential order, one after another. You use a sequence in programs when you want to perform actions one after the other. A sequence can contain any number of actions, but those actions must be in the proper order, and no action in the sequence can be skipped. Note that a sequence can contain comments, which are not considered part of the sequence itself.

Comments serve as documentation, explaining the code to the programmer and any other people who might read it. In Chapter 2 of *Programming Logic and Design*, you learned about program comments, which are statements that do not execute. You use comments in C++ programs to explain your logic to people who read your source code. The C++ compiler ignores comments.



You are responsible for including well-written, meaningful comments in all of the programs that you write. In fact, some people think that commenting your source code is as important as the source code itself.

24

You can choose from two commenting styles in C++. In the first, you type two forward slash characters // at the beginning of each line that you want the compiler to ignore. This style is useful when you only want to mark a single line as a comment. In the second style, you enclose a block of lines with the characters /* and */. This style is useful when you want to mark several lines as a comment. You may place comments anywhere in a C++ program.

The C++ program in the following example shows both styles of comments included in the Temperature program. The first six lines of the program make up a multiline block comment that explains some basic information about the program. Additionally, several single-line comments are included throughout to describe various parts of the program.

A sequence often includes **interactive input statements**, which are statements that ask, or **prompt**, the user to input data. The C++ program in the following example uses sequential statements and interactive input statements to convert a Fahrenheit temperature to its Celsius equivalent:

```
/* Temperature.cpp - This program converts a Fahrenheit
   temperature to Celsius.
   Input: Interactive
   Output: Fahrenheit temperature followed by Celsius
   temperature
*/
#include <iostream>
using namespace std;
int main()
{
    double fahrenheit;
    double celsius;
    // Prompt user
    cout << "Enter Fahrenheit temperature: ";
    // Get interactive user input
    cin >> fahrenheit;
    // Calculate celsius
    celsius = (fahrenheit - 32.0) * (5.0 / 9.0);
    // Output
    cout << "Fahrenheit temperature:" << fahrenheit << endl;
    cout << "Celsius temperature:" << celsius << endl;
    return 0;
}
```

This program is made up of sequential statements that execute one after the other. It also includes comments explaining the code. The comments are the lines enclosed within the /* and */ characters, as well as those lines that begin with //. After the variables `fahrenheit` and `celsius` are declared (using the `double` data type), the following statements execute:

```
// Prompt user
cout << "Enter Fahrenheit temperature: ";
// Get interactive user input
cin >> fahrenheit;
```

The `cout` statement is used to prompt the user for the Fahrenheit temperature so the program can convert it to Celsius. Notice that `endl` is not included in this `cout` statement. This is done to position the cursor on the same line as the displayed prompt.



Remember that `endl` causes the cursor to be displayed on the line following the output.

The next statement, `cin >> fahrenheit;`, is used to retrieve the user's input. This line consists of multiple parts. The first part, `cin`, is the name of an object that represents the user's standard input device, which is usually a keyboard. Next, you see `>>`, which is called the **extraction** or **input** operator. After `>>`, you see `fahrenheit`, which is the name of the variable that will store the data that is going to be extracted. The C++ extraction operator `>>` automatically converts input typed at the keyboard to the appropriate data type. In this example, `>>` will convert the user's input to the `double` data type because the variable `fahrenheit` is declared as data type `double`.

The next statement to execute is an assignment statement, as follows:

```
celsius = (fahrenheit - 32.0) * (5.0 / 9.0);
```

The formula that converts Fahrenheit temperatures to Celsius is used on the right side of this assignment statement. Notice the use of parentheses in the expression to control precedence. The expression is evaluated, and the resulting value is assigned to the variable named `celsius`.

Notice that the division uses the values 5.0 and 9.0. This is an example of floating-point division, which results in a value that includes a fraction. If the values 5 and 9 were used, integer division would be performed, and the fractional portion would be truncated.

The next two statements to execute in sequence are both output statements, as follows:

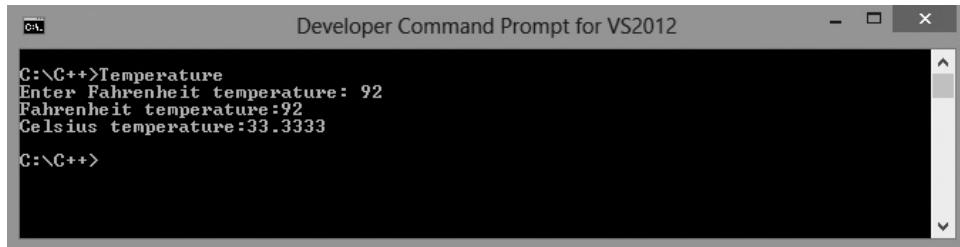
```
cout << "Fahrenheit temperature:" << fahrenheit << endl;
cout << "Celsius temperature:" << celsius << endl;
```

The statement `cout << "Fahrenheit temperature:" << fahrenheit << endl;` is used to display the string `Fahrenheit temperature:` followed by the value stored in the variable `fahrenheit`, followed by a newline character. The second output statement displays the words `Celsius temperature:` followed by the value stored in the variable `celsius`, followed by a newline character.

The last statement in this program is `return 0;`. As you learned in Chapter 1 of this book, this statement instructs the compiler to return the value 0 from the `main()` function.

This program is saved in a file named `Temperature.cpp` and is included in the student files for this chapter. You can see the output produced by the Temperature program in Figure 2-2.

Now that you have seen a complete C++ program that uses sequential statements and interactive input statements, it is time for you to begin writing your own programs.

A screenshot of a Windows command prompt window titled "Developer Command Prompt for VS2012". The window shows the output of a C++ program named "Temperature". The user has entered "Enter Fahrenheit temperature: 92" and the program has output "Fahrenheit temperature:92" and "Celsius temperature:33.3333". The command prompt prompt is "C:\>".

```
C:\>Temperature
Enter Fahrenheit temperature: 92
Fahrenheit temperature:92
Celsius temperature:33.3333
C:\>
```

Figure 2-2 Output from Temperature.cpp program

Exercise 2-5: Understanding Sequential Statements

In this exercise, you use what you have learned about sequential statements to read a scenario and then answer Questions 1–4.

Suppose you have written a C++ program that calculates the amount of paint you need to cover the walls in your family room. Two walls are 9 feet high and 19.5 feet wide. The other two walls are 9 feet high and 20.0 feet wide. The salesperson at the home improvement store told you to buy 1 gallon of paint for every 150 square feet of wall you need to paint. Suppose you wrote the following code, but your program is not compiling. Take a few minutes to study this code, and then answer Questions 1–4.

```
// This program calculates the number of gallons of paint needed.
#include <iostream>
using namespace std;
int main()
{
    double height1 = 9;
    double height2 = 9;
    int width1 = 19.5;
    double width2 = 20.0;
    double squareFeet;
    int numGallons;
    numGallons = squareFeet / 150;
    squareFeet = (width1 * height1 + width2 * height2) * 2;
    cout << "Number of Gallons: " << numGallons << endl;
    return 0;
}
```

- When you compile this program, you receive a warning message from the cl compiler that is similar to the following:

```
c:\users\jo ann\c++ pal\chapter_2\student\paint.cpp(12) :
warning C4700: uninitialized local variable 'squareFeet'
used
```

There are multiple problems with this program even though the compiler issued only one warning message, which is pointing out an uninitialized variable on line number 12 in the source code file. You would know the compiler is complaining about line 12 because of the (12) in the warning message. On the following lines, describe how to fix all of the problems you can identify.

2. You have two variables declared in this program to represent the height of your walls, `height1` and `height2`. Do you need both of these variables? If not, how would you change the program? Be sure to identify all of the changes you would make.

Lab 2-4: Understanding Sequential Statements

In this lab, you complete a C++ program provided with the data files for this book. The program calculates the amount of tax withheld from an employee's weekly salary, the tax deduction to which the employee is entitled for each dependent, and the employee's take-home pay. The program output includes state tax withheld, federal tax withheld, dependent tax deductions, salary, and take-home pay.

1. Open the source code file named `Payroll1.cpp` using Notepad or the text editor of your choice.
2. Variables have been declared and initialized for you as needed, and the input and output statements have been written. Read the code carefully before you proceed to the next step.
3. Write the C++ code needed to perform the following:
 - a. Calculate state withholding tax at 6.5 percent, and calculate federal withholding tax at 28.0.
 - b. Calculate dependent deductions at 2.5 percent of the employee's salary for each dependent.
 - c. Calculate total withholding. (Total withholding is the total state withholding combined with the total federal withholding.)
 - d. Calculate take-home pay as salary minus total withholding plus deductions.
4. Save this source code file in a directory of your choice, and then make that directory your working directory.

5. Compile the program.
6. Execute the program. You should get the following output:

State Tax: \$81.25
Federal Tax: \$350
Dependents: \$62.5
Salary: \$1250
Take-Home Pay: \$881.25

28



In Chapter 9 of this book, you will learn how to control the number of places after the decimal point when you want to output floating-point values.

7. In this program, the variables `salary` and `numDependents` are initialized with the values 1250.00 and 2. To make this program more flexible, modify it to accept interactive input for `salary` and `numDependents`. Name the modified version `Payroll12.cpp`.

CHAPTER

3

Writing Structured C++ Programs

After studying this chapter, you will be able to:

- ◎ Use structured flowcharts and pseudocode to write structured C++ programs
- ◎ Write simple modular programs in C++

In this chapter, you begin to learn how to write structured C++ programs. As you will see, creating a flowchart and writing pseudocode before you actually write the program ensures that you fully understand the program's intended design. You begin by looking at a structured flowchart and pseudocode from your text, *Programming Logic and Design, Eighth Edition*. You should do the exercises and labs in this chapter only after you have finished Chapters 2 and 3 of that book.

30

Using Flowcharts and Pseudocode to Write a C++ Program

In the first three chapters of *Programming Logic and Design*, you studied flowcharts and pseudocode for the Number-Doubling program. Figure 3-1 shows the functional, structured flowchart and pseudocode for this program.

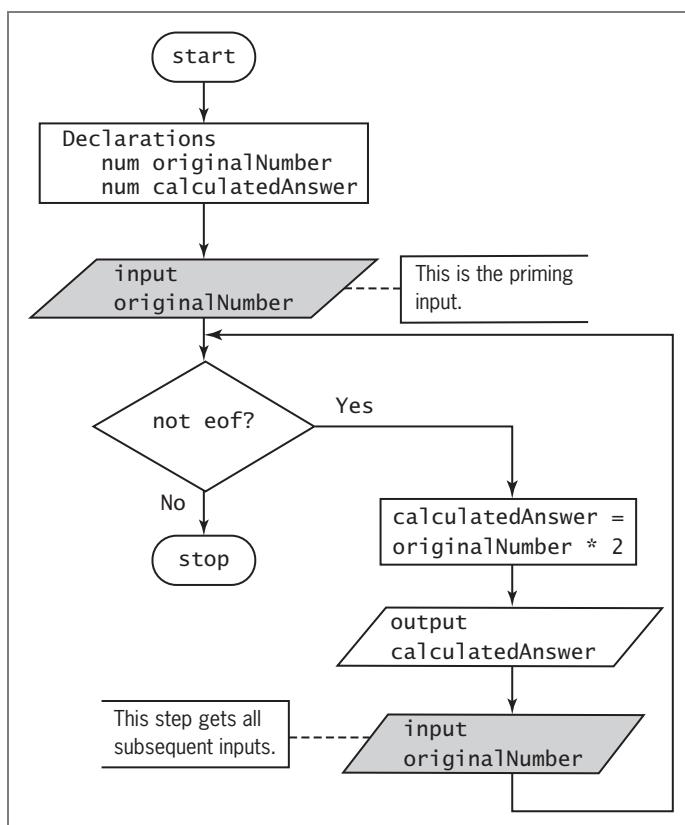


Figure 3-1 Functional, structured flowchart and pseudocode for the Number-Doubling program

By studying the flowchart and pseudocode, you can see that this program makes use of the sequence and loop structures you were introduced to in *Programming Logic and Design*. The remainder of this section walks you through the C++ code for this program. The explanations assume that you are simply reading along, but if you want, you can type the code as it is presented. The goal of this section is to help you get a feel for how flowcharts and pseudocode

can serve as a guide when you write C++ programs. You must learn more about C++ before you can expect to write this program by yourself.

In Figure 3-1, the first line of the pseudocode is the word **start**. How do you translate this pseudocode command into the C++ code that will start the Number-Doubling program? In Chapter 1 of this book, you learned that to start a C++ program, you first create a **main()** function. So to start the Number-Doubling program, you will first create a function named **main()** because it is always the first function that executes in a C++ program. Thus, the code that follows starts the Number-Doubling program by creating a **main()** function:

```
int main()
{
}
```



Notice that each opening curly brace is matched by a closing curly brace.



If you are typing the code as it is presented here, save the program in a file that has an appropriate name, such as **NumberDouble.cpp**. The complete program is also saved in a file named **NumberDouble.cpp** and is included in the student files for this chapter.

Next, you see that two variables, **originalNumber** and **calculatedAnswer**, are declared as data type **num**. The C++ code that follows adds the variable declarations, with the declarations shown in bold.

```
int main()
{
    int originalNumber;
    int calculatedAnswer;
}
```

The next line of the pseudocode instructs you to input the **originalNumber**. In other words, you need to write the input statement that primes the loop. You learned about priming read statements in Chapter 3 of *Programming Logic and Design*. In Chapter 2 of this book, you learned how to use interactive input statements in programs to allow the user to input data. You also learned to prompt the user by explaining what the program expects to receive as input. The following example includes the code that implements the priming read by displaying a prompt for the user and then retrieving the number the user wants doubled.

Note that the code in boldface has been added to the Number-Doubling program. If you were writing this code yourself, you would start by writing the code for the Number-Doubling program as shown above, and then edit it to add the boldface code shown here:

```
int main()
{
    int originalNumber;
    int calculatedAnswer;
    cout << "Enter a number to double: " << endl;
    cin >> originalNumber;
}
```

32



You have not learned enough about `while` loops to write this code yourself, but you can observe how it is done in this example. You will learn more about loops in Chapter 5 of this book.

Next, the pseudocode instructs you to begin a `while` loop with `eof` (end of file) used as the condition to exit the loop.

Since you are using interactive input in this program, it requires no `eof` marker. Instead you will use the number 0 (zero) to indicate the end of input. You use 0 because 0 doubled will always be 0. The use of 0 to indicate the end of input also requires you to change the prompt to tell the user how to end the program. Review the following code. Again, the newly added code is formatted in bold.

```
int main()
{
    int originalNumber;
    int calculatedAnswer;
    cout << "Enter a number to double or 0 to end: " << endl;
    cin >> originalNumber;
    while(originalNumber != 0)
    {
    }
}
```



Notice that a beginning curly brace ({) and an ending curly brace (}) are used in C++ to mark the beginning and end of code that executes as part of a loop.

According to the pseudocode, the body of the loop is made up of three sequential statements. The first statement calculates the `originalNumber` multiplied by 2, the second statement prints the `calculatedAnswer`, and the third statement retrieves the next `originalNumber` from the user. In C++, you actually need to add an additional statement between the curly braces that mark the body of the `while` loop. This additional statement prompts the user to enter the next number to be doubled.

In the following example, the code that makes up the body of the loop is in bold:

```
int main()
{
    int originalNumber;
    int calculatedAnswer;
    cout << "Enter a number to double or 0 to end: " << endl;
    cin >> originalNumber;
    while(originalNumber != 0)
    {
        calculatedAnswer = originalNumber * 2;
```

```
    cout << originalNumber << " doubled is "
    << calculatedAnswer << endl;
    cout << "Enter a number to double or 0 to end: " << endl;
    cin >> originalNumber;
}
}
```

The last line of the pseudocode instructs you to end the program. In C++, the closing curly brace (}) for the `main()` function signifies the end of the `main()` function and therefore the end of the program. Note that the preceding code includes two closing curly braces. The last one is the one that ends the `main()` function and the second to last one ends the `while` loop.

You are almost finished translating the pseudocode into C++ code. The code shown in bold below is the `#include` preprocessor directive and the `using` statement that you learned about in Chapter 1 of this book. Additionally, you see the `return 0;` statement. This statement instructs the compiler to return the value 0 from the `main()` function.

```
#include <iostream>
using namespace std;
int main()
{
    int originalNumber;
    int calculatedAnswer;
    cout << "Enter a number to double or 0 to end: " << endl;
    cin >> originalNumber;
    while(originalNumber != 0)
    {
        calculatedAnswer = originalNumber * 2;
        cout << originalNumber << " doubled is "
            << calculatedAnswer << endl;
        cout << "Enter a number to double or 0 to end: " << endl;
        cin >> originalNumber;
    }
    return 0;
}
```

At this point, the program is ready to be compiled. Assuming there are no syntax errors, it should execute as planned. Figure 3-2 displays the input and output of the Number-Doubling program.

```
C:\>Developer Command Prompt for VS2012
C:\>NumberDouble
Enter a number to double or 0 to end:
33
33 doubled is 66
Enter a number to double or 0 to end:
22
22 doubled is 44
Enter a number to double or 0 to end:
11
11 doubled is 22
Enter a number to double or 0 to end:
0
```

Figure 3-2 Number Double program input and output

Although you have not learned everything you need to know to write this program yourself, you can see from this example that writing the program in C++ is easier if you start with a well designed, functional, structured flowchart or pseudocode.

34

Lab 3-1: Using Flowcharts and Pseudocode to Write a C++ Program

In this lab, you use the flowchart and pseudocode in Figure 3-3 to add code to a partially created C++ program. When completed, college admissions officers should be able to use the C++ program to determine whether to accept or reject a student, based on his or her test score and class rank.

```
start
    input testScore, classRank
    if testScore >= 90 then
        if classRank >= 25 then
            output "Accept"
        else
            output "Reject"
        endif
    else
        if testScore >= 80 then
            if classRank >= 50 then
                output "Accept"
            else
                output "Reject"
            endif
        else
            if testScore >= 70 then
                if classRank >= 75 then
                    output "Accept"
                else
                    output "Reject"
                endif
            else
                output "Reject"
            endif
        endif
    endif
stop
```

Figure 3-3 Pseudocode for the College Admission program

1. Study the pseudocode in Figure 3-3.
2. Open the source code file named `CollegeAdmission.cpp` using Notepad or the text editor of your choice.
3. Declare two integer variables named `testScore` and `classRank`.

4. Write the interactive input statements to retrieve a student's test score and class rank from the user of the program. Do not forget to prompt the user for the test score and class rank.
 5. The rest of the program is written for you. Save this source code file in a directory of your choice, and then make that directory your working directory.
 6. Compile the source code file `CollegeAdmission.cpp`.
 7. Execute the program by entering 87 for the test score and 60 for the class rank. Record the output of this program.
-
8. Execute the program by entering 60 for the test score and 87 for the class rank. Record the output of this program.
-

Writing a Modular Program in C++

In Chapter 2 of your book, *Programming Logic and Design*, you learned about local and global variables and named constants. To review briefly, you declare **local variables** and local constants within the module—or, in C++ terminology, the function—that uses them. Further, you can only use a local variable or a local constant within the function in which it is declared. **Global variables** and global constants are known to the entire program. They are declared at the program level and are visible to and usable in all the functions called by the program. It is not considered a good programming practice to use global variables, so the C++ program below uses local variables (as well as local constants). A good reason for using local variables and local constants is that source code is easier to understand when variables are declared where they are used. In addition, global variables can be accessed and altered by any part of the program, which can make the program difficult to read and maintain; it also makes the program prone to error.

Recall from Chapter 2 that most programs consist of a `main()` module or function that contains the mainline logic. The `main()` function then calls other methods to get specific work done in the program. The mainline logic of most procedural programs follows this general structure:

1. Declarations of variables and constants
2. **Housekeeping tasks**, such as displaying instructions to users, displaying report headings, opening files the program requires, and inputting the first data item
3. **Detail loop tasks** that do the main work of the program, such as processing many records and performing calculations
4. **End-of-job tasks**, such as displaying totals and closing any open files

In Chapter 2 of *Programming Logic and Design*, you studied a flowchart for a modular program that prints a payroll report with global variables and constants. This flowchart is shown in Figure 3-4.

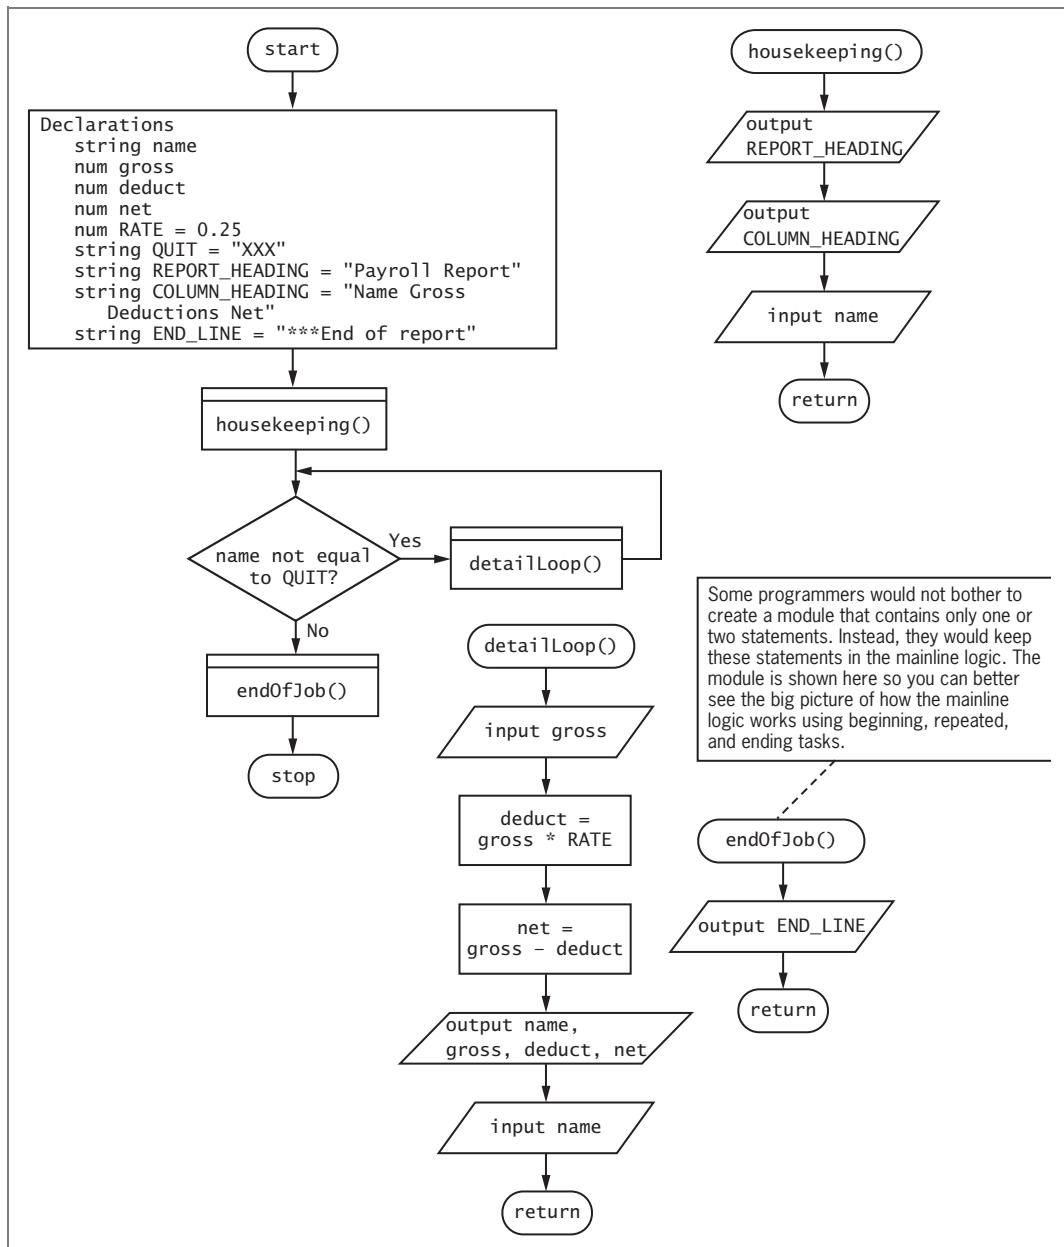


Figure 3-4 Flowchart for the Payroll Report program

In this section, you walk through the process of creating a C++ program that implements the logic illustrated in the flowchart in Figure 3-4. According to the flowchart, the program begins with the execution of the mainline method. The mainline method declares four global variables (`name`, `gross`, `deduct`, and `net`) and five global constants (`RATE`, `QUIT`, `REPORT_HEADING`, `COLUMN_HEADING`, and `END_LINE`).

The C++ code that follows shows the Payroll Report program with the `main()` function, as well as the variable and constant declarations. Notice that a `#include` preprocessor directive you have not seen before has been added at the beginning of the program. The preprocessor directive `#include <string>` allows you to use `string` objects in your C++ program. You will learn more about `string` objects in subsequent chapters.

```
#include <iostream>
#include <string>
using namespace std;
int main()
{
    // Declare variables and constants local to main()
    string name;
    double gross, deduct, net;
    const double RATE = 0.25;
    const string QUIT = "XXX";
    const string REPORT_HEADING = "Payroll Report ";
    const string END_LINE = "***End of report ";
} // End of main() function
```



If you are typing the code as it is presented here, remember to save the program in a file that is named appropriately, for example, `PayrollReport.cpp`. The complete program is also saved in a file named `PayrollReport.cpp` and is included in the student files for this chapter.



By convention, in C++ the names of constants appear in all uppercase letters. Multiple words are separated with underscores. Some programmers believe using all uppercase makes it easier for you to spot named constants in a long block of code.

Notice that one of the declarations shown in the flowchart, `String COLUMN_HEADING = "Name Gross Deductions Net"` is not included in the C++ code. Since you have not yet learned about the C++ statements needed to line up values in report format, the C++ program shown above prints information on separate lines rather than in the column format used in the flowchart.

It is important for you to understand that the variables and constants declared in the flowchart are global variables that can be used in all modules that are part of the program. However, as mentioned earlier, it is not considered a good programming practice to use global variables. The variables and constants declared in the C++ version are local, which means they can only be used in the `main()` function.

After the declarations, the flowchart makes a call to the `housekeeping()` module that prints the `REPORT_HEADING` and `COLUMN_HEADING` constants and retrieves the first employee's name entered by the user of the program. The code that follows shows how these tasks are translated to C++ statements. The added code is shown in bold.

```
#include <iostream>
#include <string>
using namespace std;
int main()
```

```
{  
    // Declare variables and constants local to main()  
    string name;  
    double gross, deduct, net;  
    const double RATE = 0.25;  
    const string QUIT = "XXX";  
    const string REPORT_HEADING = "Payroll Report";  
    const string END_LINE = "**End of report";  
    // Work done in the housekeeping() function  
    cout << REPORT_HEADING << endl;  
    cout << "Enter employee's name: ";  
    cin >> name;  
} // End of main() function
```

Since it is not considered a good programming practice to use global variables, all of the variables and constants declared for this program are local variables that are available only in the `main()` function. If you were to create an additional function for the housekeeping tasks, that function would not have access to the `name` variable to store an employee's name. So, for now, the C++ programs that you write will have only one function (module), the `main()` function. Additional modules, such as the `housekeeping()` module, will be simulated through the use of comments. As shown in the preceding code, the statements that would execute as part of a `housekeeping()` function have been grouped together in the C++ program and preceded by a comment. You will learn how to create additional functions and pass data to functions in Chapter 9 of this book.

In the flowchart, the next statement to execute after the `housekeeping()` module finishes its work is a `while` loop in the `main` module that continues to execute until the user enters `XXX` when prompted for an employee's name. Within the loop, the `detailLoop()` module is called. The work done in the `detailLoop()` consists of retrieving an employee's gross pay; calculating deductions; calculating net pay; printing the employee's name, gross pay, deductions, and net pay on the user's screen; and retrieving the name of the next employee. The following code shows the C++ statements that have been added to the `PayrollReport` program to implement this logic. The added statements are shown in bold.

```
#include <iostream>  
#include <string>  
using namespace std;  
int main()  
{  
    // Declare variables and constants local to main()  
    string name;  
    double gross, deduct, net;  
    const double RATE = 0.25;  
    const string QUIT = "XXX";  
    const string REPORT_HEADING = "Payroll Report";  
    const string END_LINE = "**End of report ";  
    // Work done in the housekeeping() function  
    cout << REPORT_HEADING << endl;  
    cout << "Enter employee's name: ";  
    cin >> name;  
    while(name != QUIT)
```

```

{
    // work done in the detailLoop() function
    cout << "Enter employee's gross pay: ";
    cin >> gross;
    deduct = gross * RATE;
    net = gross - deduct;
    cout << "Name: " << name << endl;
    cout << "Gross Pay: " << gross << endl;
    cout << "Deductions: " << deduct << endl;
    cout << "Net Pay: " << net << endl;
    cout << "Enter employee's name: ";
    cin >> name;
}
} // End of main() function

```

The `while` loop in the C++ program compares the name entered by the user with the value of the constant named `QUIT`. As long as the name is not equal to `XXX` (the value of `QUIT`), the loop executes. The statements that make up the simulated `detailLoop()` method include retrieving the employee's gross pay; calculating deductions and net pay; printing the employee's name, gross pay, deductions, and net pay; and retrieving the name of the next employee to process.

In the flowchart, when a user enters `XXX` for the employee's name, the program exits the `while` loop and then calls the `endOfJob()` module. The `endOfJob()` module is responsible for printing the value of the `END_LINE` constant. When the `endOfJob()` module finishes, control returns to the mainline module, and the program stops. The completed C++ program is shown next with the additional statements shown in bold.

```

#include <iostream>
#include <string>
using namespace std;
int main()
{
    // Declare variables and constants local to main()
    string name;
    double gross, deduct, net;
    const double RATE = 0.25;
    const string QUIT = "XXX";
    const string REPORT_HEADING = "Payroll Report";
    const string END_LINE = "***End of report ";
    // Work done in the housekeeping() function
    cout << REPORT_HEADING << endl;
    cout << "Enter employee's name: ";
    cin >> name;
    while(name != QUIT)
    {
        // work done in the detailLoop() function
        cout << "Enter employee's gross pay: ";
        cin >> gross;
        deduct = gross * RATE;
        net = gross - deduct;
        cout << "Name: " << name << endl;
        cout << "Gross Pay: " << gross << endl;

```

```
    cout << "Deductions: " << deduct << endl;
    cout << "Net Pay: " << net << endl;
    cout << "Enter employee's name: ";
    cin >> name;
}
// work done in the endOfJob() function
cout << END_LINE;
return 0;
} // End of main() function
```

40

A screenshot of a Windows command prompt window titled "Developer Command Prompt for VS2012". The window shows the output of a C++ program named "PayrollReport". The output is:
C:\>PayrollReport
Payroll Report
Enter employee's name: William
Enter employee's gross pay: 1500
Name: William
Gross Pay: 1500
Deductions: 375
Net Pay: 1125
Enter employee's name: XXX
**End of report
C:\>
The window has a standard Windows title bar and scroll bars.

Figure 3-5 Output of the Payroll Report program when the input is William and 1500

This program is now complete. Figure 3-5 shows the program’s output in response to the input William (for the name), and 1500 (for the gross).

Lab 3-2: Writing a Modular Program in C++

In this lab, you add the input and output statements to a partially completed C++ program. When completed, the user should be able to enter a year, a month, and a day. The program then determines if the date is valid. Valid years are those that are greater than 0, valid months include the values 1 through 12, and valid days include the values 1 through 31.

1. Open the source code file named `BadDate.cpp` using Notepad or the text editor of your choice.
2. Notice that variables have been declared for you.
3. Write the simulated `housekeeping()` function that contains the prompts and input statements to retrieve a year, a month, and a day from the user.
4. Include the output statements in the simulated `endOfJob()` function. The format of the output is as follows:

month/day/year is a valid date.

or

month/day/year is an invalid date.

5. Save this source code file in a directory of your choice, and then make that directory your working directory.
6. Compile the source code file `BadDate.cpp`.
7. Execute the program entering the following date: month = **5**, day = **32**, year = **2014**. Record the output of this program.
8. Execute the program entering the following date: month = **9**, day = **21**, year = **2002**. Record the output of this program.

4

CHAPTER

Writing Programs that Make Decisions

After studying this chapter, you will be able to:

- ④ Use relational and logical Boolean operators to make decisions in a program
- ④ Compare `string` objects
- ④ Write decision statements in C++, including an `if` statement, an `if else` statement, nested `if` statements, and the `switch` statement
- ④ Use decision statements to make multiple comparisons by using AND logic and OR logic

You should complete the exercises and labs in this chapter only after you have finished Chapter 4 of your book, *Programming Logic and Design, Eighth Edition*, by Joyce Farrell. In this chapter, you practice using the relational and logical operators in C++ to write Boolean expressions. You also learn the C++ syntax for decision statements, including the `if` statement, the `if else` statement, nested `if` statements, and `switch` statements. Finally, you learn to write C++ statements that make multiple comparisons.

Boolean Operators

You use Boolean operators in expressions when you want to compare values. When you use a **Boolean operator** in an expression, the evaluation of that expression results in a value that is `true` or `false`. In C++, you can subdivide the Boolean operators into two groups: relational operators and logical operators. We begin the discussion with the relational operators.

Relational Operators

In the context of programming, the term **relational** refers to the connections, or relationships, between values. For example, one value might be greater than another, less than another, or equal to the other value. The terms *greater than*, *less than*, and *equal to* all refer to a relationship between two values. As with all Boolean operators, a relational operator allows you to ask a question that results in a `true` or `false` answer. Depending on the answer, your program will execute different statements that can perform different actions.

Table 4-1 lists the relational operators used in C++.

Operator	Meaning
<code><</code>	Less than
<code><=</code>	Less than or equal to
<code>></code>	Greater than
<code>>=</code>	Greater than or equal to
<code>==</code>	Equal to (two equal signs with no space between them)
<code>!=</code>	Not equal to

Table 4-1 Relational operators

To see how to use relational operators, suppose you declare two variables: an `int` named `number1` that you initialize with the value 10 and another `int` variable named `number2` that

you initialize with the value 15. The following code shows the declaration statements for these variables:

```
int number1 = 10;
int number2 = 15;
```

The following code samples illustrate how relational operators are used in expressions:

- `number1 < number2` evaluates to `true` because 10 is less than 15.
- `number1 <= number2` evaluates to `true` because 10 is less than or equal to 15.
- `number1 > number2` evaluates to `false` because 10 is not greater than 15.
- `number1 >= number2` evaluates to `false` because 10 is not greater than or equal to 15.
- `number1 == number2` evaluates to `false` because 10 is not equal to 15.
- `number1 != number2` evaluates to `true` because 10 is not equal to 15.

Logical Operators

You can use another type of Boolean operator, **logical operators**, when you need to ask more than one question but you want to receive only one answer. For example, in a program, you may want to ask if a number is between the values 1 and 10. This actually involves two questions. You need to ask if the number is greater than 1 *and* if the number is less than 10. Here, you are asking two questions but you want only one answer—either yes (`true`) or no (`false`).

Logical operators are useful in decision statements because, like relational expressions, they evaluate to `true` or `false`, thereby permitting decision making in your programs.

Table 4-2 lists the logical operators used in C++.

Operator	Name	Description
<code>&&</code>	AND	All expressions must evaluate to <code>true</code> for the entire expression to be <code>true</code> ; this operator is written as two & symbols with no space between them.
<code> </code>	OR	Only one expression must evaluate to <code>true</code> for the entire expression to be <code>true</code> ; this operator is written as two symbols with no space between them.
<code>!</code>	NOT	This operator reverses the value of the expression; if the expression evaluates to <code>false</code> , then reverse it so that the expression evaluates to <code>true</code> .

Table 4-2 Logical operators

To see how to use the logical operators, suppose you declare two variables: an `int` named `number1` that you initialize with the value 10, and another `int` variable named `number2` that you initialize with the value 15 as in the previous example.

The following code samples illustrate how you can use the logical operators along with the relational operators in expressions:

- `(number1 > number2) || (number1 == 10)` evaluates to `true` because the first expression evaluates to `false`, 10 is not greater than 15, and the second expression evaluates to `true`, 10 is equal to 10. Only one expression needs to be `true` using OR logic for the entire expression to be `true`.
- `(number1 > number2) && (number1 == 10)` evaluates to `false` because the first expression is `false`, 10 is not greater than 15, and the second expression is `true`, 10 is equal to 10. Using AND logic, both expressions must be `true` for the entire expression to be `true`.
- `(number1 != number2) && (number1 == 10)` evaluates to `true` because both expressions are `true`; that is, 10 is not equal to 15 and 10 is equal to 10. Using AND logic, if both expressions are `true`, then the entire expression is `true`.
- `!(number1 == number2)` evaluates to `true` because the expression evaluates to `false`, 10 is not equal to 15. The `!` operator then reverses `false`, which results in a `true` value.

Relational and Logical Operator Precedence and Associativity

Like the arithmetic operators discussed in Chapter 2, the relational and logical operators are evaluated according to specific rules of associativity and precedence. Table 4-3 shows the precedence and associativity of the operators discussed thus far in this book.

Operator Name	Symbol	Order of Precedence	Associativity
Parentheses	<code>()</code>	First	Left to right
Unary	<code>- + !</code>	Second	Right to left
Multiplication, division, and modulus	<code>* / %</code>	Third	Left to right
Addition and subtraction	<code>+ -</code>	Fourth	Left to right
Relational	<code>< > <= >=</code>	Fifth	Left to right
Equality	<code>== !=</code>	Sixth	Left to right
AND	<code>&&</code>	Seventh	Left to right
OR	<code> </code>	Eighth	Left to right
Assignment	<code>= += -= *= /= %=</code>	Ninth	Right to left

Table 4-3 Order of precedence and associativity

As shown in Table 4-3, the AND operator has a higher precedence than the OR operator, meaning its Boolean values are evaluated first. Also notice that the relational operators have higher precedence than the equality operators and both the relational and equality operators have higher precedence than the AND and OR operators. All of these operators have left-to-right associativity.



Some symbols appear in Table 4-3 more than once because they have more than one meaning. For example, when the (-) operator is used before a number or a variable that contains a number, it is interpreted as the unary (-) operator. When the (-) operator is used between operands, it is interpreted as the subtraction operator.

To see how to use the logical operators and the relational operators in expressions, first assume that the variables `number1` and `number2` are declared and initialized as shown in the following code:

```
int number1 = 10;
int number2 = 15;
```

Now, you write the following expression in C++:

```
number1 == 8 && number2 == number1 || number2 == 15
```

Looking at Table 4-3, you can see that the equality operator (==) has a higher level of precedence than the AND operator (&&) and the AND operator (&&) has a higher level of precedence than the OR operator (||). Also, notice that there are three == operators in the expression; thus, the left-to-right associativity rule applies. Figure 4-1 illustrates the order in which the operators are used.

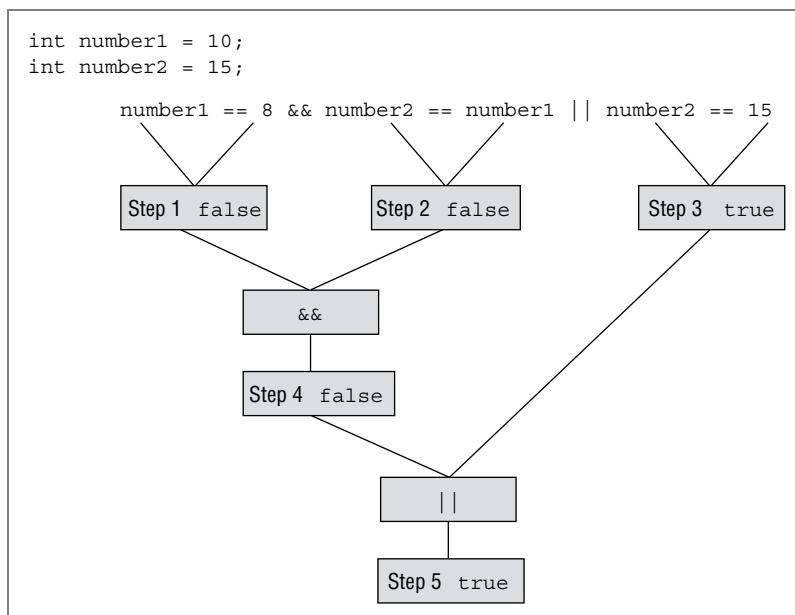


Figure 4-1 Evaluation of expression using relational and logical operators



Remember that you can change the order of precedence by using parentheses.

48

As you can see in Figure 4-1, it takes five steps, following the rules of precedence and associativity, to determine the value of the expression.

As you can see in Figure 4-2, when parentheses are added, it still takes five steps, but the order of evaluation is changed and the result is also changed.

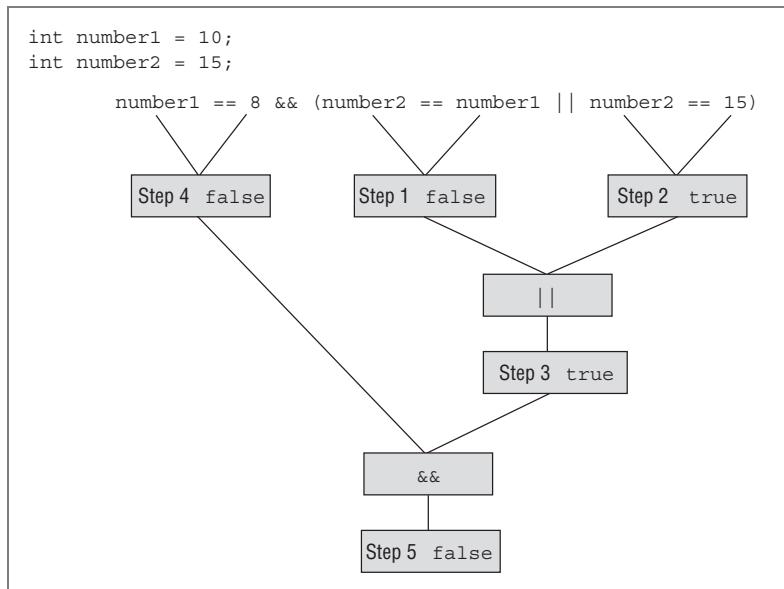


Figure 4-2 Evaluation of expression using relational and logical operators with parentheses

Comparing strings

In C++, you use the same relational operators when you compare `string` objects that you use to compare numeric data types such as `ints` and `doubles`.

The following code shows how to use the equality operator to compare two `string` objects and also to compare one `string` object and one `string` constant:

```
string s1 = "Hello";
string s2 = "World";
// The following test evaluates to false because "Hello" is not
// the same as "World".
if (s1 == s2)
    // code written here executes if true
```

```
else
    // code written here executes if false
// The following test evaluates to true because "Hello" is the
// same as "Hello".
if (s1 == "Hello")
    // code written here executes if true
else
    // code written here executes if false
// The following test evaluates to false because "Hello" is not
// the same as "hello".
if (s1 == "hello")
    // code written here executes if true
else
    // code written here executes if false
```



Two strings are equal when their contents are the same.



C++ is **case sensitive**, which means that C++ does not consider a lowercase *h* to be equal to an uppercase *H* because their ASCII values are different. Lowercase *h* has an ASCII value of 104, and uppercase *H* has an ASCII value of 72. A table of ASCII values can be found in “Appendix A—Understanding Numbering Systems and Computer Codes” in *Programming Logic and Design, Eighth Edition*.

The following code shows how to use the other relational operators to compare two `string` objects and also to compare one `string` object and one `string` constant:

```
string s1 = "Hello";
string s2 = "World";
// The following test evaluates to false because "Hello" is not
// greater than "World".
if (s1 > s2)
    // code written here executes if true
else
    // code written here executes if false
// The following test evaluates to true because "Hello" is the
// same as "Hello".
if (s1 <= "Hello")
    // code written here executes if true
else
    // code written here executes if false
```

When you compare `strings`, C++ compares the ASCII values of the individual characters in the `string` to determine if one `string` is greater than, less than, or equal to another, in terms of alphabetizing the text in the `strings`. As shown in the preceding code, the `string` object `s1`, whose value is "Hello", is not greater than the `string` object `s2`, whose value is "World", because *World* comes after *Hello* in alphabetical order.

The following code sample shows additional examples of using the relational operators with two `string` objects:

```
string s1 = "whole";
string s2 = "whale";
// The next statement evaluates to true because the contents of
// s1, "whole", are greater than the contents of s2, "whale".
if (s1 > s2)
    // code written here executes if true
else
    // code written here executes if false
// The next statement evaluates to true because the contents of
// s2, "whale", are less than the contents of s1, "whole".
if (s2 < s1)
    // code written here executes if true
else
    // code written here executes if false
// The next statement evaluates to true because the contents of
// s1, "whole", are the same as the string constant, "whole".
if (s1 == "whole")
    // code written here executes if true
else
    // code written here executes if false
```

Decision Statements

Every decision in a program is based on whether an expression evaluates to `true` or `false`. Programmers use decision statements to change the flow of control in a program. **Flow of control** means the order in which statements are executed. Decision statements are also known as branching statements, because they cause the computer to make a decision, choosing from one or more branches (or paths) in the program.

There are different types of decision statements in C++. We will begin with the `if` statement.

The `if` Statement

The `if` statement is a single-path decision statement. As you learned in *Programming Logic and Design*, `if` statements are also referred to as “single alternative” or “single-sided” statements.

When we use the term **single-path**, we mean that if an expression evaluates to `true`, your program executes one or more statements, but if the expression evaluates to `false`, your program will not execute these statements. There is only one defined path—the path taken if the expression evaluates to `true`. In either case, the statement following the `if` statement is executed.

The **syntax**, or set of rules, for writing an `if` statement in C++ is as follows:

```
if(expression)
    statementA;
```

Note that when you type the keyword `if` to begin an `if` statement, you follow it with an expression placed within parentheses.

When the compiler encounters an **if** statement, the expression within the parentheses is evaluated. If the expression evaluates to **true**, then the computer executes **statementA**. If the expression in parentheses evaluates to **false**, then the computer will not execute **statementA**. Remember that whether the expression evaluates to **true** and executes **statementA**, or the expression evaluates to **false** and does not execute **statementA**, the statement following the **if** statement executes next.

Note that a C++ statement, such as an **if** statement, can be either a simple statement or a block statement. A block statement is made up of multiple C++ statements. C++ defines a block as statements placed within a pair of curly braces. If you want your program to execute more than one statement as part of an **if** statement, you must enclose the statements in a pair of curly braces or only one statement will execute. The following example illustrates an **if** statement that uses the relational operator (**<**) to test if the value of the variable **customerAge** is less than 65. You will see the first curly brace in the fourth line and the second curly brace in the third to last line.

```
int customerAge = 53;
int discount = 10, numUnder_65 = 0;
if(customerAge < 65)
{
    discount = 0;
    numUnder_65 += 1;
}
cout << "Discount : " << discount << endl;
cout << "Number of customers under 65 is: " << numUnder_65 << endl;
```

In the preceding code, the variable named **customerAge** is initialized to the value 53. Because 53 is less than 65, the expression **customerAge < 65** evaluates to **true** and the block statement executes. The block statement is made up of the two assignment statements within the curly braces: **discount = 0;** and **numUnder_65 += 1;**. If the expression evaluates to **false**, the block statement does not execute. In either case, the next statement to execute is the output statement **cout << "Discount : " << discount << endl;**.

Notice that you do not include a semicolon at the end of the line with the **if** and the expression to be tested. Doing so is not a syntax error, but it can create a logic error in your program. A **logic error** causes your program to produce incorrect results. In C++, the semicolon (**;**) is called the null statement and is considered a legal statement. The **null** statement is a statement that does nothing. Examine the following code:

```
if (customerAge < 65); // semicolon here is not correct
{
    discount = 0;
    numUnder_65 += 1;
}
```

If you write an **if** statement as shown in the preceding code, your program will test the expression **customerAge < 65**. If it evaluates to **true**, the null statement executes, which means your program does nothing, and then the statement **discount = 0;** executes because this is the next statement following the **if** statement. This does not cause a logic error in your program, but consider what happens when the expression in the **if** statement evaluates to

false. If false, the null statement does not execute, but the statement `discount = 0;` will execute because it is the next statement after the if statement.

The following code uses an if statement to test two string objects for equality:

```
string dentPlan = "Y";
double grossPay = 500.00;
if (dentPlan == "Y")
    grossPay = grossPay - 23.50;
```

In this example, if the value of the string object named `dentPlan` and the string constant "Y" are the same value, the expression evaluates to true, and the `grossPay` calculation assignment statement executes. If the expression evaluates to false, the `grossPay` calculation assignment statement does not execute.

Exercise 4-1: Understanding if Statements

In this exercise, you use what you have learned about writing if statements in C++ to study a complete C++ program that uses if statements. Take a few minutes to study the code that follows, and then answer Questions 1–4.

```
// VotingAge.cpp - This program determines if a
// person is eligible to vote.
#include <iostream>
#include <string>
using namespace std;
int main()
{
    int myAge = 17;
    string ableToVote = "Yes";
    const int VOTING_AGE = 18;
    if(myAge < VOTING_AGE)
        ableToVote = "No";
    cout << "My Age: " << myAge << endl;
    cout << "Able To Vote: " << ableToVote << endl;
    return 0;
}
```

1. What is the exact output when this program executes?

2. What is the exact output if the value of `myAge` is changed to 19?

3. What is the exact output if the value of `myAge` is changed to 18 and the expression in the `if` statement is changed to `myAge <= VOTING_AGE`?
-
-

4. What is the exact output if the value of `myAge` is changed to 18 and the variable named `ableToVote` is initialized with the value "No" rather than the value "Yes"?
-
-

Lab 4-1: Understanding `if` Statements

In this lab, you complete a prewritten C++ program for a carpenter who creates personalized house signs. The program is supposed to compute the price of any sign a customer orders, based on the following facts:

- The charge for all signs is a minimum of \$35.00.
- The first five letters or numbers are included in the minimum charge; there is a \$4 charge for each additional character.
- If the sign is made of oak, add \$20.00. No charge is added for pine.
- Black or white characters are included in the minimum charge; there is an additional \$15 charge for gold-leaf lettering.

1. Open the file named `HouseSign.cpp` using Notepad or the text editor of your choice.
2. You need to declare variables for the following, and initialize them where specified:
 - A variable for the cost of the sign initialized to 0.00.
 - A variable for the color of the characters initialized to "gold".
 - A variable for the wood type initialized with the value "oak".
 - A variable for the number of characters initialized with the value 8.
3. Write the rest of the program using assignment statements and `if` statements as appropriate. The output statements are written for you.
4. Compile the program.
5. Execute the program. Your output should be: The charge for this sign is \$82.

Note that you cannot control the number of places that appear after the decimal point until you learn more about C++ in Chapter 9 of this book.

The `if else` Statement

The `if else` statement is a dual-path or dual-alternative decision statement. That is, your program will take one of two paths as a result of evaluating an expression in an `if else` statement.



Do not include a semicolon at the end of the line containing the keyword `if` and the expression to be tested, or on the line with the keyword `else`. As you learned earlier, doing so is not a syntax error, but it can create a logic error in your program.

The syntax for writing an `if else` statement in C++ is as follows:

```
if (expression)
    statementA;
else
    statementB;
```

When the compiler encounters an `if else` statement, the expression in the parentheses is evaluated. If the expression evaluates to `true`, then the computer executes `statementA`. If the expression in parentheses evaluates to `false`, then the computer executes `statementB`. Both `statementA` and `statementB` can be simple statements or block statements. Regardless of which path is taken in a program, the statement following the `if else` statement is the next one to execute.

The following code sample illustrates an `if else` statement written in C++:

```
int hoursWorked = 45;
double rate = 15.00;
double grossPay;
string overtime = "Yes";
const int HOURS_IN_WEEK = 40;
const double OVERTIME_RATE = 1.5;
if (hoursWorked > HOURS_IN_WEEK)
{
    overtime = "Yes";
    grossPay = HOURS_IN_WEEK * rate +
        (hoursWorked - HOURS_IN_WEEK) * OVERTIME_RATE * rate;
}
else
{
    overtime = "No";
    grossPay = hoursWorked * rate;
}
cout << "Overtime: " << overtime << endl;
cout << "Gross Pay: $" << grossPay << endl;
```



`HOURS_IN_WEEK` is a constant that is initialized with the value 40, and `OVERTIME_RATE` is a constant that is initialized with the value 1.5.

In the preceding code, the value of the variable named `hoursWorked` is tested to see if it is greater than `HOURS_IN_WEEK`.

You use the greater-than relational operator (`>`) to make the comparison. If the expression `hoursWorked > HOURS_IN_WEEK` evaluates to `true`, then the first block statement executes. This first block statement contains one statement that assigns the string constant "Yes" to the variable named `overtime`, and another statement that calculates the employee's gross pay, including overtime pay, and assigns the calculated value to the variable named `grossPay`.

If the expression `hoursWorked > HOURS_IN_WEEK` evaluates to `false` then a different path is followed, and the second block statement following the keyword `else` executes. This block statement contains one statement that assigns the string constant "No" to the variable named `overtime`, and another statement that calculates the employee's gross pay with no overtime and assigns the calculated value to the variable named `grossPay`.

Regardless of which path is taken in this code, the next statement to execute is the output statement `cout << "Overtime: " << overtime << endl;` immediately followed by the output statement `cout << "Gross Pay: $" << grossPay << endl;`.

Exercise 4-2: Understanding `if else` Statements

In this exercise, you use what you have learned about writing `if else` statements in C++ to study a complete C++ program that uses `if else` statements. This program was written to calculate customer charges for a telephone company. The telephone company charges 25 cents per minute for calls outside of the customer's area code that last over 10 minutes. All other calls are 10 cents per minute. Take a few minutes to study the code that follows, and then answer Questions 1–4.

```
// Telephone.cpp - This program determines telephone call
// charges.
#include <iostream>
using namespace std;
int main()
{
    int custAC, custNumber;
    int calledAC, calledNumber;
    int callMinutes;
    double callCharge;
    const int MAX_MINS = 10;
    const double CHARGE_1 = .25;
    const double CHARGE_2 = .10;
    custAC = 847;
    custNumber = 5551234;
    calledAC = 630;
    calledNumber = 5557890;
    callMinutes = 50;
    if(calledAC != custAC && callMinutes > MAX_MINS)
        callCharge = callMinutes * CHARGE_1;
    else
        callCharge = callMinutes * CHARGE_2;
    cout << "Customer Number: " << custAC << "-" << custNumber
        << endl;
```

```
cout << "Called Number: " << calledAC << "-"
    << calledNumber << endl;
cout << "The charge for this call is $" << callCharge
    << endl;
return 0;
}
```

56

1. What is the exact output when this program executes?

2. What is the exact output if the value of `callMinutes` is changed to 20?

3. What is the exact output if the expression in the `if` statement is changed to `callMinutes >= MAX_MINS`?

4. What is the exact output if the variable named `calledAC` is assigned the value 847 rather than the value 630?

Lab 4-2: Understanding `if` `else` Statements

In this lab, you complete a prewritten C++ program that computes the largest and smallest of three integer values. The three values are -50, 53, 78.

1. Open the file named `LargeSmall.cpp` using Notepad or the text editor of your choice.
2. Two variables named `largest` and `smallest` are declared for you. Use these variables to store the largest and smallest of the three integer values. You must decide what other variables you will need and initialize them if appropriate.
3. Write the rest of the program using assignment statements, `if` statements, or `if` `else` statements as appropriate. There are comments in the code that tell you where you should write your statements. The output statements are written for you.
4. Compile the program.
5. Execute the program. Your output should be:

The largest value is 78
The smallest value is -50

Nested if Statements

You can nest `if` statements to create a multipath decision statement. When you nest `if` statements, you include an `if` statement within another `if` statement. This is helpful in programs in which you want to provide more than two possible paths.



Do not include a semicolon at the end of the lines with expressions to be tested or on the line with the keyword `else`.

The syntax for writing a nested `if` statement in C++ is as follows:

```
if(expressionA)
    statementA;
else if(expressionB)
    statementB;
else
    statementC;
```

This is called a nested `if` statement because the second `if` statement is a part of the first `if` statement. This is easier to see if the example is changed as follows:

```
if(expressionA)
    statementA;
else
    if(expressionB)
        statementB;
    else
        statementC;
```

Now you will see how a nested `if` statement works. If `expressionA`, which is enclosed in parentheses, evaluates to `true`, then the computer executes `statementA`. If `expressionA` evaluates to `false`, then the computer will evaluate `expressionB`. If `expressionB` evaluates to `true`, then the computer will execute `statementB`. If `expressionA` and `expressionB` both evaluate to `false`, then the computer will execute `statementC`. Regardless of which path is taken in this code, the statement following the `if` `else` statement is the next one to execute.

The C++ code sample that follows illustrates a nested `if` statement.

```
if(empDept <= 3)
    supervisorName = "Dillon";
else if(empDept <= 7)
    supervisorName = "Escher";
else
    supervisorName = "Fontana";
cout << "Supervisor: " << supervisorName << endl;
```

When you read the preceding code, you can assume that a department number is never less than 1. If the value of the variable named `empDept` is less than or equal to the value 3 (in the range of values from 1 to 3), then the value "Dillon" is assigned to the variable named `supervisorName`. If the value of `empDept` is not less than or equal to 3, but it is less than or equal to 7 (in the range of values from 4 to 7), then the value "Escher" is assigned to the

variable named `supervisorName`. If the value of `empDept` is not in the range of values from 1 to 7, then the value "Fontana" is assigned to the variable named `supervisorName`. As you can see, there are three possible paths this program could take when the nested `if` statement is encountered. Regardless of which path the program takes, the next statement to execute is the output statement

```
cout << "Supervisor: " << supervisorName << endl;
```

58

Exercise 4-3: Understanding Nested `if` Statements

In this exercise, you use what you have learned about writing nested `if` statements. This program was written for the Woof Wash dog-grooming business to calculate a total charge for services rendered. Woof Wash charges \$12 for a bath, \$10 for a trim cut, and \$7 to clip nails. Take a few minutes to study the code that follows, and then answer Questions 1–3.

```
// WoofWash.cpp - This program determines if a doggy
// service is provided and prints the charge.
#include <iostream>
#include <string>
using namespace std;
int main()
{
    string service;
    const string SERVICE_1 = "bath";
    const string SERVICE_2 = "cut";
    const string SERVICE_3 = "trim nails";
    double charge;
    const double BATH_CHARGE = 12.00;
    const double CUT_CHARGE = 10.00;
    const double NAIL_CHARGE = 7.00;
    cout << "Enter service: ";
    cin >> service;
    if(service == SERVICE_1)
        charge = BATH_CHARGE;
    else if(service == SERVICE_2)
        charge = CUT_CHARGE;
    else if(service == SERVICE_3)
        charge = NAIL_CHARGE;
    else
        charge = 0.00;
    if(charge > 0.00)
        cout << "The charge for a doggy " << service << " is $" 
            << charge << endl;
    else
        cout << "We do not perform the " << service
            << " service." << endl;
    return 0;
}
```

1. What is the exact output when this program executes if the user enters "bath"?
-
-

2. What is the exact output when this program executes if the user enters "shave"?
-
-

3. What is the exact output when this program executes if the user enters "BATH"?
-
-

Lab 4-3: Understanding Nested `if` Statements

In this lab, you complete a prewritten C++ program that calculates an employee's productivity bonus and prints the employee's name and bonus. Bonuses are calculated based on an employee's productivity score as shown in Table 4-4. A productivity score is calculated by first dividing an employee's transactions dollar value by the number of transactions and then dividing the result by the number of shifts worked.

Productivity Score	Bonus
<=30	\$50
31–69	\$75
70–199	\$100
>= 200	\$200

Table 4-4 Employee productivity scores and bonuses

1. Open the file named `EmployeeBonus.cpp` using Notepad or the text editor of your choice.
2. Variables have been declared for you, and the input statements and output statements have been written. Read them over carefully before you proceed to the next step.
3. Design the logic, and write the rest of the program using a nested `if` statement.
4. Compile the program.
5. Execute the program and enter the following as input:

Employee's first name: Kim Smith

Number of shifts: 25

Number of transactions: 75

Transaction dollar value: 40000.00

6. Your output should be:
Employee Name: Kim Smith
Employee Bonus: \$50.0

You cannot control the number of places that appear after the decimal point until you learn more about C++ in Chapter 9 of this book.

60

The switch Statement

The `switch` statement is similar to a nested `if` statement because it is also a multipath decision statement. A `switch` statement offers the advantage of being easier for you to read than nested `if` statements, and a `switch` statement is also easier for you, the programmer, to maintain. You use the `switch` statement in situations when you want to compare an expression with several integer constants.

The syntax for writing a `switch` statement in C++ is as follows:

```
switch(expression)
{
    case constant: statement(s);
    case constant: statement(s);
    case constant: statement(s);
    default:         statement(s);
}
```

You begin writing a `switch` statement with the keyword `switch`. Then, within parentheses, you include an expression that evaluates to an integer value. Cases are then defined within the `switch` statement by using the keyword `case` as a label, and including an integer value after this label. For example, you could include an integer constant such as 10 or an arithmetic expression that evaluates to an integer such as $10/2$. The computer evaluates the expression in the `switch` statement and then compares it to the integer values following the `case` labels. If the expression and the integer value match, then the computer executes the `statement(s)` that follow until it encounters a `break` statement or a closing curly brace. The `break` statement causes an exit from the `switch` statement. You can use the keyword `default` to establish a case for values that do not match any of the integer values following the `case` labels. Note also that all of the cases, including the `default` case, are enclosed within curly braces.



If you omit a `break` statement in a `case`, all the code up to the next `break` statement or a closing curly brace is executed. This is probably not what you intend.

The following code sample illustrates the use of the `switch` statement in C++:

```
int deptNum;
string deptName;
deptNum = 2;
switch(deptNum)
{
    case 1: deptName = "Marketing";
              break;
    case 2: deptName = "Development";
              break;
    case 3: deptName = "Sales";
              break;
    default:deptName = "Unknown";
              break;
}
cout << "Department: " << deptName << endl;
```

In the preceding example, when the program encounters the `switch` statement, the value of the variable named `deptNum` is 2. The value 2 matches the integer constant 2 in the second case of the `switch` statement. Therefore, the string constant "Development" is assigned to the `string` object named `deptName`. A `break` statement is encountered next, which causes the program to exit from the `switch` statement. The statement following the `switch` statement

`cout << "Department: " << deptName << endl;` executes next.

If the `break` statements in the preceding example were omitted and the value of `deptNum` was 2, all of the statements up to the closing curly brace would execute. The output would be: Department: Unknown. In this example, omitting the `break` statements would be considered a logic error.

Exercise 4-4: Using a switch Statement

In this exercise, you use what you have learned about the `switch` statement to study some C++ code, and then answer Questions 1–5.

First, examine the following code:

```
int numValue = 10;
int answer = 0;
switch(numValue)
{
    case 5: answer += 5;
    case 10: answer += 10;
    case 15: answer += 15;
              break;
    case 20: answer += 20;
    case 25: answer += 25;
    default: answer = 0;
              break;
}
cout << "Answer: " << answer << endl;
```

1. What is the value of `answer` if the value of `numValue` is 10?

2. What is the value of `answer` if the value of `numValue` is 20?

3. What is the value of `answer` if the value of `numValue` is 5?

4. What is the value of `answer` if the value of `numValue` is 17?

5. Is the `break` statement in the `default` case needed? Explain.

Lab 4-4: Using a `switch` Statement

In this lab, you complete a prewritten C++ program that calculates an employee’s end-of-year bonus and prints the employee’s name, yearly salary, performance rating, and bonus. In this program, bonuses are calculated based on employees’ annual salary and their performance rating. The rating system is contained in Table 4-5.

Rating	Bonus
1	25 percent of annual salary
2	15 percent of annual salary
3	10 percent of annual salary
4	None

Table 4-5 Employee ratings and bonuses

1. Open the file named `EmployeeBonus2.cpp` using Notepad or the text editor of your choice.
2. Variables have been declared for you, and the input statements and output statements have been written. Read them over carefully before you proceed to the next step.
3. Design the logic, and write the rest of the program using a `switch` statement.
4. Compile the program.
5. Execute the program entering the following as input:

Employee’s name: Jeanne Hanson

Employee’s salary: 70000.00

Employee’s performance rating: 2

6. Confirm that your output matches the following:

Employee Name: Jeanne Hanson
Employee Salary: \$70000
Employee Rating: 2
Employee Bonus: \$10500

Using Decision Statements to Make Multiple Comparisons

When you write programs, you must often write statements that include multiple comparisons. For example, you may want to determine that two conditions are **true** before you decide which path your program will take. In the next sections, you learn how to implement AND logic in a program by using the AND (**&&**) logical operator. You also learn how to implement OR logic using the OR (**||**) logical operator.

Using AND Logic

When you write C++ programs, you can use the AND operator (**&&**) to make multiple comparisons in a single decision statement. Remember when using AND logic that all expressions must evaluate to **true** for the entire expression to be **true**.

The C++ code that follows illustrates a decision statement that uses the AND operator (**&&**) to implement AND logic:

```
string medicalPlan = "Y";
string dentalPlan = "Y";
if((medicalPlan == "Y") && (dentalPlan == "Y"))
    cout << "Employee has medical insurance" <<
          " and also has dental insurance." << endl;
else
    cout << "Employee may have medical insurance or may " <<
          "have dental insurance, but does not have both " <<
          "medical and dental insurance." << endl;
```

In this example, the variables named `medicalPlan` and `dentalPlan` have both been initialized to the string constant "Y". When the expression `medicalPlan == "Y"` is evaluated, the result is **true**. When the expression `dentalPlan == "Y"` is evaluated, the result is also **true**. Because both expressions evaluate to **true**, the entire expression `medicalPlan == "Y" && dentalPlan == "Y"` evaluates to **true**. Because the entire expression is **true**, the output generated is "Employee has medical insurance and also has dental insurance."

If you initialize either of the variables `medicalPlan` or `dentalPlan` with a value other than "Y", then the expression `medicalPlan == "Y" && dentalPlan == "Y"` evaluates to **false**, and the output generated is "Employee may have medical insurance or may have dental insurance, but does not have both medical and dental insurance."

Using OR Logic

You can use OR logic when you want to make multiple comparisons in a single decision statement. Of course, you must remember when using OR logic that only one expression must evaluate to `true` for the entire expression to be `true`.

64

The C++ code that follows illustrates a decision statement that uses the OR operator (`||`) to implement OR logic:

```
string medicalPlan = "Y";
string dentalPlan = "N";
if(medicalPlan == "Y" || dentalPlan == "Y")
    cout << "Employee has medical insurance or dental " <<
        "insurance or both." << endl;
else
    cout << "Employee does not have medical insurance " <<
        "and also does not have dental insurance." << endl;
```

In this example, the variable named `medicalPlan` is initialized with the string constant "Y", and the variable named `dentalPlan` is initialized to the string constant "N". When the expression `medicalPlan == "Y"` is evaluated, the result is `true`. When the expression `dentalPlan == "Y"` is evaluated, the result is `false`. The expression `medicalPlan == "Y" || dentalPlan == "Y"` evaluates to `true` because when using OR logic, only one of the expressions must evaluate to `true` for the entire expression to be `true`. Because the entire expression is `true`, the output generated is "Employee has medical insurance or dental insurance or both."

If you initialize both of the variables `medicalPlan` and `dentalPlan` with the string constant "N", then the expression `medicalPlan == "Y" || dentalPlan == "Y"` evaluates to `false`, and the output generated is "Employee does not have medical insurance and also does not have dental insurance."

Exercise 4-5: Making Multiple Comparisons in Decision Statements

In this exercise, you use what you have learned about OR logic to study a complete C++ program that uses OR logic in a decision statement. This program was written for a marketing research firm that wants to determine if a customer prefers Coke or Pepsi over some other drink. Take a few minutes to study the code that follows, and then answer Questions 1–4.

```
// CokeOrPepsi.cpp - This program determines if a customer
// prefers to drink Coke or Pepsi or some other drink.
#include <iostream>
#include <string>
using namespace std;
int main()
{
    string customerFirstName; // Customer's first name
    string customerLastName; // Customer's last name
    string drink = ""; // Customer's favorite drink
    cout << "Enter customer's first name: ";
```

```
cin >> customerFirstName;
cout << "Enter customer's last name: ";
cin >> customerLastName;
cout << "Enter customer's drink preference: ";
cin >> drink;
if(drink == "Coke" || drink == "Pepsi")
{
    cout << "Customer First Name: " << customerFirstName
        << endl;
    cout << "Customer Last Name: " << customerLastName
        << endl;
    cout << "Drink: " << drink << endl;
}
else
    cout << customerFirstName << " " << customerLastName
        << " does not prefer Coke or Pepsi." << endl;
return 0;
}
```

1. What is the exact output when this program executes if the customer's name is Chas Matson and the drink is Coke?

2. What is the exact output when this program executes if the customer's name is Chas Matson and the drink is Pepsi?

3. What is the exact output from this program when

```
if(drink == "Coke" || drink == "Pepsi")
```

is changed to

```
if(drink == "Coke" && drink == "Pepsi")
```

and the customer's name is Chas Matson and the drink is Pepsi?

4. What is the exact output from this program when

```
if(drink == "Coke" || drink == "Pepsi")
```

is changed to

```
if(drink == "Coke" || drink == "Pepsi" || drink == "coke"  
|| drink == "pepsi")
```

and the customer's name is Chas Matson, and the drink is coke? What does this change allow a user to enter?

Lab 4-5: Making Multiple Comparisons in Decision Statements

In this lab, you complete a partially written C++ program for an airline that offers a 25 percent discount to passengers who are 6 years old or younger and the same discount to passengers who are 65 years old or older. The program should request a passenger's name and age and then print whether the passenger is eligible or not eligible for a discount.

1. Open the file named `Airline.cpp` using Notepad or the text editor of your choice.
2. Variables have been declared and initialized for you, and the input statements have been written. Read them carefully before you proceed to the next step.
3. Design the logic deciding whether to use AND or OR logic. Write the decision statement to identify when a discount should be offered and when a discount should not be offered.
4. Be sure to include output statements telling whether or not the customer is eligible for a discount.
5. Compile the program.
6. Execute the program, entering the following as input:
 - a. Customer Name: Will Moriarty
Customer Age: 11
What is the output? _____
 - b. Customer Name: James Chung
Customer Age: 64
What is the output? _____
 - c. Customer Name: Darlene Sanchez
Customer Age: 75
What is the output? _____

- d. Customer Name: Ray Sanchez

Customer Age: 60

What is the output? _____

- e. Customer Name: Tommy Sanchez

Customer Age: 6

What is the output? _____

67

- f. Customer Name: Amy Patel

Customer Age: 8

What is the output? _____

5

CHAPTER

Writing Programs Using Loops

After studying this chapter, you will be able to:

- ◎ Use the increment (`++`) and decrement (`--`) operators in C++
- ◎ Recognize how and when to use `while` loops in C++, including how to use a counter and how to use a sentinel value to control a loop
- ◎ Use `for` loops in C++
- ◎ Write a `do while` loop in C++
- ◎ Include nested loops in applications
- ◎ Accumulate totals by using a loop in a C++ application
- ◎ Use a loop to validate user input in an application

In this chapter, you learn how to use C++ to program three types of loops: a `while` loop, a `do while` loop, and a `for` loop. You also learn how to nest loops, how to use a loop to help you accumulate a total in your programs, and how to use a loop to validate user input. You should do the exercises and labs in this chapter after you have finished Chapter 5 in your book, *Programming Logic and Design, Eighth Edition*, by Joyce Farrell. In that chapter, you learned that loops change the flow of control in a program by allowing a programmer to direct the computer to execute a statement or a group of statements multiple times. But before you start learning about the loops in C++, it is helpful to learn about two additional operators, the increment and decrement operators.

The Increment (`++`) and Decrement (`--`) Operators

You will often use the increment and decrement operators when your programs require loops. These operators provide a concise, efficient method for adding 1 to (incrementing) or subtracting 1 from (decrementing) an lvalue. An **lvalue** is an area of memory in which a value that your program needs may be stored. In C++ code, you place an lvalue on the left side of an assignment statement. Recall that an assignment statement stores a value at a memory location that is associated with a variable, and you place a variable name on the left side of an assignment statement.



The *l* in *lvalue* stands for *left*.



The increment and decrement operators may be used only with integer data types.

For example, the C++ assignment statement

```
number = 10;
```

assigns the value 10 to the variable named `number`. This causes the computer to store the value 10 at the memory location associated with `number`. Because the increment and decrement operators add 1 to or subtract 1 from an lvalue, the statement `number++;` is equivalent to `number = number + 1;`, and the statement `number--;` is equivalent to `number = number - 1;.` Each expression changes or writes to the memory location associated with the variable named `number`.

Both the increment and decrement operators have prefix and postfix forms. Which form you use depends on when you want to increment or decrement the value stored in the variable. When you use the **prefix form**, as in `++number`, you place the operator in front of the name of the variable. This increments or decrements the lvalue immediately. When you use the **postfix form**, as in `number++`, you place the operator after the name of the variable. This increments or decrements the lvalue after it is used.

The example that follows illustrates the use of both forms of the increment operator in C++:

```
x = 5;  
y = x++; // Postfix form  
// y is assigned the value of x,  
// then x is incremented.  
// Value of y is 5.  
// Value of x is 6.  
  
x = 5;  
y = ++x; // Prefix form  
// x is incremented first, then  
// the value of x is assigned to y.  
// Value of y is 6.  
// Value of x is 6.
```

You might understand the postfix form better if you think of the statement `y = x++;` as being the same as the following:

```
x = 5;  
y = x;  
x = x + 1;
```

To understand the prefix form better, think of `y = ++x;` as being the same as the following:

```
x = 5;  
x = x + 1;  
y = x;
```

Exercise 5-1: Using the Increment and Decrement Operators

In this exercise, you use what you have learned about increment and decrement operators in C++ to answer Questions 1–4.

1. Examine the following code:

```
x = 6;  
y = ++x;
```

After this code executes, what is the value of x? _____ y? _____

2. Examine the following code:

```
x = 6;  
y = x++;
```

After this code executes, what is the value of x? _____ y? _____

3. Examine the following code:

```
x = 6;  
y = --x;
```

After this code executes, what is the value of x? _____ y? _____

4. Examine the following code:

```
x = 6;  
y = x--;
```

After this code executes, what is the value of x? _____ y? _____

72

Writing a **while** Loop in C++

As you learned in *Programming Logic and Design*, three steps must occur in every loop:

1. You must initialize a variable that will control the loop. This variable is known as the **loop control variable**.
2. You must compare the loop control variable to some value, known as the **sentinel value**, which decides whether the loop continues or stops. This decision is based on a Boolean comparison. The result of a **Boolean** comparison is always a **true** or **false** value.
3. Within the loop, you must alter the value of the loop control variable.

You also learned that the statements that are part of a loop are referred to as the **loop body**. In C++, the loop body may consist of a single statement or a block statement.



Remember that a block statement is several statements within a pair of curly braces.

The statements that make up a loop body may be any type of statement, including assignment statements, decision statements, or even other loops. The C++ syntax for writing a **while** loop is as follows:

```
while(expression)  
    statement;
```

Notice there is no semicolon after the ending parenthesis. Placing a semicolon after the ending parenthesis is not a syntax error, but it is a logic error. It results in an **infinite loop**, which is a loop that never stops executing the statements in its body. It never stops executing because the semicolon is a statement called the **null** statement and is interpreted as “do nothing.” Think of a **while** loop with a semicolon after the ending parenthesis as meaning “while the condition is true, do nothing forever.”

The **while** loop allows you to direct the computer to execute the statement in the body of the loop as long as the expression within the parentheses evaluates to **true**. Study the example that follows, which illustrates a **while** loop that uses a block statement as its loop body:

```
const int NUM_TIMES = 3;  
int num = 0;  
while(num < NUM_TIMES)  
{  
    cout << "Welcome to C++ Programming." << endl;  
    num++;  
}
```

In this example, a block statement is used because the loop body contains more than one statement.

The first statement causes the text "Welcome to C++ Programming." to appear on the user's screen. The second statement, `num++`, is important because it causes `num`, the loop control variable, to be incremented. When the loop is first encountered the comparison `num < NUM_TIMES` is made for the first time when the value of `num` is 0. The 0 is compared to and found to be less than 3, which means the condition is `true`, and the text "Welcome to C++ Programming." is displayed for the first time. The next statement, `num++;` causes 1 to be added to the value of `num`. The second time the comparison is made, the value of `num` is 1, which is still less than 3, and causes the text to appear a second time followed by adding 1 to the value of `num`. The third comparison also results in a `true` value because the value of `num` is now 2, and 2 is still less than 3; as a result, the text appears a third time and `num` is incremented again. The fourth time the comparison is made, the value of `num` is 3, which is not less than 3; as a result, the program exits the loop.

The loop in the next code example produces the same results as the previous example. The text "Welcome to C++ Programming." is displayed three times.

```
const int NUM_TIMES = 3;
int num = 0;
while(num++ < NUM_TIMES)
    cout << "Welcome to C++ Programming." << endl;
```

Be sure you understand why the postfix increment operator is used in the expression `num++ < NUM_TIMES`.

The first time this comparison is made, the value of `num` is 0. The 0 is then compared to and found to be less than 3, which means the condition is `true`, and the text "Welcome to C++ Programming." is displayed.



When you use the postfix increment operator, the value of `num` is not incremented until after the comparison is made.

The second time the comparison is made, the value of `num` is 1; because 1 is still less than 3, the text appears a second time. The third comparison also results in a `true` value because the value of `num` is now 2, and 2 is still less than 3; as a result, the text appears a third time. The fourth time the comparison is made, the value of `num` is 3, which is not less than 3; as a result, the program exits the loop.



Failing to learn the difference between the prefix and postfix forms of the increment and decrement operators can result in serious program errors.

If the prefix increment operator is used in the expression `++num < NUM_TIMES`, the loop executes twice instead of three times. This occurs because the first time this comparison is made, `num` is incremented before the comparison is done. This results in

`num` having a value of 1 the first time "Welcome to C++ Programming." is displayed and a value of 2 the second time it is displayed. Then, when the value of `num` is 3, the condition is `false`, causing the program to exit the loop. This time, "Welcome to C++ Programming." is not displayed.

74

Exercise 5-2: Using a `while` Loop

In this exercise, you use what you have learned about writing `while` loops. Study the following code, and then answer Questions 1–4.

```
const int NUM_LOOPS = 8;
int numTimes = NUM_LOOPS;
while(numTimes++ < NUM_LOOPS)
    cout << "Value of numTimes is " << numTimes << endl;
```

1. What is the loop control variable? What is the sentinel value?

2. What is the output?

3. What is the output if the code is changed to `while(numTimes++ <= NUM_LOOPS)`?

4. What is the output if the code is changed to `while(++numTimes <= NUM_LOOPS)`?

Using a Counter to Control a Loop

In Chapter 5 of *Programming Logic and Design*, you learned that you can use a counter to control a `while` loop. With a counter, you set up the loop to execute a specified number of times. Also recall that a `while` loop will execute zero times if the expression used in the comparison immediately evaluates to `false`. In that case, the computer does not execute the body of the loop at all.

Chapter 5 of *Programming Logic and Design* discusses a counter-controlled loop that controls how many times the word "Hello" is printed. Take a look at the following pseudocode for this counter-controlled loop:

```
start
    Declarations
        num count = 0
    while count < 4
        print "Hello"
        count = count + 1
    endwhile
    output "Goodbye"
stop
```



Incrementing the counter variable is an important statement. Each time through the loop, the **count** variable must be incremented or the expression **count < 4** would never be **false**. This would result in an infinite loop.

The counter for this loop is a variable named **count**, which is assigned the value 0. The Boolean expression, **count < 4**, is tested to see if the value of **count** is less than 4. If **true**, the loop executes. If **false**, the program exits the loop. If the loop executes, the program displays the word "Hello", and then adds 1 to the value of **count**. Given this pseudocode, the loop body executes four times, and the word "Hello" is displayed four times.

This is what the code looks like when you translate the pseudocode to C++:

```
int count = 0;
while(count < 4)
{
    cout << "Hello" << endl;
    count++;
}
```

First, the variable **count** is assigned a value of 0 and is used as the counter variable to control the **while** loop. The **while** loop follows and includes the Boolean expression, **count < 4**, within parentheses. The counter-controlled loop executes a block statement that is marked by an opening curly brace and a closing curly brace. The statements in the loop body display the word "Hello" and then increment **count**, which adds 1 to the counter variable.

Exercise 5-3: Using a Counter-Controlled **while** Loop

In this exercise, you use what you have learned about counter-controlled loops. Study the following code, and then answer Questions 1–4.



Remember that **number2 += number1;** is the same as **number2 = number2 + number1;**.

```
int number1 = 0;
int number2 = 0;
while(number1 < 6)
    number1++;
    number2 += number1;
```

1. What is the value of **number1** when the loop exits? _____
2. What is the value of **number2** when the loop exits? _____

3. If the statement `number1++` is changed to `++number1`, what is the value of `number1` when the loop exits? _____
4. What could you do to force the value of `number2` to be 21 when the loop exits?

76

Lab 5-1: Using a Counter-Controlled `while` Loop

In this lab, you use a counter-controlled `while` loop in a C++ program provided with the data files for this book. When completed, the program should print the numbers 0 through 10, along with their values multiplied by 2 and by 10. The data file contains the necessary variable declarations and some output statements.

1. Open the source code file named `Multiply.cpp` using Notepad or the text editor of your choice.
 2. Write a counter-controlled `while` loop that uses the loop control variable to take on the values 0 through 10. Remember to initialize the loop control variable before the program enters the loop.
 3. In the body of the loop, multiply the value of the loop control variable by 2 and by 10. Remember to change the value of the loop control variable in the body of the loop.
 4. Save the source code file in a directory of your choice, and then make that directory your working directory.
 5. Compile the source code file `Multiply.cpp`.
 6. Execute the program. Record the output of this program.
-
-

Using a Sentinel Value to Control a Loop

As you learned in Chapter 1 of *Programming Logic and Design*, a sentinel value is a value such as "Y" or "N" that a user must supply to stop a loop. To learn about sentinel values in C++, you will look at a program discussed in Chapter 5 of *Programming Logic and Design* and in Chapter 3 of this book. The program displays a payroll report for a small company. This program includes a `while` loop and uses a sentinel value to determine when the loop executes or when the loop is exited. The pseudocode is shown in Figure 5-1.

```

Declarations
    string name
    num gross
    num deduct
    num net
    num RATE = 0.25
    string QUIT = "XXX"
    string REPORT_HEADING = "Payroll Report"
    string COLUMN_HEADING = "Name Gross Deductions Net"
    string END_LINE = "**End of report"
    housekeeping()
    while not name = QUIT
        detailLoop()
    endwhile
    endOfJob()
stop

housekeeping()
    output REPORT_HEADING
    output COLUMN_HEADING
    input name
return

detailLoop()
    input gross
    deduct = gross * RATE
    net = gross - deduct
    output name, gross, deduct, net
    input name
return

endOfJob()
    output END_LINE
return

```

Figure 5-1 Pseudocode for a payroll report program

Note that a **priming read** is included in the `housekeeping()` method in the pseudocode shown in Figure 5-1. Recall that you perform a priming read before a loop executes to input a value that is then used to control the loop. When a priming read is used, the program must perform another read within the loop body to get the next input value. You can see the priming read, the loop, and the last output statement portion of the pseudocode translated to C++ in the following code sample:

```

cout << "Enter employee's name or XXX to quit: ";
cin >> name;
while(name != QUIT)
{
    // This is the work done in the detailLoop() function
    cout << "Enter employee's gross pay: ";
    cin >> gross;
    deduct = gross * RATE;
}

```

```
net = gross - deduct;
cout << "Name: " << name << endl;
cout << "Gross Pay: " << gross << endl;
cout << "Deductions: " << deduct << endl;
cout << "Net Pay: " << net << endl;
cout << "Enter employee's name or XXX to quit: ";
cin >> name;
}
// This is the work done in the endOfJob() method
cout << END_LINE;
```

In this code example, the variable named `name` is the loop control variable. It is assigned a value when the program instructs the user to `Enter employee's name or XXX to quit:` and reads the user's response. The loop control variable is tested with `name != QUIT`. If the user enters a name (any value other than `XXX`, which is the constant value of `QUIT`), then the test expression is `true` and the statements within the loop body execute. If the user enters `XXX` (the constant value of `QUIT`), which is the sentinel value, then the test expression is `false` and the loop is exited.



It is important for you to understand that lowercase `xxx` and uppercase `XXX` are different values.

The first statement in the loop instructs the user to enter an employee's gross pay. The program then retrieves the user's input and stores it in the variable named `gross`. The employee's deductions are calculated next and stored in the variable named `deduct` followed by the program calculating the employee's net pay and storing the value in the variable named `net`. Next, the program outputs the name of the employee followed by the employee's gross pay, deductions, and net pay.

The last two statements in the loop prompt the user for a new name and then retrieve the user's input and store it in the variable named `name`. This is the statement that changes the value of the loop control variable. The loop body ends when program control returns to the top of the loop, where the Boolean expression in the `while` statement is tested again. If the user enters the next employee's name at the last prompt, then the loop is entered again, and a new gross pay is input, followed by calculations that determine new values for deductions and net pay. Next, the program displays the name of the employee followed by this employee's gross pay, deductions, and net pay. The program then prompts the user to enter a new value for `name`. If the user enters `XXX`, then the test expression is `false`, and the loop body does not execute. When the loop is exited, the next statement to execute displays `***End of report` (the constant value of `END_LINE`.)

Exercise 5-4: Using a Sentinel Value to Control a `while` Loop

In this exercise, you use what you have learned about sentinel values. Study the following code, and then answer Questions 1–5.

```
int numToPrint, counter;
cout << "How many pages do you want to print? ";
cin >> numToPrint;
counter = 1;
while(counter <= numToPrint);
{
    cout << "Page Number " << counter << endl;
    counter++;
}
cout << "Value of counter is " << counter << endl;
```

1. What is the output if the user enters a 3?
-

2. What is the problem with this code, and how can you fix it?
-

3. Assuming you fix the problem, if the user enters 50 as the number of pages to print, what is the value of `counter` when the loop exits?
-

4. Assuming you fix the problem, if the user enters 0 as the number of pages to print, how many pages will print?
-

5. What is the output if the curly braces are deleted?
-

Lab 5-2: Using a Sentinel Value to Control a `while` Loop

In this lab, you write a `while` loop that uses a sentinel value to control a loop in a C++ program provided with the data files for this book. You also write the statements that make up the body of the loop. The source code file already contains the necessary variable declarations and output statements. You designed this program for the Hollywood Movie Rating Guide in Chapter 5, Exercise 15, in *Programming Logic and Design*. Each theater patron enters a value from 0 to 4 indicating the number of stars the patron awards to the Guide’s featured movie of the week. The program executes continuously until the theater manager enters a negative number to quit. At the end of the program, you should display the average star rating for the movie.

1. Open the source code file named `MovieGuide.cpp` using Notepad or the text editor of your choice.
2. Write the `while` loop using a sentinel value to control the loop, and write the statements that make up the body of the loop. The output statements within the loop have already been written for you.

3. Save this source code file in a directory of your choice, and then make that directory your working directory.
4. Compile the source code file `MovieGuide.cpp`.
5. Execute the program. Input the following: 0 3 4 4 1 1 2 -1
6. Record the average star rating for the movie.

Writing a for Loop in C++

In Chapter 5 of *Programming Logic and Design* you learned that a `for` loop is a **definite** loop; this means this type of loop will execute a definite number of times. The following is the syntax for a `for` loop in C++:

```
for(expression1; expression2; expression3)
    statement;
```

In C++, the `for` loop consists of three expressions that are separated by semicolons and enclosed within parentheses. The `for` loop executes as follows:

- The first time the `for` loop is encountered, the first expression is evaluated. Usually, this expression initializes a variable that is used to control the `for` loop.
- Next, the second expression is evaluated. If the second expression evaluates to `true`, the loop statement executes. If the second expression evaluates to `false`, the loop is exited.
- After the loop statement executes, the third expression is evaluated. The third expression usually increments or decrements the variable that you initialized in the first expression.
- After the third expression is evaluated, the second expression is evaluated again. If the second expression still evaluates to `true`, the loop statement executes again, and then the third expression is evaluated again.
- This process continues until the second expression evaluates to `false`.

The following code sample illustrates a C++ `for` loop. Notice the code uses a block statement in the `for` loop.

```
int number = 0;
int count;
const int NUM_LOOPS = 10;
for(count = 0; count < NUM_LOOPS; count++)
{
    number += count;
    cout << "Value of number is: " << number << endl;
}
```

In this `for` loop example, the variable named `count` is initialized to 0 in the first expression. The second expression is a Boolean expression that evaluates to `true` or `false`. When the expression `count < NUM_LOOPS` is evaluated the first time, the value of `count` is 0 and the result is `true`. The loop body is then entered. This is where a new value is computed and assigned to the variable named `number` and then displayed. The first time through the loop, the output is as follows: `Value of number is: 0`.

After the output is displayed, the third expression in the **for** loop is evaluated; this adds 1 to the value of **count**, making the new value of **count** equal to 1. When expression two is evaluated a second time, the value of **count** is 1. The program then tests to see if the value of **count** is less than **NUM_LOOPs**. This results in a **true** value and causes the loop body to execute again where a new value is computed for **number** and then displayed. The second time through the loop, the output is as follows: **Value of number is: 1.**

Next, expression three is evaluated; this adds 1 to the value of **count**. The value of **count** is now 2. Expression two is evaluated a third time and again is **true** because 2 is less than **NUM_LOOPs**. The third time through, the loop body changes the value of **number** to 3 and then displays the new value. The output is as follows: **Value of number is: 3.**

This process continues until the value of **count** becomes 10. At this time, 10 is not less than **NUM_LOOPs**, so the second expression results in a **false** value, and causes the program to exit the **for** loop.

The counter-controlled loop that displays the word "Hello" four times (which you studied in the "Using a Counter to Control a Loop" section of this chapter) can be rewritten using a **for** loop instead of the **while** loop. In fact, when you know how many times a loop will execute, it is considered a good programming practice to use a **for** loop instead of a **while** loop.

To rewrite the **while** loop as a **for** loop, you can delete the assignment statement **counter = 0;** because you initialize **counter** in expression one. You can also delete **counter++;** from the loop body because you increment **counter** in expression three. The program continues to print the word "Hello" in the body of the loop. The following code sample illustrates this **for** loop:

```
int counter;
for(counter = 0; counter < 4; counter++)
{
    cout << "Hello" << endl;
}
```



The curly braces are not required because now the loop body contains just one statement. However, it is considered a good programming practice to include them, as it makes the code more readable and may help prevent an error later if additional statements are added to the body of the loop.

Exercise 5-5: Using a **for** Loop

In this exercise, you use what you have learned about **for** loops. Study the following code, and then answer Questions 1–4.

```
const int NUM_LOOPs = 12;
int numTimes;
for(numTimes = 1; numTimes <= NUM_LOOPs; numTimes++)
{
    cout << "Value of numTimes is: " << numTimes << endl;
    numTimes++;
}
```

Answer the following four questions with *True* or *False*.

1. This loop executes 12 times._____
2. This loop could be written as a `while` loop._____
3. Changing the `<=` operator to `<` will make no difference in the output._____
4. This loop executes six times._____

Lab 5-3: Using a `for` Loop

In this lab, you work with the same C++ program you worked with in Lab 5-1. As in Lab 5-1, the completed program should print the numbers 0 through 10, along with their values multiplied by 2 and by 10. However, in this lab you should accomplish this using a `for` loop instead of a counter-controlled `while` loop.

1. Open the source code file named `NewMultiply.cpp` using Notepad or the text editor of your choice.
2. Write a `for` loop that uses the loop control variable to take on the values 0 through 10.
3. In the body of the loop, multiply the value of the loop control variable by 2 and by 10.
4. Save this source code file in a directory of your choice, and then make that directory your working directory.
5. Compile the source code file `NewMultiply.cpp`.
6. Execute the program. Is the output the same?

Writing a `do while` Loop in C++

In Chapter 5 of *Programming Logic and Design* you learned about the `do until` loop. C++ does not support a `do until` loop, but it does have a `do while` loop. The `do while` loop uses logic that can be stated as “*do a* while *b* is true.” This is similar to a `while` loop; however, there is a difference. When you use a `while` loop, the condition is tested before the statements in the loop body execute. When you use a `do while` loop, the condition is tested after the statements in the loop body execute once. As a result, you should choose a `do while` loop when your program logic requires the body of the loop to execute at least once. The body of a `do while` loop continues to execute as long as the expression evaluates to `true`. The `do while` syntax is as follows:

```
do
    statement;
while(expression);
```

The following `do while` loop is a revised version of the `while` loop you saw earlier, which prints the word "Hello" four times. In this version, the loop is rewritten as a `do while` loop.

```
int counter = 0;
const int NUM_LOOPS = 4;
do
{
    cout << "Hello" << endl;
    counter++;
} while(counter < NUM_LOOPS);
```

In this example, block statements are used in `do while` loops just as they are in `while` and `for` loops. When this loop is entered, the word "Hello" is printed, the value of `counter` is incremented, and then the value of `counter` is compared with the value 4. Notice that the word "Hello" will always be printed at least once because the loop control variable, `counter`, is compared to the value 4 at the bottom of the loop.

Exercise 5-6: Using a do while Loop

In this exercise, you use what you have learned about `do while` loops. Study the following code, and then answer Questions 1–4.

```
const int NUM_TIMES = 3;
int counter = 0;
do
{
    counter++;
    cout << "Strike " << counter << endl;
}while(counter < NUM_TIMES);
```

1. How many times does this loop execute? _____

2. What is the output of this program?

3. Is the output different if you change the order of the statements in the body of the loop, so that `counter++` comes after the output statement? _____

4. What is the loop control variable?

Lab 5-4: Using a do while Loop

In this lab, you work with the same C++ program you worked with in Labs 5-1 and 5-3. As in those earlier labs, the completed program should print the numbers 0 through 10, along with their values multiplied by 2 and by 10. However, in this lab you should accomplish this using a `do while` loop.



By writing the same program three different ways in this chapter, you have seen that a single problem can be solved in different ways.

1. Open the source code file named `NewestMultiply.cpp` using Notepad or the text editor of your choice.
 2. Write a `do while` loop that uses the loop control variable to take on the values 0 through 10.
 3. In the body of the loop, multiply the value of the loop control variable by 2 and by 10.
 4. Save this source code file in a directory of your choice, and then make that directory your working directory.
 5. Compile the source code file `NewestMultiply.cpp`.
 6. Execute the program. Is the output the same?
-

Nesting Loops

As the logic of your programs becomes more complex, you may find that you need to use nested loops. That is, you may need to include a loop within another loop. You have learned that when you use nested loops in a program, you must use multiple control variables to control the separate loops.

In Chapter 5 of *Programming Logic and Design*, you studied the design logic for a program that produces a quiz answer sheet. Some of the declarations and a section of the pseudocode for this program are as follows:

```
num PARTS = 5
num QUESTIONS = 3
string PART_LABEL = "Part "
string LINE = ". _____"
string QUIT = "ZZZ"
output quizName
partCounter = 1
while partCounter <= PARTS
    output PART_LABEL, partCounter
    questionCounter = 1
    while questionCounter <= QUESTIONS
        output questionCounter, LINE
        questionCounter = questionCounter + 1
    endwhile
    partCounter = partCounter + 1
endwhile
output "Enter next quiz name or ", QUIT, " to quit"
input quizName
```

This pseudocode includes two loops. The outer loop uses the loop control variable, `partCounter`, to control the loop using the sentinel value, 5 (constant value of `PARTS`). The inner loop uses the control variable `questionCounter` to keep track of the number of lines to print for the questions in a part of the quiz. Refer to Chapter 5 in *Programming Logic and Design* for a line-by-line description of the pseudocode. When you are sure you understand the logic, take a look at the code sample that follows. This code sample shows some of the C++ code for the Answer Sheet program.

```
int partCounter;
int questionCounter;
const int PARTS = 5;
const int QUESTIONS = 3;
const string PART_LABEL = "Part ";
const string LINE = ". _____";
partCounter = 1;
while(partCounter <= PARTS)
{
    cout << PART_LABEL << partCounter;
    questionCounter = 1;
    while(questionCounter <= QUESTIONS)
    {
        cout << questionCounter << LINE;
        questionCounter++;
    }
    partCounter++;
}
```

The entire C++ program is saved in a file named `AnswerSheet.cpp`. This file is included with the data files for this book. You may want to study the source code, compile it, and execute the program to experience how nested loops behave.

Exercise 5-7: Nesting Loops

In this exercise, you use what you have learned about nesting loops. Study the following code, and then answer Questions 1–4.

```
int sum = 0;
const int MAX_ROWS = 5, MAX_COLS = 7;
int rows, columns;
for(rows = 0; rows < MAX_ROWS; rows++)
    for(columns = 0; columns < MAX_COLS; columns++)
        sum += rows + columns;
cout << "Value of sum is " << sum << endl;
```

1. How many times does the outer loop execute?

2. How many times does the inner loop execute?

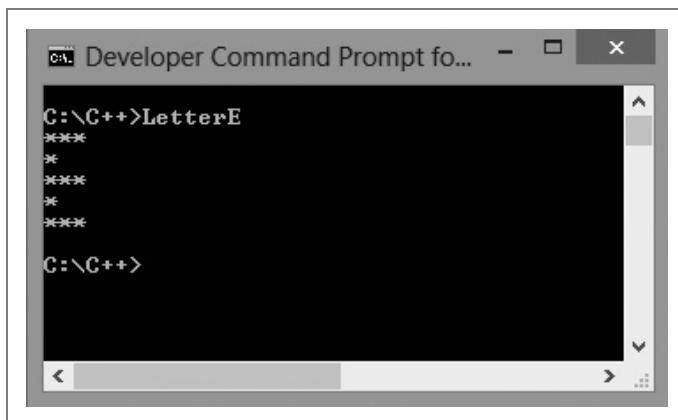
3. What is the value of `sum` printed by `cout`?

4. What would happen if you changed `rows++` and `columns++` to `++rows` and `++columns`?

86

Lab 5-5: Nesting Loops

In this lab, you add nested loops to a C++ program provided with the data files for this book. The program should print the outline of the letter *E*, as shown in Figure 5-2. The letter *E* is printed using asterisks, three across and five down. This program uses `cout << "*";` to print an asterisk and `cout << " ";` to print a space.



```
Developer Command Prompt for Localhost: MINGW32 Terminal
C:\>C++>LetterE
*** *
*** *
*** *
*** *
*** *
C:\>
```

Figure 5-2 Letter E printed by `LetterE.cpp`.

1. Open the source code file named `LetterE.cpp` using Notepad or the text editor of your choice.
2. Write the nested loops to control the number of rows and the number of columns that make up the letter *E*.
3. In the loop body, use a nested `if` statement to decide when to print an asterisk and when to print a space. The output statements have been written, but you must decide when and where to use them.
4. Save this source code file in a directory of your choice, and then make that directory your working directory.
5. Compile the source code file `LetterE.cpp`.

6. Execute the program. Your letter *E* should look like the letter *E* in Figure 5-2.
7. Modify the program to change the number of rows from five to seven and the number of columns from three to five. What does the letter *E* look like now?

Accumulating Totals in a Loop

You can use a loop in C++ to accumulate a total as your program executes. For example, assume that your computer science instructor has asked you to design and write a program that she can use to calculate an average score for the midterm exam she gave last week. To find the average test score, you need to add all the students' test scores, and then divide that sum by the number of students who took the midterm.

Note that the logic for this program should include a loop that will execute for each student in the class. In the loop, you get a student's test score as input and add that value to a total. After you get all of the test scores and accumulate the sum of all the test scores, you divide that sum by the number of students. You should plan to ask the user to input the number of student test scores that will be entered because your instructor wants to reuse this program using a different number of students each time it is executed.

As you review your work, you realize that the program will accumulate a sum within the loop and that you will also need to keep a count for the number of students. You learned in *Programming Logic and Design* that you add 1 to a counter each time a loop executes and that you add some other value to an accumulator. For this program, that other value added to the accumulator is a student's test score.



If `testTotal` is not initialized, it may contain an unknown value referred to as a “garbage” value. The C++ compiler issues a warning message for uninitialized variables.



You must calculate the average outside of the loop, not inside the loop. The only way you could calculate the average inside the loop is to do it each time the loop executes, but this is inefficient.

The following C++ code sample shows the loop required for this program. Notice the loop body includes an accumulator and a counter.

```
int numStudents, stuCount, testScore;
double testTotal, average;
// Get user input to control loop
cout << "Enter number of students: ";
cin >> numStudents;
// Initialize accumulator variable to 0
testTotal = 0;
// Loop for each student
for(stuCount = 0; stuCount < numStudents; stuCount++)
```

```
{  
    // Input student test score  
    cout << "Enter student's score: ";  
    cin >> testScore;  
    // Accumulate total of test scores  
    testTotal += testScore;  
}  
// Calculate average test score  
average = testTotal / stuCount;
```

88

If a user entered a 0, meaning 0 students took the midterm, the `for` loop would not execute because the value of `numStudents` is 0, and the value of `stuCount` is also 0.

In the code, you use the `cout` statement to ask your user to tell you how many students took the test. Then you use `cin` to retrieve the user input and store it in the variable named `numStudents`. Next, the accumulator, `testTotal`, is initialized to 0.

After the accumulator is initialized, the code uses a `for` loop and the loop control variable, `stuCount`, to control the loop. A `for` loop is a good choice because, at this point in the program, you know how many times the loop should execute. You use the `for` loop's first expression to initialize `stuCount`, and then the second expression is evaluated to see if `stuCount` is less than `numStudents`. If this is `true`, the body of the loop executes, using `cout` and `cin` again, this time asking the user to enter a test score and retrieving the input and storing it in `testScore`.

Next, you must add the value of `testScore` to the accumulator, `testTotal`. The loop control variable, `stuCount`, is then incremented, and the incremented value is tested to see if it is less than `numStudents`. If this is `true` again, the loop executes a second time. The loop continues to execute until the value of `stuCount < numStudents` is `false`. Outside the `for` loop, the program calculates the average test score by dividing `testTotal` by `stuCount`.

The entire C++ program is saved in a file named `TestAverage.cpp`. You may want to study the source code, compile it, and execute the program to experience how accumulators and counters behave.

Exercise 5-8: Accumulating Totals in a Loop

In this exercise, you use what you have learned about using counters and accumulating totals in a loop. Study the following code, and then answer Questions 1–4. The complete program is saved in the file named `Rainfall.cpp`. You may want to compile and execute the program to help you answer the questions.

```
const int DAYS_IN_WEEK = 7;  
for(counter = 1; counter <= DAYS_IN_WEEK; counter++)  
{  
    cout << "Enter rainfall amount for Day " + counter << ":";  
    cin >> rainfall;  
    cout << "Day " << counter << "rainfall amount is " <<  
        rainfall << " inches" << endl;  
    sum += rainfall;  
}  
// calculate average  
average = sum / DAYS_IN_WEEK;
```

1. What happens when you compile this program if the variable `sum` is not initialized with the value 0?

2. Could you replace `sum += rainfall;` with `sum = sum + rainfall;` ?

3. The variables `sum`, `rainfall`, and `average` should be declared to be what data type to calculate the most accurate average rainfall?

4. Could you replace `DAYS_IN_WEEK` in the statement `average = sum / DAYS_IN_WEEK;` with the variable named `counter` and still get the desired result? Explain.

Lab 5-6: Accumulating Totals in a Loop

In this lab, you add a loop and the statements that make up the loop body to a C++ program provided with the data files for this book. When completed, the program should calculate two totals: the number of left-handed people and the number of right-handed people in your class. Your loop should execute until the user enters the character `X` instead of `L` for left-handed or `R` for right-handed.

The inputs for this program are as follows: `R R R L L L R L R R L X`

Variables have been declared for you, and the input and output statements have been written.

1. Open the source code file named `LeftOrRight.cpp` using Notepad or the text editor of your choice.
 2. Write a loop and a loop body that allows you to calculate a total of left-handed and right-handed people in your class.
 3. Save this source code file in a directory of your choice, and then make that directory your working directory.
 4. Compile the source code file `LeftOrRight.cpp`.
 5. Execute the program using the data listed above. Record the results.
-

Using a Loop to Validate Input

In Chapter 5 of *Programming Logic and Design*, you learned that you cannot count on users to enter valid data in programs that ask them to enter data. You also learned that you should validate input from your user so you can avoid problems caused by invalid input.

If your program requires a user to enter a specific value, such as a "Y" or an "N" in response to a question, then your program should validate that your user entered an exact match to either "Y" or "N". You must also decide what action to take in your program if the user's input is not either "Y" or "N". As an example of testing for an exact match, consider the following code:

90

```
string answer;
cout << "Do you want to continue? Enter Y or N. ";
cin >> answer;
while(answer != "Y" && answer != "N")
{
    cout << "Invalid Response. Please type Y or N. ";
    cin >> answer;
}
```

In the example, the variable named `answer` contains your user's answer to the question "Do you want to continue? Enter Y or N." In the expression that is part of the `while` loop, you test to see if your user really did enter a "Y" or an "N". If not, the program enters the loop, tells the user he or she entered invalid input and then requests that he or she type "Y" or "N". The expression in the `while` loop is tested again to see if the user entered valid data this time. If not, the loop body executes again and continues to execute until the user enters valid input.

You can also verify user input in a program that requests a user to enter numeric data. For example, your program could ask a user to enter a number in the range of 1 to 4. It is very important to verify this numeric input, especially if your program uses the input in arithmetic calculations. What would happen if the user entered the word *one* instead of the numeral *1*? Or, what would happen if the user entered *100*? More than likely, your program will not run correctly. The following code illustrates how you can verify that a user enters correct numeric data.

```
int answer;
const int MIN_NUM = 1;
const int MAX_NUM = 4;
cout << "Please enter a number in the range of " << MIN_NUM <<
    " to " << MAX_NUM << ": ";
cin >> answer;
while(answer < MIN_NUM || answer > MAX_NUM)
{
    cout << "Number must be between " << MIN_NUM << " and " <<
        MAX_NUM << ". Please try again: ";
    cin >> answer;
}
```

Exercise 5-9: Validating User Input

In this exercise, you use what you have learned about validating user input to answer Questions 1–3.

1. You plan to use the following statement in a C++ program to validate user input:

```
while(inputString == "")
```

What would your user enter to cause this test to be true?

2. You plan to use the following statement in a C++ program to validate user input:

```
while(userAnswer == "N" || userAnswer == "n")
```

What would a user enter to cause this test to be true?

3. You plan to use the following statement in a C++ program to validate user input:

```
while(userAnswer < 1 || userAnswer > 10)
```

What would a user enter to cause this test to be true?

Lab 5-7: Validating User Input

In this lab, you will make additions to a C++ program provided with the data files for this book. The program is a guessing game. A random number between 1 and 10 is generated in the program. The user enters a number between 1 and 10, trying to guess the correct number. If the user guesses correctly, the program congratulates the user, and then the loop that controls guessing numbers exits; otherwise, the program asks the user if he or she wants to guess again. If the user enters a "Y", he or she can guess again. If the user enters "N", the loop exits. You can see that the "Y" or an "N" is the sentinel value that controls the loop. The entire program has been written for you. You need to add code that validates correct input, which is "Y" or "N", when the user is asked if he or she wants to guess a number, and a number in the range of 1 through 10 when the user is asked to guess a number.

1. Open the source code file named `GuessNumber.cpp` using Notepad or the text editor of your choice.
2. Write loops that validate input at all places in the code where the user is asked to provide input. Comments have been included in the code to help you identify where these loops should be written.
3. Save this source code file in a directory of your choice, and then make that directory your working directory.

4. Compile the source code file `GuessNumber.cpp`.
5. Execute the program. See if you can guess the randomly generated number. Execute the program several times to see if the random number changes. Also, test the program to verify that incorrect input is handled correctly. On your best attempt, how many guesses did you have to take to guess the correct number?_____

6

CHAPTER

Using Arrays in C++ Programs

After studying this chapter, you will be able to:

- ◎ Use arrays in C++ programs
- ◎ Search an array for a value
- ◎ Use parallel arrays in a C++ program

You should do the exercises and labs in this chapter after you have finished Chapter 6 of *Programming Logic and Design, Eighth Edition*, by Joyce Farrell. In this chapter, you learn how to use C++ to declare and initialize arrays. You then access the elements of an array to assign values and process them within your program. You also learn why it is important to stay within the bounds of an array. In addition, you study some programs written in C++ that implement the logic and design presented in *Programming Logic and Design*.

Array Basics

An **array** is a group of data items in which all items have the same data type, are referenced using the same variable name, and are stored in consecutive memory locations. To reference an individual element in an array, you use a subscript. Think of a **subscript** as the position number of a value within an array. It is important for you to know that in C++, subscript values begin with 0 (zero) and end with $n-1$, where n is the number of items stored in the array. You might be tempted to think that the first value in an array would be element number 1, but in fact it would be element number 0. The fifth element in an array would be element number 4.

To use an array in a C++ program, you must first learn how to declare an array, initialize an array with predetermined values, access array elements, and stay within the bounds of an array. In the next section you will focus on declaring arrays.

Declaring Arrays

Before you can use an array in a C++ program, you must first **declare** it. That is, you must give it a name and specify the data type for the data that will be stored in it. In some cases, you also specify the number of items that will be stored in the array. The following code shows how to declare two arrays, one named `cityPopulations` that will be used to store four `ints`, and another named `cities` that will be used to store four `strings`:

```
int cityPopulations[4];  
string cities[4];
```

As shown, you begin by specifying the data type of the items that will be stored in the array. The data type is followed by the name of the array and then a pair of square brackets. Within the square brackets, you see an integer value that specifies the number of elements this array can hold.

As shown in Figure 6-1, the compiler allocates enough consecutive memory locations to store four elements of data type `int` for the array named `cityPopulations`. If `cityPopulations[0]` is stored at memory address 1000, then the address of `cityPopulations[1]` is 1004 because each `int` requires 4 bytes of memory. Similarly, `cityPopulations[3]` is at address 1012.

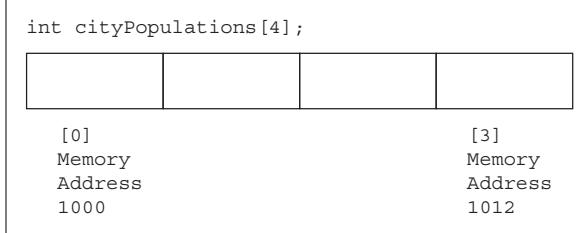


Figure 6-1 Memory allocation for the `cityPopulations` array



If an array is declared to store items of data type `double`, 8 bytes are allocated for each item in the array.

The `cityPopulations` array provides an example of how memory is allocated for arrays that contain primitive data types. Memory allocation is different for arrays of `string`s because a `string` is an object in C++, not a primitive data type. Using the Visual C++ compiler, 28 bytes of memory are allocated for a `string` object. This represents the size of the object.



The amount of memory allocated for a `string` object may be different on your system.

Additional memory for the characters that make up the `string` is allocated **dynamically** (as your program runs) when a value is stored in the `string` object. This is shown in Figure 6-2.

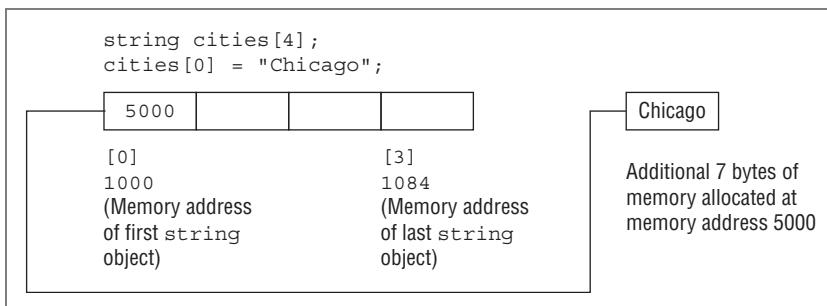


Figure 6-2 Memory allocation for the `cities` array

In Figure 6-2, the compiler allocates enough consecutive memory locations to store four `string` objects for the array named `cities`. If the address of `cities[0]` is 1000, the address of `cities[1]` is 1028, and the address of `cities[3]` is 1084 because each reference requires 28

bytes of memory. When a string is assigned to `cities[0]`, memory is dynamically allocated at another memory address. This address is then stored in the array. When the statement `cities[0] = "Chicago";` executes, the memory dynamically allocated for Chicago begins at address 5000, and then that address (5000) is stored in the first element of the array. An example of creating `string` objects is presented later in this chapter.

96



Dynamic memory allocation is an advanced topic that you will learn more about when you take additional courses in C++ programming.

Initializing Arrays

In C++, array elements are not automatically initialized to any value when the array is declared. Therefore, all of the elements of an array contain **garbage values**, which means they contain the values last stored at the memory location assigned to your array. Because these values are not useful in your program, it is a good idea to initialize arrays to all zeros for arrays that store numbers and to empty strings for arrays that store strings. Two double quotes with no space between, "", is the empty string in C++.



If you initialize an array when you declare it, you do not specify the array's size within the square brackets. Instead, you use an empty pair of square brackets.

You can and will sometimes want to initialize arrays with values that you choose. This can be done when you declare the array. To initialize an array when you declare it, use curly braces to surround a comma-delimited list of data items, as shown in the following example:

```
int cityPopulations[] = {9500000, 871100, 23900, 40100};  
string cities[] = {"Chicago", "Detroit", "Batavia", "Lima"};
```



These initializations create four elements for each array and assign values to them.

You can also use assignment statements to provide values for array elements after an array is declared, as in the following example:

```
cityPopulations[0] = 9500000;  
cities[0] = "Chicago";
```

A loop is often used to assign values to the elements in an array, as shown here:

```
int loopIndex;  
for(loopIndex = 0; loopIndex < 3; loopIndex++)  
{  
    cityPopulations[loopIndex] = 12345;  
    cities[loopIndex] = "AnyCity";  
}
```

The first time this loop is encountered, `loopIndex` is assigned the value 0. Because 0 is less than 3, the body of the loop executes, assigning the value 12345 to `cityPopulations[0]` and the value "AnyCity" to `cities[0]`. Next, the value of `loopIndex` is incremented and takes on the value 1. Because 1 is less than 3, the loop executes a second time and the value 12345 is assigned to `cityPopulations[1]`, and "AnyCity" is assigned to `cities[1]`. Each time the loop executes, the value of `loopIndex` is incremented. This allows you to access a different location in the array each time the body of the loop executes.

Accessing Array Elements

You need to access individual locations in an array when you assign a value to an array element, print its value, change its value, assign the value to another variable, and so forth. In C++, you use an integer expression placed in square brackets to indicate which element in the array you want to access. This integer expression is the subscript.



Remember that subscript values begin with 0 (zero) in C++.

The following C++ program declares two arrays of data type `double`, initializes an array of data type `double`, copies values from one array to another, changes several values stored in the array named `target`, and prints the values of the arrays named `source` and `target`. You can compile and execute this program if you like. It is stored in the file named `ArrayTest.cpp`.

```
#include <iostream>
#include <string>
using namespace std;
int main()
{
    double target[3];
    double source[] = {1.0, 5.5, 7.9};
    int loopIndex;
    // Copy values from source to target
    for(loopIndex = 0; loopIndex < 3; loopIndex++)
        target[loopIndex] = source[loopIndex];
    // Assign values to two elements of target
    target[0] = 2.0;
    target[1] = 4.5;
    // Print values stored in source and target
    for(loopIndex = 0; loopIndex < 3; loopIndex++)
    {
        cout << "Source " << source[loopIndex] << endl;
        cout << "Target " << target[loopIndex] << endl;
    }
}
```

Staying Within the Bounds of an Array

As a C++ programmer, you must be careful to ensure the subscript values you use to access array elements are within the legal bounds. Unlike most other programming languages, the C++ compiler does *not* check to make sure that a subscript used in your program is greater than or equal to 0 and less than the length of the array. For example, suppose you declare an array named `numbers` as follows:

```
int numbers[10];
```

In this case, the C++ compiler does not check to make sure the subscripts you use to access this array are integer values between 0 and 9.



When using a loop to access the elements in an array, be sure that the test you use to terminate the loop keeps you within the legal bounds, 0 to $n-1$, where n is the number of items stored in the array.

However, when you execute your program, if you use an array subscript that is not in the legal bounds, a garbage value is accessed. As an example, consider the highlighted operator in the following code, which is taken from the previous C++ program example:

```
double source[] = {1.0, 5.5, 7.9};  
int loopIndex;  
for(loopIndex = 0; loopIndex < 3; loopIndex++)
```

If you change the highlighted operator to `<=`, as shown below, your program will compile with no errors:

```
for(loopIndex = 0; loopIndex <= 3; loopIndex++)
```

A problem arises, however, when your program runs because the loop will execute when the value of `loopIndex` is 3. When you access the array element `source[3]`, you are outside the bounds of the array because there is no such element in this array, and a garbage value is stored at this location in the array.

Using Constants with Arrays

It is a good programming practice to use a named constant to help you stay within the bounds of an array when you write programs that declare and access arrays. The following example shows how to use a named constant:



It is a C++ convention to use all capital letters when naming constants.

```
const int NUM_ITEMS = 3;  
double target[NUM_ITEMS];  
int loopIndex;  
for(loopIndex = 0; loopIndex < NUM_ITEMS; loopIndex++)  
    target[loopIndex] = loopIndex + 10;
```

Exercise 6-1: Array Basics

Use what you have learned about declaring and initializing arrays to complete the following:

1. Write array declarations for each of the following:
 - a. Six grade point averages.

 - b. Seven last names.

 - c. 10 ages.

2. Declare and initialize arrays that store the following:
 - a. The whole numbers 2, 4, 6, 8, 10

 - b. The last names Carlson, Matthews, and Cooper

 - c. The prices 15.00, 122.00, and 7.50

3. Write an assignment statement that assigns the value 32 to the first element of the array of integers named `customerNumber`.

4. An array of `ints` named `numbers` is stored at memory location 4000. Where is `numbers[1]`? Where is `numbers[4]`?

Lab 6-1: Array Basics

In this lab, you complete a partially prewritten C++ program that uses an array. The program prompts the user to interactively enter eight batting averages, which the program stores in an array. The program should then find the minimum and maximum batting average stored in the array as well as the average of the eight batting averages. The data file provided for this lab includes the input statement and some variable declarations. Comments are included in the file to help you write the remainder of the program.

1. Open the source code file named `BattingAverage.cpp` using Notepad or the text editor of your choice.
2. Write the C++ statements as indicated by the comments.
3. Save this source code file in a directory of your choice, and then make that directory your working directory.
4. Compile the source code file `BattingAverage.cpp`.

5. Execute the program. Enter the following batting averages: **.299, .157, .242, .203, .198, .333, .270, .190**. The minimum batting average should be .157, and the maximum batting average should be .333. The average should be .2365.

100

Searching an Array for an Exact Match

One of the programs described in *Programming Logic and Design* uses an array to hold valid item numbers for a mail-order business. The idea is that when a customer orders an item, you can determine if the customer ordered a valid item number by searching through the array for that item number. This program relies on a technique called setting a flag to verify that an item exists in an array. The pseudocode and the C++ code for this program are shown in Figure 6-3.

```
start
    Declarations
        num item
        num SIZE = 6
        num VALID_ITEMS[SIZE] = 106, 108, 307,
            405, 457, 688
        num sub
        string foundIt
        num badItemCount = 0
        string MSG_YES = "Item available"
        string MSG_NO = "Item not found"
        num FINISH = 999
    getReady()
    while item <> FINISH
        findItem()
    endwhile
    finishUp()
stop

getReady()
    output "Enter item number or ", FINISH, " to quit"
    input item
return

findItem()
    foundIt = "N"
    sub = 0
    while sub < SIZE
        if item = VALID_ITEMS[sub] then
            foundIt = "Y"
        endif
        sub = sub + 1
    endwhile
    if foundIt = "Y" then
        output MSG_YES
    else
        output MSG_NO
        badItemCount = badItemCount + 1
    endif
    output "Enter next item number or ", FINISH, " to quit"
    input item
return

finishUp()
    output badItemCount, " items had invalid numbers"
return
```

Figure 6-3 Pseudocode and C++ code for the Mail Order program (continues)

(continued)

```

#include <iostream>
#include <string>
using namespace std;

int main()
{
    // Declare variables
    int item, badItemCount = 0;
    const int SIZE = 6;
    int VALID_ITEMS[] = {106, 108, 307, 405, 457, 688};
    int sub;
    bool foundIt = false;
    const string MSG_YES = "Item Available";
    const string MSG_NO = "Item not found";
    const int FINISH = 999;

    // This is the work done in the getReady() function
    // Get user input
    cout << "Enter item number or " << FINISH << " to quit ";
    cin >> item;

    // Loop through array
    while(item != FINISH)
    {
        // This is the work done in the findItem() function
        foundIt = false;
        sub = 0;
        while(sub < SIZE)
        {
            // Test to see if this item is a valid item
            if(item == VALID_ITEMS[sub])
                foundIt = true;    // Set flag to true
            sub += 1;             // Increment loop index
        }
        // Test value of foundIt
        if(foundIt == true)
            cout << MSG_YES << endl;
        else
        {
            cout << MSG_NO << endl;
            badItemCount++;
        }
        cout << "Enter item number or " << FINISH << " to quit ";
        cin >> item;
    }
    // This is the work done in the finishUp() function
    cout << badItemCount << " items had invalid numbers." << endl;
    return 0;
}

```

101

Figure 6-3 Pseudocode and C++ code for the Mail Order program

Notice that the equality operator (==) is used when comparing the `int` value in the first `if` statement and the `bool` value in the second `if` statement.



The program can be found in the file named `MailOrder.cpp`. You may want to compile and execute the program to see how it operates.

As shown in Figure 6-3, when you translate the pseudocode to C++, you make a few changes. In both the pseudocode and the C++ code, the variable named `foundIt` is the flag. However,

in the C++ code, you assign the value `false` instead of the string constant "N" to the variable named `foundIt`. This is because the variable named `foundIt` is declared as a variable of the `bool` type. The `bool` data type is one of the primitive data types found in C++ and is only used to store `true` and `false` values.

102



In Figure 6-3, the array declaration and the lines of code that set the flag are highlighted.



The variable named `foundIt` could be declared as a `string` as it was in the pseudocode. It is better to use a `bool` in C++ for `true` and `false` values.

Exercise 6-2: Searching an Array for an Exact Match

In this exercise you use what you have learned about searching an array for an exact match. Study the following code, and then answer Questions 1–4. Note that this code may contain errors.

```
string apples[] = {"Gala", "Rome", "Fuji", "Delicious"};
int foundIt = false, i;
const int MAX_APPLES = 4;
string inApple;
cout << "Enter apple type:";
cin >> inApple;
for(i = 0; i <= MAX_APPLES; i++)
{
    if(inApple == apples[i])
    {
        foundIt = true;
    }
}
```

1. Is the `for` loop written correctly?

If not, how can you fix it?

2. Which variable is the flag?

3. Is the flag variable declared correctly?

If not, what should you do to fix it?

-
4. Is the comparison in the `if` statement done correctly?
-

If not, how can you fix it?

103

Lab 6-2: Searching an Array for an Exact Match

In this lab, you use what you have learned about searching an array to find an exact match to complete a partially prewritten C++ program. The program uses an array that contains valid names for 10 cities in Michigan. You ask the user to enter a city name; your program then searches the array for that city name. If it is not found, the program should print a message that informs the user the city name is not found in the list of valid cities in Michigan.

The data file provided for this lab includes the input statements and the necessary variable declarations. You need to use a loop to examine all the items in the array and test for a match. You also need to set a flag if there is a match and then test the flag variable to determine if you should print the "Not a city in Michigan." message. Comments in the code tell you where to write your statements. You can use the Mail Order program in this chapter as a guide.

1. Open the source code file named `MichiganCities.cpp` using Notepad or the text editor of your choice.
2. Study the prewritten code to make sure you understand it.
3. Write a loop statement that examines the names of cities stored in the array.
4. Write code that tests for a match.
5. Write code that, when appropriate, prints the message "Not a city in Michigan.".
6. Save this source code file in a directory of your choice, and then make that directory your working directory.
7. Compile the source code file `MichiganCities.cpp`.
8. Execute the program using the following as input:

Chicago
Brooklyn
Watervliet
Acme

Parallel Arrays

As you learned in *Programming Logic and Design*, you use parallel arrays to store values and to maintain a relationship between the items stored in the arrays. Figure 6-4 shows that the student ID number stored in `stuID[0]` and the grade stored in `grades[0]` are related—student 56 received a grade of 99.

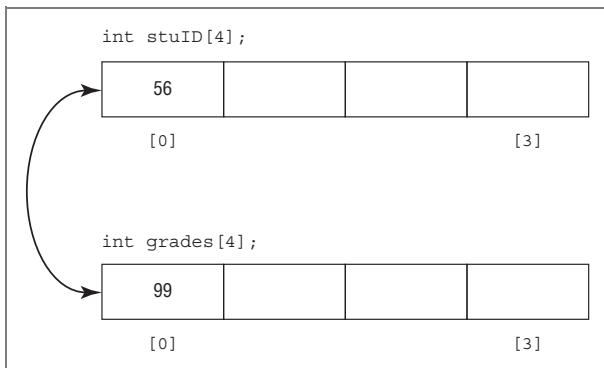


Figure 6-4 Parallel arrays

This relationship is established by using the same subscript value when accessing each array. Note that as the programmer, you must maintain this relationship in your code by always using the same subscript. C++ does not create or maintain the relationship.



Parallel arrays have the same number of elements, but they do not necessarily store items of the same data type.

One of the programs discussed in *Programming Logic and Design* is an expanded version of the Mail Order program discussed in the “Searching an Array for an Exact Match” section earlier in this chapter. In this expanded program, you need to determine the price of the ordered item and print the item number along with the price. You use parallel arrays to help you organize the data for the program. One array, `VALID_ITEMS`, contains six valid item numbers. The other array, `VALID_PRICES`, contains six valid prices. Each price is in the same position as the corresponding item number in the other array. When a customer orders an item, you search the `VALID_ITEMS` array for the customer’s item number. When the item number is found, you use the price stored in the same location of the `VALID_PRICES` array, and then output the item number and the price. The complete C++ program is stored in the file named `MailOrder2.cpp`. The pseudocode and C++ code that search the `VALID_ITEMS` array, use a price from the `VALID_PRICES` array, and then print the ordered item and its price are shown in Figure 6-5.



In Figure 6-5, the declaration for the second array and the lines of code that use the variable named `price` are highlighted.

```
start
    Declarations
        num item
        num price
        num SIZE = 6
        num VALID_ITEMS[SIZE] = 106, 108, 307,
            405, 457, 688
        num VALID_PRICES[SIZE] = 0.59, 0.99,
            4.50, 15.99, 17.50, 39.00
        num_sub
        string foundIt
        num badItemCount = 0
        string MSG_YES = "Item available"
        string MSG_NO = "Item not found"
        num FINISH = 999
    getReady()
    while item <> FINISH
        findItem()
    endwhile
    finishUp()
stop

getReady()
    output "Enter item number or ", FINISH, " to quit"
    input item
return

findItem()
    foundIt = "N"
    sub = 0
    while sub < SIZE
        if item = VALID_ITEMS[sub] then
            foundIt = "Y"
            price = VALID_PRICES[sub]
        endif
        sub = sub + 1
    endwhile
    if foundIt = "Y" then
        output MSG_YES
        output "The price of ", item, " is ", price
    else
        output MSG_NO
        badItemCount = badItemCount + 1
    endif
    output "Enter next item number or ", FINISH, " to quit"
    input item
return

finishUp()
    output badItemCount, " items had invalid numbers"
return
```

Figure 6-5 Pseudocode and C++ code for the Mail Order2 program (continues)

(continued)

106

```
#include <iostream>
#include <string>
using namespace std;
int main()
{
    int item, badItemCount = 0;
    double price;
    const int SIZE = 6;
    int VALID_ITEMS[] = {106, 108, 307, 405, 457, 688};
    double VALID_PRICES[] = {0.59, 0.99, 4.50, 15.99, 17.50, 39.00};
    int sub;
    bool foundIt = false;
    const string MSG_YES = "Item Available";
    const string MSG_NO = "Item not found";
    const int FINISH = 999;

    // This is the work done in the getReady() function
    cout << "Enter next item number or " << FINISH << " to quit ";
    cin >> item;

    while(item != FINISH)
    {
        // This is the work done in the findItem() function
        foundIt = false;
        sub = 0;
        while(sub < SIZE)
        {
            if(item == VALID_ITEMS[sub])
            {
                foundIt = true;
                price = VALID_PRICES[sub];
            }
            sub++;
        }
        if(foundIt == true)
        {
            cout << MSG_YES << endl;
            cout << "The price of " << item <<
                " is " << price << endl;
        }
        else
        {
            cout << MSG_NO << endl;
            badItemCount++;
        }
        cout << "Enter next item number or " << FINISH << " to quit ";
        cin >> item;
    }
    // This is the work done in the finishUp() function
    cout << badItemCount << " items had invalid numbers" << endl;
    return 0;
}
```

Figure 6-5 Pseudocode and C++ code for the Mail Order2 program

Exercise 6-3: Parallel Arrays

In this exercise, you use what you have learned about parallel arrays. Study the following code, and then answer Questions 1–4. Note that this code may contain errors.

```
string cities[] = "Chicago", "Rome", "Paris", "London";
int populations[] = 2695598, 2760665, 2190777, 7805417;
const int MAX_CITIES = 4;
int foundIt;
int i, x;
```

```
string inCity;
cout << "Enter city name: ";
cin >> inCity;
for(i = 0; i = MAX_CITIES; ++i)
{
    if(inCity == cities[i])
    {
        foundIt = i;
    }
}
cout << "Population for " << cities[foundIt] <<
    " is " << populations[foundIt] << endl;
```

1. Are the arrays declared and initialized correctly?

If not, how can you fix them?

2. Is the **for** loop written correctly?

If not, how can you fix it?

3. As written, how many times will the **for** loop execute?

4. How would you describe the purpose of the statement **foundIt = i;**?

Lab 6-3: Parallel Arrays

In this lab, you use what you have learned about parallel arrays to complete a partially completed C++ program. The program is described in Chapter 6, Exercise 7, in *Programming Logic and Design*. The program should either print the name and price for a coffee add-in from the Jumpin' Jive Coffee Shop or it should print the message "Sorry, we do not carry that.".

Read the problem description carefully before you begin. The data file provided for this lab includes the necessary variable declarations and input statements. You need to write the part of the program that searches for the name of the coffee add-in(s) and either prints the name and price of the add-in or prints the error message if the add-in is not found. Comments in the code tell you where to write your statements. You can use the expanded Mail Order2 program shown in Figure 6-5 as a guide.

1. Open the source code file named **JumpinJive.cpp** using Notepad or the text editor of your choice.
2. Study the prewritten code to make sure you understand it.
3. Write the code that searches the array for the name of the add-in ordered by the customer.

4. Write the code that prints the name and price of the add-in or the error message, and then write the code that prints the cost of the total order.
5. Save this source code file in a directory of your choice, and then make that directory your working directory.
6. Compile the source code file `JumpinJive.cpp`.
7. Execute the program using the following data and record the output:

Cream
Caramel
Whiskey
chocolate
Chocolate
Cinnamon
Vanilla



Remember that C++ is case sensitive, which means it distinguishes between uppercase letters and lowercase letters. This means, for example, that *cinnamon* is not the same as *Cinnamon*.

7

CHAPTER

File Handling and Applications

After studying this chapter, you will be able to:

- ◎ Understand computer files and perform file operations
- ◎ Work with sequential files and control break logic

In this chapter, you learn how to open and close files in C++, how to use C++ to read data from and write data to a file in a program, and how to work with sequential files in a C++ program. You should do the exercises and labs in this chapter after you have finished Chapter 7 in *Programming Logic and Design, Eighth Edition*.

110

File Handling

Business applications are often required to manipulate large amounts of data that are stored in one or more files. As you learned in Chapter 7 of *Programming Logic and Design*, data is organized in a hierarchy. At the lowest level of the hierarchy is a **field**, which is a group of characters. On the next level up is a **record**, which is a group of related fields. For example, you could write a program that processes employee records with each employee record consisting of three fields: the employee's first name, the employee's last name, and the employee's department number.

In C++, to use the data stored in a file, the program must first open the file and then read the data from the file. You use prewritten classes that are part of a C++ library to accomplish this. In the next section, you learn how to use these classes to open a file, close a file, read data from a file, and write data to a file.

Using Input and Output Classes

To use the classes that you need to perform file input (`ifstream`) and output (`ofstream`), you must add the following preprocessor directive at the beginning of your C++ program:

```
#include <fstream>
```

The classes included in the `fstream` class are prewritten for you by the C++ development team and allow you to simplify your programming tasks by creating objects using these classes. You can then use the attributes and methods of these objects in your C++ programs.

Opening a File for Reading

To open a file and read data into a C++ program, you declare a file input stream object (variable) as shown below:

```
ifstream data_in;
```



You can think of `ifstream` as the data type for input files. In Chapter 10, you will learn that `ifstream` is actually a class and `data_in` is an object.

You then use the `open` method (function) to specify the name of the file to open. Look at the following example:

```
data_in.open("inputFile.dat");
```

In the example, you use the `ifstream` object (`data_in`), followed by a dot (.), followed by the name of the method (`open`) to open a file for input. Within the parentheses, you place the name of the file you want to open within a pair of double quotes. This statement opens the file named `inputFile.dat` for reading. This means that the program can now read data from the file. In this example, the file named `inputFile.dat` must be saved in the same folder as the C++ program that is using the file. To open a file that is saved in a different folder, a path must be specified as in this example:

```
data_in.open("C:\myC++Programs\Chapter7\inputFile.dat");
```

Reading Data from an Input File

Once you have opened the input file, you are ready to read the data in the file. You can use the extraction operator `>>` just as you used it with `cin`.

Assume that the input file for a program is organized so an employee's first name is on one line, followed by his last name on the next line, followed by his salary on the third, as follows:

Tim

Moriarty

4000.00

To allow the program to read this data, you would write the following C++ code:

```
ifstream data_in;
string firstName, lastName;
double salary;
data_in.open("inputFile.dat");
data_in >> firstName;
data_in >> lastName;
data_in >> salary;
```

The first line in the example declares `data_in`, an `ifstream` object. The second line declares two `string` variables named `firstName` and `lastName`. The third line declares a `double` named `salary`. Next, `data_in` is used with the `open` method to open the file named `inputFile.dat`. The extraction operator `>>` is then used three times to read the three lines of input from the file (`inputFile.dat`) associated with `data_in`. After this code executes, the variable named `firstName` contains the value *Tim*, the variable named `lastName` contains the value *Moriarty*, and the variable named `salary` contains the value *4000.00*.

Reading Data Using a Loop and EOF

In a program that has to read large amounts of data, it is usually best to have the program use a loop. In the loop, the program continues to read from the file until EOF (end of file) is encountered. In C++, the `eof` method (function) returns a `true` value when EOF is reached and a `false` value when EOF has not been reached. The `eof` method can be used with `cin` for keyboard input or with an `ifstream` object, such as `data_in`.

The C++ code that follows shows how to use the `eof` method as part of a loop:

```
ifstream data_in;
string firstName;
data_in.open("inputFile.dat");
data_in >> firstName;
while(!(data_in.eof()))
{
    cout << firstName;
    data_in >> firstName;
}
```



For keyboard input, you enter the EOF character by typing Ctrl+Z (Windows) or Ctrl+D (UNIX/Linux).

In this example, a priming read is used on the fourth line, `data_in >> firstName;`, and then the `eof` method is used as the expression to be tested in the `while` loop. As long as the value returned by `eof` is not equal to `true`, the expression is `true` and the loop is entered. Notice the negation operator `!` is used to reverse the value returned by the `eof` method. As soon as EOF is encountered, the test becomes `false` and the program exits the loop. The parentheses that surround the `eof` method call are used to control precedence.

Opening a File for Writing

To write data from a C++ program to an output file, the program must first open a file. Similar to input files, you must first declare a file stream object (variable). For output files, you declare an output stream object as shown below:

```
ofstream data_out;
```



You can think of `ofstream` as the data type for output files. In Chapter 10, you will learn that `ofstream` is actually a class and `data_out` is an object.

You then use the `open` method (function) to specify the name of the file to open. Look at the following example:

```
data_out.open("outputFile.dat");
```

In the example, you use the `ofstream` object (`data_out`), followed by a dot (.), followed by the name of the method (`open`) to open a file for output. Within the parentheses, you place the name of the file you want to open within a pair of double quotes. This statement opens the file named `outputFile.dat` for writing. This means that the program can now write data to the file. In this example, the file named `outputFile.dat` is saved in the same folder as the C++ program that is using the file. To open a file that is saved in a different folder, a path must be specified as in this example:

```
data_out.open("C:\myC++Programs\Chapter7\outputFile.dat");
```

Writing Data to an Output File

Once you have opened the output file, you are ready to write data to the file. You can use the insertion operator `<<` just as you used it with `cout`.

As an example, assume that an employee's `firstName`, `lastName`, and `salary` have been read from an input file as in the previous example, and that the employee is to receive a 15 percent salary increase that is calculated as follows:

```
const double INCREASE = 1.15;
double newSalary;
newSalary = salary * INCREASE;
```

You now want to write the employee's `firstName`, `lastName`, and `newSalary` to the output file named `outputFile.dat`. The code that follows accomplishes this task:

```
ofstream data_out;
data_out.open("outputFile.dat");
data_out << firstName << endl;
data_out << lastName << endl;
data_out << newSalary << endl;
data_out.close();
```

The last line in the previous example uses `data_out` and the `close` method to close the output file. It is a good programming practice to close input and output files when they are no longer needed by a program. The C++ program `fileIOTest.cpp` shown in Figure 7-1 implements the file input and output operations discussed in this section.

```
// fileIOTest.cpp
// Input: inputFile.dat
// Output: outputFile.dat

#include <fstream>
#include <iostream>
#include <string>
using namespace std;
int main()
{
    string firstName;
    string lastName;
    double salary;
    double new_salary;
    const double INCREASE = 1.15;
    ifstream data_in;
    ofstream data_out;

    // Open input and output file
    data_in.open("inputFile.dat");
    data_out.open("outputFile.dat");

    // Read record from input file
    data_in >> firstName;
    data_in >> lastName;
    data_in >> salary;
    // Calculate new salary
    new_salary = salary * INCREASE;
    // Write record to output file
    data_out << firstName << endl;
    data_out << lastName << endl;
    data_out << new_salary << endl;
    // Close files
    data_in.close();
    data_out.close();
    return 0;
} // End of main function
```

Figure 7-1 Reading and writing file data

When writing code that opens files and writes to files, you need to be aware of two potential problems: first, the program might try to open a nonexistent file; and second, it might try to read beyond the EOF marker. You will learn how to handle these error situations as you learn more about C++. They are beyond the scope of this book.

Exercise 7-1: Opening Files and Performing File Input

In this exercise, you use what you have learned about opening a file and getting input into a program from a file. Study the following code, and then answer Questions 1–3.

```
1 ofstream data_in;
2 data_in.open(myDVDFile.dat);
3 string dvdName, dvdPrice, dvdShelf;
4 data_in >> dvdShelf;
5 data_in >> dvdPrice;
6 data_in >> dvdName;
```

Figure 7-2 Code for Exercise 7-1

1. Describe the error on line 1, and explain how to fix it.

2. Describe the error on line 2, and explain how to fix it.

3. Consider the following data from the input file `myDVDFile.dat`:

Fargo 8.00 1A
Amadeus 20.00 2C
Casino 7.50 3B

- a. What value is stored in the variable named `dvdName`?

-
- b. What value is stored in the variable named `dvdPrice`?

-
- c. What value is stored in the variable named `dvdShelf`?

-
- d. If there is a problem with the values of these variables, what is the problem and how could you fix it?

Lab 7-1: Opening Files and Performing File Input

In this lab, you open a file and read input from that file in a prewritten C++ program. The program should read and print the names of flowers and whether they are grown in shade or sun. The data is stored in the input file named `flowers.dat`.

1. Open the source code file named `Flowers.cpp` using Notepad or the text editor of your choice.
2. Declare the variables you will need.

- 116
3. Write the C++ statements that will open the input file `flowers.dat` for reading.
 4. Write a `while` loop to read the input until EOF is reached.
 5. In the body of the loop, print the name of each flower and where it can be grown (sun or shade).
 6. Save this source code file in a directory of your choice, and then make that directory your working directory.
 7. Compile the source code file `Flowers.cpp`.
 8. Execute the program.

Understanding Sequential Files and Control Break Logic

As you learned in Chapter 7 of *Programming Logic and Design* a **sequential file** is a file in which records are stored one after another in some order. The records in a sequential file are organized based on the contents of one or more fields, such as ID numbers, part numbers, or last names.

A **single-level control break** program reads data from a sequential file and causes a break in the logic based on the value of a single variable. In Chapter 7 of *Programming Logic and Design*, you learned about techniques you can employ to implement a single-level control break program. Be sure you understand these techniques before you continue on with this chapter. The program described in Chapter 7 of *Programming Logic and Design* that produces a report of customers by state is an example of a single-level control break program. This program reads a record for each client, keeps a count of the number of clients in each state, and prints a report. As shown in Figure 7-3, the report generated by this program includes clients' names, cities, and states, along with a count of the number of clients in each state.

Company Clients by State of Residence			
Name	City	State	
Albertson	Birmingham	Alabama	
Davis	Birmingham	Alabama	
Lawrence	Montgomery	Alabama	
		Count for Alabama	3
Smith	Anchorage	Alaska	
Young	Anchorage	Alaska	
Davis	Fairbanks	Alaska	
Mitchell	Juneau	Alaska	
Zimmer	Juneau	Alaska	
		Count for Alaska	5
Edwards	Phoenix	Arizona	
		Count for Arizona	1

Figure 7-3 Control break report with totals after each state

Each client record is made up of the following fields: Name, City, and State. Note the following example records, each made up of three lines:

Albertson
Birmingham
Alabama
Lawrence
Montgomery
Alabama
Smith
Anchorage
Alaska

Remember that input records for a control break program are usually stored in a data file on a storage device, such as a disk, and the records are sorted according to a predetermined control break variable. For example, the control break variable for this program is `state`, so the input records would be sorted according to `state`.

Figure 7-4 includes the pseudocode for the Client By State program, and Figure 7-5 includes the C++ code that implements the program.

```
Start
  Declarations
    InputFile inFile
    string TITLE = "Company Clients by State of Residence"
    string COL_HEADS = "Name    City    State"
    string name
    string city
    string state
    num count = 0
    String oldState
    getReady()
    while not eof
      produceReport()
    endwhile
    finishUp()
  stop

  getReady()
    output TITLE
    output COL_HEADS
    open inFile "ClientsByState.dat"
    input name, city, state from inFile
    oldState = state
  return
```

Figure 7-4 Client By State program pseudocode (*continues*)

(continued)

118

```
produceReport()
    if state <> oldState then
        controlBreak()
    endif
    output name, city, state
    count = count + 1
    input name, city, state from inFile
return

controlBreak()
    output "Count for ", oldState, count
    count = 0
    oldState = state
return

finishUp()
    output "Count for ", oldState, count
    close inFile
return
```

Figure 7-4 Client By State program pseudocode

```
1 // ClientByState.cpp - This program creates a report that lists
2 // clients with a count of the number of clients for each state.
3 // Input: client.dat
4 // Output: Report
5
6 #include <iostream>
7 #include <iostream>
8 #include <string>
9 using namespace std;
10
11 int main()
12 {
13     // Declarations
14     ifstream fin;
15     const string TITLE = "\n\nCompany Clients by State of Residence\n\n";
16     string name = "", city = "", state = "";
17     int count = 0;
18     string oldState = "";
19     bool done;
```

Figure 7-5 Client By State program written in C++ (continues)

(continued)

```
21     fin.open("client.dat");
22     // This is the work done in the getReady() function
23     cout << TITLE << endl;
24     fin >> name;
25     if(!(fin.eof()))
26     {
27         fin >> city;
28         fin >> state;
29         done = false;
30         oldState = state;
31     }
32     else
33         done = true;
34     while(done == false)
35     {
36         // This is the work done in the produceReport() function
37         if(state != oldState)
38         {
39             // This is the work done in the controlBreak() function
40             cout << "\t\t\tCount for " << oldState << " " << count << endl;
41             count = 0;
42             oldState = state;
43         }
44         cout << name << " " << city << " " << state << endl;
45         count++;
46         fin >> name;
47         if(!(fin.eof()))
48         {
49             fin >> city;
50             fin >> state;
51             done = false;
52         }
53         else
54             done = true;
55     }
56     // This is the work done in the finishUp() function
57     cout << "\t\t\tCount for " << oldState << " " << count << endl;
58     fin.close();
59
60     return 0;
61 } // End of main() class.
```

Figure 7-5 Client By State program written in C++

As you can see in Figure 7-5, the C++ program begins on line 1 with comments that describe what the program does. (The line numbers shown in this program are not part of the C++ code. They are included for reference only.) The program also includes comments that describe the program's input and output. Next comes the C++ code that defines the `main()` function (line 11).

Within the `main()` function, lines 14 through 19 declare variables and initialize them when appropriate. Line 21 opens the input file named `client.dat`. Lines 23 through 33 include the work done in the `getReady()` function, which includes printing the heading for the report this program generates and performing a priming read. You learned about performing a priming read in Chapter 3 of this book and in Chapter 3 of *Programming Logic and Design*.

Notice that the C++ code in the priming read (lines 24 through 28) is a little different than the pseudocode. An `if` statement is used on line 25 to test if a client's name was read from the input file or if EOF was encountered. If EOF is not encountered, the result of this test will be `true`, causing the execution of the input statements that read the city and state from the input file. The `bool` value `false` is also assigned to the variable named `done` on line 29, followed by assigning the current value of `state` to the variable named `oldState` on line 30. Remember that the variable `state` serves as the control break variable. If EOF is encountered, the result of this test will be `false`, causing the `bool` value `true` to be assigned to the variable named `done` on line 33. The `bool` variable named `done` is used later in the program to control the `while` loop.

Next comes the `while` loop (line 34), which continues to execute as long as the value of the `bool` variable `done` is `false`. The body of the `while` loop contains the work done in the `produceReport()` function. First, on line 37, an `if` statement tests the control break variable, `state`, to determine if the record the program is currently working with has the same state as the previous record's state. If it does not, this indicates the beginning of a new state. As a result, the program performs the work done in the `controlBreak()` function (lines 40 through 42). The work of the `controlBreak()` function does the following:

1. Prints the value of the variable named `count` that contains the count of clients in the current state (line 40)
2. Assigns the value 0 to the variable named `count` to prepare for the next state
3. Assigns the value of the variable named `state` to the variable named `oldState` to prepare for the next state

If the record the program is currently working with has the same state as the previous record's state, the `controlBreak()` function's work is not performed. Whether or not the current record's state is the same state as the previous record's state, the next statement to execute (line 44) prints the client's name, city, and state. Then the variable named `count` is incremented on line 45 followed by the program reading the next client's record on lines 46 through 51 using the same technique as the priming read.

The condition in the `while` loop on line 34 is then tested again, causing the loop to continue executing until the value of the variable named `done` is `true`. The variable named `done` is assigned the value `true` when the program encounters EOF when reading from the input file on line 54.

When the `while` loop is exited, the last section of the program executes. This consists of the work done in the `finishUp()` function and comprises:

- Printing the value of the variable named `count` (which is the count of the clients in the last state in the input file) on line 57
- Closing the input file (line 58)

Exercise 7-2: Accumulating Totals in Single-Level Control Break Programs

In this exercise, you will use what you have learned about accumulating totals in a single-level control break program. Study the following code, and then answer Questions 1–4.

```
if(sectionNum != oldSectionNum)
{
    System.out.println("Section Number " + oldSectionNum);
    totalSections = 1;
    oldSectionNum = sectionNum;
}
```

1. What is the control break variable?

-
2. The value of the control break variable should never be changed. True or false?
-

3. Is `totalSections` being calculated correctly?
-

If not, how can you fix the code?

4. In a control break program, it doesn't matter if the records in the input file are in a specified order. True or false?
-

Lab 7-2: Accumulating Totals in Single-Level Control Break Programs

In this lab, you will use what you have learned about accumulating totals in a single-level control break program to complete a C++ program. The program should produce a report for a supermarket manager to help her keep track of hours worked by her part-time employees. The report should include the day of the week and the number of hours worked for each employee for each day of the week and the total hours for the day of the week. The report should look similar to the one shown in Figure 7-6.

The screenshot shows a terminal window titled "Developer Command Prompt for VS2012". The window displays the output of a C++ program named "SuperMarket". The program prompts the user to enter the day of the week or a control break command ("done") and the hours worked. It then prints the day name, the hours worked, and the cumulative total for each day. The process repeats for all days of the week, ending with a total for the week and a final total.

```
C:\>C++>SuperMarket

WEEKLY HOURS WORKED

Enter day of week or done to quit: Monday
Enter hours worked: 6
Monday 6
Day Total 6

Enter a day of week or done to quit: Tuesday
Enter hours worked: 2
Tuesday 2
Tuesday 3
Day Total 5

Enter a day of week or done to quit: Wednesday
Enter hours worked: 5
Wednesday 5
Wednesday 3
Day Total 8

Enter a day of week or done to quit: Thursday
Enter hours worked: 6
Thursday 6
Day Total 6

Enter a day of week or done to quit: Friday
Enter hours worked: 3
Friday 3
Friday 5
Day Total 8

Enter a day of week or done to quit: Saturday
Enter hours worked: 7
Saturday 7
Saturday 7
Day Total 21

Enter a day of week or done to quit: Sunday
Enter hours worked: 0
Sunday 0
Day Total 0

Enter a day of week or done to quit: done
Day Total 0

C:\>
```

Figure 7-6 Super Market program report

The student file provided for this lab includes the necessary variable declarations and input and output statements. You need to implement the code that recognizes when a control break should occur. You also need to complete the control break code. Be sure to accumulate the daily totals for all days in the week. Comments in the code tell you where to write your code. You can use the Client By State program in this chapter as a guide for this new program.

1. Open the source code file named `SuperMarket.cpp` using Notepad or the text editor of your choice.
2. Study the prewritten code to understand what has already been done.

3. Write the control break code, including the code for the `dayChange()` function, in the `main()` function.
4. Save this source code file in a directory of your choice, and then make that directory your working directory.
5. Compile the source code file, `SuperMarket.cpp`.
6. Execute the program using the following input values:
Monday—6 hours (employee 1)
Tuesday—2 hours (employee 1), 3 hours (employee 2)
Wednesday—5 hours (employee 1), 3 hours (employee 2)
Thursday—6 hours (employee 1)
Friday—3 hours (employee 1), 5 hours (employee 2)
Saturday—7 hours (employee 1), 7 hours (employee 2), 7 hours (employee 3)
Sunday—0 hours

8

CHAPTER

Advanced Array Techniques

After studying this chapter, you will be able to:

- ◎ Explain the need to sort data
- ◎ Swap data values in a program
- ◎ Create a bubble sort in C++
- ◎ Work with multidimensional arrays

In this chapter, you review why you would want to sort data, how to use C++ to swap two data values in a program, how to create a bubble sort in a C++ program, and how to use multidimensional arrays. You should do the exercises and labs in this chapter after you have finished Chapter 8 in *Programming Logic and Design, Eighth Edition*.

126

Sorting Data

Data records are always stored in some order, but possibly they are not in the order in which you want to process or view them in your program. When this is the case, you need to give your program the ability to arrange (sort) records in a useful order. For example, the inventory records you need to process might be stored in product-number order, but you might need to produce a report that lists products from lowest cost to highest cost. That means your program needs to sort the records by cost.

Sorting makes searching for records easier and more efficient. A human can usually find what he or she is searching for by simply glancing through a group of data items, but a program must look through a group of data items one by one, making a decision about each one. When searching unsorted records for a particular data value, a program must examine every single record until it either locates the data value or determines that it does not exist. However, when searching sorted records, the program can quickly determine when to stop searching, as shown in the following step-by-step scenario:

1. The records used by your program are sorted by product number.
2. The user is searching for product number 12367.
3. The program locates the record for product number 12368 but has not yet found product number 12367.
4. The program determines that the record for product number 12367 does not exist, and therefore stops searching through the list.

Many search algorithms require that data be sorted before it can be searched. (An **algorithm** is a plan for solving a problem.) You can choose from many algorithms for sorting and searching for data. In *Programming Logic and Design*, you learned how to swap data values in an array, and you also learned about the bubble sort. Both of these topics are covered in this book.

Swapping Data Values

When you swap values, you place the value stored in one variable into a second variable, and then you place the value that was originally stored in the second variable in the first variable. You must also create a third variable to temporarily hold one of the values you want to swap so a value is not lost. For example, if you try to swap values using the following code, you will lose the value of `score2`:

```
int score1 = 90;  
int score2 = 85;
```

```
score2 = score1; // The value of score2 is now 90  
score1 = score2; // The value of score1 is also 90
```

However, if you use a variable to temporarily hold one of the values, the swap is successful. This is shown in the following code:

```
int score1 = 90;  
int score2 = 85;  
int temp;  
temp = score2; // The value of temp is 85  
score2 = score1; // The value of score2 is 90  
score1 = temp; // The value of score1 is 85
```

127

Exercise 8-1: Swapping Values

In this exercise, you use what you have learned about swapping values to answer the following question.

1. Suppose you have declared and initialized two `string` variables, `product1` and `product2`, in a C++ program. Now, you want to swap the values stored in `product1` and `product2` but only if the value of `product1` is greater than the value of `product2`. Write the C++ code that accomplishes this task. The declarations are as follows:

```
string product1 = "hammer";  
string product2 = "wrench";
```

Lab 8-1: Swapping Values

In this lab, you complete a C++ program that swaps values stored in three `int` variables and determines maximum and minimum values. The C++ file provided for this lab contains the necessary variable declarations, as well as the input and output statements. You want to end up with the smallest value stored in the variable named `first` and the largest value stored in the variable named `third`. You need to write the statements that compare the values and swap them if appropriate. Comments included in the code tell you where to write your statements.

1. Open the source code file named `Swap.cpp` using the text editor of your choice.
2. Write the statements that test the first two integers, and swap them if necessary.
3. Write the statements that test the second and third integer, and swap them if necessary.
4. Write the statements that test the first and second integers again, and swap them if necessary.

5. Save this source code file in a directory of your choice, and then make that directory your working directory.
6. Compile the source code file `Swap.cpp`.
7. Execute the program using the following sets of input values, and record the output:

101	22	-23
630	1500	9
21	2	2

Using a Bubble Sort

A bubble sort is one of the easiest sorting techniques to understand. However, while it is logically simple, it is not very efficient. If the list contains n values, the bubble sort will make $n - 1$ passes over the list. For example, if the list contains 100 values, the bubble sort will make 99 passes over the data. During each pass, it examines successive overlapped pairs and swaps or exchanges those values that are out of order. After one pass over the data, the heaviest (largest) value sinks to the bottom and is in the correct position in the list.

In *Programming Logic and Design*, you learned several ways to refine the bubble sort. One way is to reduce unnecessary comparisons by ignoring the last value in the list in the second pass through the data, because you can be sure that it is already positioned correctly. On the third pass, you can ignore the last two values in the list because you know they are already positioned correctly. Thus, in each pass, you can reduce the number of items to be compared, and possibly swapped, by one.

Another refinement to the bubble sort is to eliminate unnecessary passes over the data in the list. When items in the array to be sorted are not entirely out of order, it may not be necessary to make $n - 1$ passes over the data because after several passes, the items may already be in order. You can add a flag variable to the bubble sort, and then test the value of that flag variable to determine whether any swaps have been made in any single pass over the data. If no swaps have been made, you know that the list is in order; therefore, you do not need to continue with additional passes.

You also learned about using a constant for the size of the array to make your logic easier to understand and your programs easier to change and maintain. Finally, you have learned how to sort a list of varying size by counting the number of items placed in the array as you read in items.

All of these refinements are included in the pseudocode for the Score Sorting program in Figure 8-1. The C++ code that implements the Score Sorting logic is provided in Figure 8-2. The line numbers shown in Figure 8-2 are not part of the C++ code. They are provided for reference only.

```
start
    num SIZE = 100
    num scores[SIZE]
    num x
    num y
    num temp
    num numberOfEls = 0
    num comparisons
    num QUIT = 999
    String didSwap
    fillArray()
    sortArray()
    displayArray()
stop

fillArray()
    x = 0
    output "Enter a score or ", QUIT, " to quit "
    input scores[x]
    x = x + 1
    while x < SIZE AND scores[x] <> QUIT
        output "Enter a score or ", QUIT, " to quit "
        input scores[x]
        x = x + 1
    endwhile
    numberOfEls = x
    comparisons = numberOfEls - 1
return

void sortArray()
    x = 0
    didSwap = "Yes"
    while didSwap = "Yes"
        x = 0
        didSwap = "No"
        while x < comparisons
            if scores[x] > scores[x + 1] then
                swap()
                didSwap = "Yes"
            endif
            x = x + 1
        endwhile
    endwhile
return
```

Figure 8-1 Pseudocode for the Score Sorting program (*continues*)

(continued)

130

```
void swap()
    temp = scores[x + 1]
    scores[x + 1] = scores[x]
    scores[x] = temp
return

void displayArray()
    x = 0
    while x < numberofEls
        output scores[x]
        x = x + 1
    endwhile
return
```

Figure 8-1 Pseudocode for the Score Sorting program

```
1 // StudentScores.cpp - This program interactively reads a
2 // variable number of student test scores, stores the
3 // scores in an array, and then sorts the scores in
4 // ascending order.
5 // Input: Interactive
6 // Output: Sorted list of student scores.
7
8 #include <iostream>
9 using namespace std;
10
11 int main()
12 {
13     // Declare variables.
14     // Maximum size of array
15     const int SIZE = 100;
16     // Array of student scores.
17     int scores[SIZE];
18     int x;
19     int temp;
20     // Actual number of elements in array.
21     int numberofEls = 0;
22     int comparisons;
23     const int QUIT = 999;
24     bool didSwap;
25
26     // Work done in the fillArray function
27     x = 0;
28     cout << "Enter a score or " << QUIT << " to quit ";
29     cin >> scores[x];
30     x++;
```

Figure 8-2 C++ code for the Score Sorting program (continues)

(continued)

```

31     while(x < SIZE && scores[x-1] != QUIT)
32     {
33         cout << "Enter a score or " << QUIT << " to quit ";
34         cin >> scores[x];
35         x++;
36     } // End of input loop.
37     numberOFEls = x-1;
38     comparisons = numberOFEls -1;
39
40     // Work done in the sortArray() function
41     didSwap = true; // Set flag to true.
42     // Outer loop controls number of passes over data.
43     while(didSwap == true) // Test flag.
44     {
45         x = 0;
46         didSwap = false;
47         // Inner loop controls number of items to compare.
48         while(x < comparisons)
49         {
50             if(scores[x] > scores[x+1]) // Swap?
51             {
52                 // Work done in the swap() function
53                 temp = scores[x + 1];
54                 scores[x+1] = scores[x];
55                 scores[x] = temp;
56                 didSwap = true;
57             }
58             x++; // Get ready for next pair.
59         }
60         comparisons--;
61     }
62
63
64     // Work done in the displayArray() function
65     x = 0;
66     while(x < numberOFEls)
67     {
68         cout << scores[x] << endl;
69         x++;
70     }
71     return 0;
72 } // End of main() function.

```

Figure 8-2 C++ code for the Score Sorting program

The `main()` Function

As shown in Figure 8-2, the `main()` function (line 11) declares variables and performs the work of the program. The variables include the following:

- A constant named `SIZE`, initialized with the value 100, which represents the maximum number of items this program can sort
- An array of data type `int` named `scores` that is used to store up to a maximum of `SIZE` (100) items to be sorted

- 132
- An `int` variable named `x` that is used as the array subscript
 - an `int` variable named `temp` that is used to swap the values stored in the array
 - An `int` named `numberOfEls` that is used to hold the actual number of items stored in the array
 - An `int` named `comparisons` that is used to control the number of comparisons that should be done
 - An `int` constant named `QUIT`, initialized to 999, that is used to control the `while` loop
 - A `bool` named `didSwap` that is used as a flag to indicate when a swap has taken place

After these variables are declared, the work done in the `fillArray()` function begins on line 27. The `fillArray()` work is responsible for filling up the array with items to be sorted. On line 41, the work done in the `sortArray()` function begins. This work is responsible for sorting the items stored in the `scores` array. Lastly, the work done in the `displayArray()` function begins on line 65 and is responsible for displaying the sorted scores on the user's screen.

The `fillArray()` Function

The work done in the `fillArray()` function, which begins on line 27 in Figure 8-2, is responsible for (1) storing the data in the array, and (2) counting the actual number of elements placed in the array. The `fillArray()` function assigns the value 0 to the variable named `x` and then performs a priming read (lines 28 and 29) to retrieve the first student score from the user and to also store the score in the array named `scores` at location `x` on line 29. Notice the array subscript variable `x` is initialized to 0 on line 27 because the first position in an array is position 0. Also, notice the variable named `x` is incremented on line 30 because it is used to count the number of scores entered by the user of the program.

On line 31, the condition that controls the `while` loop is tested. The `while` loop executes as long as the number of scores input by the user (represented by the variable named `x`) is less than `SIZE(100)` and as long as the user has not entered 999 (the value of the constant `QUIT`) for the student score. If `x` is less than `SIZE` and the user does not want to quit, there is enough room in the array to store the student score. In that case, the program retrieves the next student score, and stores the score in the array named `scores` at location `x` on line 34. The program then increments the value of `x` (line 35) to get ready to store the next student score in the array. The loop continues to execute until the user enters the value 999 or until there is no more room in the array.

When the program exits the loop, the value of `x-1` is assigned to the variable named `numberOfEls` on line 37. Notice that `x` is used as the array subscript and that its value is incremented every time the `while` loop executes, including when the user enters the value 999 in order to quit; therefore, `x` represents the number of student scores the user entered *plus one*. On line 38 the value of `numberOfEls -1` is assigned to the variable named `comparisons` and represents the maximum number of elements the bubble sort will compare on a pass over the data stored in the array. It ensures that the program does not attempt to compare item `x` with item `x+1`, when `x` is the last item in the array.

The `sortArray()` Function

The work done in the `sortArray()` function begins on line 41 and uses a refined bubble sort to rearrange the student scores in the array named `scores` to be in ascending order. Refer to Figure 8-1, which includes the pseudocode, and Figure 8-2, which includes the C++ code that implements the `sortArray()` function.

Line 41 initializes the flag variable `didSwap` to `true`, because, at this point in the program, it is assumed that items will need to be swapped.

The outer loop (line 43), `while(didSwap == true)`, controls the number of passes over the data. This logic implements one of the refinements discussed earlier—eliminating unnecessary passes over the data. As long as `didSwap` is `true`, the program knows that swaps have been made and that, therefore, the data is still out of order. Thus, when `didSwap` is `true`, the program enters the loop. The first statement in the body of the loop (line 45) is `x = 0;.` The program assigns the value 0 to `x` because `x` is used as the array subscript. Recall that in C++, the first subscript in an array is number 0.

Next, to prepare for comparing the elements in the array, line 46 assigns the value `false` to `didSwap`. This is necessary because the program has not yet swapped any values in the array on this pass. The inner loop begins on line 48. The test, `x < comparisons`, controls the number of pairs of values in the array the program compares on one pass over the data. This implements another of the refinements discussed earlier—reducing unnecessary comparisons. The last statement in the outer loop (line 60), `comparisons--;`, decrements the value of `comparisons` by 1 each time the outer loop executes. The program decrements `comparisons` because, when a complete pass is made over the data, it knows an item is positioned in the array correctly. Comparing the value of `comparisons` with the value of `x` in the inner loop reduces the number of necessary comparisons made when this loop executes.

On line 50, within the inner loop, adjacent items in the array are accessed and compared using the subscript variable `x` and `x+1`. The adjacent array items are compared to see if the program should swap them. If the values should be swapped, the program executes the statements that make up the work done in the `swap()` function on lines 53 through 55, which uses the technique discussed earlier to rearrange the two values in the array. Next, line 56 assigns `true` to the variable named `didSwap`. The last task performed by the inner loop (line 58) is adding 1 to the value of the subscript variable `x`. This ensures that the next time through the inner loop, the program will compare the next two adjacent items in the array. The program continues to compare two adjacent items and possibly swap them as long as the value of `x` is less than the value of `comparisons`.

The `displayArray()` Function

In the `displayArray()` function, you print the sorted array on the user's screen. Figure 8-1 shows the pseudocode for this function. The C++ code is shown in Figure 8-2.

The work done in the `displayArray()` function begins on line 65 of Figure 8-2. Line 65 assigns the value 0 to the subscript variable, `x`. This is done before the `while` loop is entered because the first item stored in the array is referenced using the subscript value 0. The loop in

lines 66 through 70 prints all of the values in the array named `scores` by incrementing the value of the subscript variable, `x`, each time the loop body executes. When the loop exits, the statement `return 0;` (line 71) executes and ends the program.

134

Exercise 8-2: Using a Bubble Sort

In this exercise, you use what you have learned about sorting data using a bubble sort. Study the following C++ code, and then answer Questions 1–4.

```
int numbers[] = {432, -5, 54, -10, 36, 9, 65, 21, 24};  
const int NUM_ITEMS = 9;  
int j, k, temp;  
int numPasses = 0, numCompares = 0, numSwaps = 0;  
for(j = 0; j < NUM_ITEMS - 1; j++)  
{  
    numPasses++;  
    for(k = 0; k < NUM_ITEMS - 1; k++)  
    {  
        numCompares++;  
        if(numbers[k] > numbers[k+1])  
        {  
            numSwaps++;  
            temp = numbers[k+1];  
            numbers[k+1] = numbers[k];  
            numbers[k] = temp;  
        }  
    }  
}
```

1. Does this code perform an ascending sort or a descending sort? How do you know?

2. Exactly how many passes are made over the data in the array? Specify a number.

3. How many comparisons are made? Specify a number.

4. Do the variables named `numPasses`, `numCompares`, and `numSwaps` accurately keep track of the number of passes, compares, and swaps made in this bubble sort? Explain your answer.

Lab 8-2: Using a Bubble Sort

In this lab, you complete a C++ program that uses an array to store data for the village of Marengo. The program is described in Chapter 8, Exercise 5, in *Programming Logic and Design*. The program should allow the user to enter each household size and determine the mean and median household size in Marengo. The program should output the mean and median household size in Marengo. The file provided for this lab contains the necessary

variable declarations and input statements. You need to write the code that sorts the household sizes in ascending order using a bubble sort and then prints the mean and median household size in Marengo. Comments in the code tell you where to write your statements.

1. Open the source code file named `HouseholdSize.cpp` using Notepad or the text editor of your choice.
2. Write the bubble sort.
3. Output the mean and median household size in Marengo.
4. Save this source code file in a directory of your choice, and then make that directory your working directory.
5. Compile the source code file `HouseholdSize.cpp`.
6. Execute the program with the following input and record the output:
Household sizes: 4, 1, 2, 4, 3, 3, 2, 2, 2, 4, 5, 6

135

Using Multidimensional Arrays

As you learned in Chapter 8 of *Programming Logic and Design*, an array whose elements are accessed using a single subscript is called a **one-dimensional array** or a **single-dimensional array**. You also learned that a **two-dimensional array** stores elements in two dimensions and requires two subscripts to access elements.

In Chapter 8 of *Programming Logic and Design*, you saw how useful two-dimensional arrays can be when you studied the example of owning an apartment building with five floors with each floor having studio, one-bedroom, and two-bedroom apartments. The rent charged for these apartments depends on which floor the apartment is located as well as the number of bedrooms the apartment has. Table 8-1 shows the rental amounts.

Floor	Studio apartment	1-bedroom apartment	2-bedroom apartment
0	350	390	435
1	400	440	480
2	475	530	575
3	600	650	700
4	1000	1075	1150

Table 8-1 Rent schedule based on floor and number of bedrooms

In C++, declaring a two-dimensional array to store the rents shown in Table 8-1 requires two sets of square brackets. The first set of square brackets holds the number of rows in the array, and the second set of square brackets holds the number of columns. The declaration is shown below:

```
const int FLOORS = 5;
const int BEDROOMS = 3;
double rent[FLOORS][BEDROOMS];
```

The declaration shows the array's name, `rent`, followed by two sets of square brackets. The number of rows is included in the first set of square brackets using the constant value of `FLOORS(5)`, and the number of columns is included in the second set of square brackets using the constant value of `BEDROOMS(3)`.

136

As shown below, you can also initialize a two-dimensional array when you declare it by enclosing all of the values within a pair of curly braces and also enclosing the values (separated by commas) for each row within curly braces. Also, notice that each group of values within curly braces is separated by commas.

```
double rent[][] = {{350, 390, 435},  
                   {400, 440, 480},  
                   {475, 530, 575},  
                   {600, 650, 700},  
                   {1000, 1075, 1150}};
```

To access individual elements in the `rent` array, two subscripts are required as shown below.

```
double myRent;  
myRent = rent[3][1];
```

The first subscript (3) determines the row, and the second subscript (1) determines the column. In the assignment statement, `myRent = rent[3][1]`, the value 650 is assigned to the variable named `myRent`.



Remember that in C++, array subscripts begin with 0.

Figure 8-3 shows the pseudocode for a program that continuously displays rents for apartments based on renter requests for bedrooms and floor, and Figure 8-4 shows the C++ code that implements the program.

```
start  
  Declarations  
    num RENTS_BY_FLR_AND_BDRMS[5][3] = {350, 390, 435},  
                                            {400, 440, 480},  
                                            {475, 530, 575},  
                                            {600, 650, 700},  
                                            {1000, 1075, 1150}  
    num floor  
    num bedrooms  
    num QUIT = 99  
  getReady()  
  while floor <> QUIT  
    determineRent()  
  endwhile  
  finish()  
stop
```

Figure 8-3 Pseudocode for a program that determines rent (continues)

(continued)

```
getReady()
    output "Enter floor "
    input floor
return

determineRent()
    output "Enter number of bedrooms "
    input bedrooms
    output "Rent is $",
        RENTS_BY_FLR_AND_BDRMS[floor][bedrooms]
    output "Enter floor "
    input floor
return

finish()
    output "End of program"
return
```

Figure 8-3 Pseudocode for a program that determines rent

```
#include <iostream>
using namespace std;

int main()
{
    // Declare variables.
    const int FLOORS = 5;
    const int BEDROOMS = 3;
    double rent[FLOORS][BEDROOMS] = {{350, 390, 435},
                                      {400, 440, 480},
                                      {475, 530, 575},
                                      {600, 650, 700},
                                      {1000, 1075, 1150}};

    int floor;
    int bedroom;
    int QUIT = 99;

    // Work done in the getReady() function
    cout << "Enter floor or 99 to quit: ";
    cin >> floor;
    while(floor != QUIT)
    {
        // Work done in the determineRent() function
        cout << "Enter number of bedrooms: ";
        cin >> bedroom;
        cout << "Rent is $" << rent[floor][bedroom] << endl;
        cout << "Enter floor or 99 to quit: ";
        cin >> floor;
    }
    // Work done in the finish() function
    cout << "End of program" << endl;
    return 0;
} // End of main() function
```

Figure 8-4 C++ code for a program that determines rent

Exercise 8-3: Using Multidimensional Arrays

In this exercise, you use what you have learned about using multidimensional arrays to answer Questions 1–3.

1. A two-dimensional array declared as `int myNums[6][2];` has how many rows?

2. A two-dimensional array declared as `int myNums[6][2];` has how many columns?

3. Consider the following array declaration, `int myNums[6][2];`
Are the following C++ statements legal?

```
number = myNums[5][3]; _____
number = myNums[0][1]; _____
number = myNums[1][2]; _____
```

Lab 8-3: Using Multidimensional Arrays

139

In this lab, you will complete a C++ program that uses a two-dimensional array to store data for the Building Block Day Care Center. The program is described in Chapter 8, Exercise 10, in *Programming Logic and Design*. The day care center charges varying weekly rates depending on the age of the child and the number of days per week the child attends. The weekly rates are shown in Table 8-2.

Age in Years	Days Per Week				
	1	2	3	4	5
0	30.00	60.00	88.00	115.00	140.00
1	26.00	52.00	70.00	96.00	120.00
2	24.00	46.00	67.00	89.00	110.00
3	22.00	40.00	60.00	75.00	88.00
4 or more	20.00	35.00	50.00	66.00	84.00

Table 8-2 Weekly rates for Lab 8-3

The program should allow the user to enter the age of the child and the number of days per week the child will be at the day care center. The program should output the appropriate weekly rate. The file provided for this lab contains all of the necessary variable declarations, except the two-dimensional array. You need to write the input statements and the code that initializes the two-dimensional array, determines the weekly rate, and prints the weekly rate. Comments in the code tell you where to write your statements.

1. Open the source code file named `DayCare.cpp` using Notepad or the text editor of your choice.
2. Declare and initialize the two-dimensional array.
3. Write the C++ statements that retrieve the age of the child and the number of days the child will be at the day care center.
4. Determine and print the weekly rate.
5. Save this source code file in a directory of your choice, and then make that directory your working directory.
6. Compile the source code file `DayCare.cpp`.
7. Execute the program.

Advanced Modularization Techniques

After studying this chapter, you will be able to:

- ◎ Write functions that require no parameters
- ◎ Write functions that require a single parameter
- ◎ Write functions that require multiple parameters
- ◎ Write functions that return values
- ◎ Pass entire arrays and single elements of an array to a function
- ◎ Pass arguments to functions by reference and by address
- ◎ Overload functions
- ◎ Use C++ built-in functions

In Chapter 2 of this book, you learned that local variables are variables that are declared within the function that uses them. You also learned that most programs consist of a main function that contains the mainline logic and then calls other functions to get specific work done in the program.

142

In this chapter, you learn more about functions in C++. You learn how to write functions that require no parameters, how to write functions that require a single parameter, how to write functions that require multiple parameters, and how to write functions that return a value. You also learn how to pass an array to a function, how to pass arguments by reference and by address, how to overload a function, and how to use some of the built-in functions found in C++. To help you learn about functions, you will study some C++ programs that implement the logic and design presented in *Programming Logic and Design, Eighth Edition*, by Joyce Farrell.

You should do the exercises and labs in this chapter after you have finished Chapter 9 of *Programming Logic and Design*.

Writing Functions with No Parameters

To begin learning about functions, review the C++ code for the Customer Bill program shown in Figure 9-1. Notice the line numbers in front of each line of code in this program. These line numbers are not actually part of the program but are included for reference only.

```
1  /* Program Name: CustomerBill.cpp
2   Function: This program uses a function to print a company name and address
3   and then prints a customer name and balance.
4   Input: Interactive
5   Output: Company name and address, customer name and balance
6 */
7 #include <iostream>
8 #include <string>
9 void nameAndAddress(); // function declaration
10 using namespace std;
11 int main()
12 {
13     // Declare variables local to main()
14     string firstName;
15     string lastName;
16     double balance;
17
18     // Get interactive input
19     cout << "Enter customer's first name: ";
20     cin >> firstName;
21     cout << "Enter customer's last name: ";
22     cin >> lastName;
23     cout << "Enter customer's balance: ";
24     cin >> balance;
```

Figure 9-1 C++ code for Customer Bill program (continues)

(continued)

```
25      // Call nameAndAddress() function
26      nameAndAddress();
27      // Output customer name and address
28      cout << "Customer Name: " << firstName << " " << lastName << endl;
29      cout << "Customer Balance: " << balance << endl;
30
31      return 0;
32 } // End of main() function
33
34
35 void nameAndAddress()
36 {
37     // Declare and initialize local, constant Strings
38     const string ADDRESS_LINE1 = "ABC Manufacturing";
39     const string ADDRESS_LINE2 = "47 Industrial Lane";
40     const string ADDRESS_LINE3 = "Wild Rose, WI 54984";
41
42     // Output
43     cout << ADDRESS_LINE1 << endl;
44     cout << ADDRESS_LINE2 << endl;
45     cout << ADDRESS_LINE3 << endl;
46 } // End of nameAndAddress() function
```

Figure 9-1 C++ code for Customer Bill program

On lines 1–6, you see a multiline comment that describes the Customer Bill program followed by two `#include` preprocessor directives on lines 7 and 8 that give the program the ability to perform input and output and the ability to use `strings`. On line 9, you see the function declaration for the `nameAndAddress()` function. As with variables in C++, you must declare a function before you can call the function in your program. A **function declaration** (also known as a **function prototype**) should specify the data type of the value the function returns, the name of the function, and the data type of each of its arguments. On line 9 you see that the `nameAndAddress()` function returns nothing (`void`) and expects no arguments.

The program begins execution with the `main()` function, which is shown on line 11. This function contains the declaration of three variables (lines 14, 15, and 16), `firstName`, `lastName`, and `balance`, which are local to the `main()` function. Next, on lines 19 through 24, interactive input statements retrieve values for `firstName`, `lastName`, and `balance`.

The function `nameAndAddress()` is then called on line 27, with no arguments listed within its parentheses. Remember that **arguments**, which are sometimes called **actual parameters**, are data items sent to functions. There are no arguments for the `nameAndAddress()` function because this function requires no data. You will learn about passing arguments to functions later in this chapter. The two statements on lines 29 and 30 in the `main()` function are print statements that output the customer's `firstName`, `lastName`, and `balance`.

Next, on line 35, you see the header for the `nameAndAddress()` function. The **header** begins with the `void` keyword, followed by the function name, which is `nameAndAddress`. As you

learned in Chapter 1 of this book, the `void` keyword indicates that the `nameAndAddress()` function does not return a value. You learn more about functions that return values later in this chapter.

Also, notice there are no formal parameters within the parentheses. Remember that **formal parameters** are the variables in the function header that accept the values from the actual parameters. (You will learn about writing functions that accept parameters in the next section of this chapter.) In the next part of the Customer Bill program, you see three constants that are local to the `nameAndAddress()` function: `ADDRESS_LINE1`, `ADDRESS_LINE2`, and `ADDRESS_LINE3`. These constants are declared and initialized on lines 38, 39, and 40, and then printed on lines 43, 44, and 45.

When the input to this program is Ed Gonzales (name) and 352.39 (balance), the output is shown in Figure 9-2.



```
C:\>CustomerBill
Enter customer's first name: Ed
Enter customer's last name: Gonzales
Enter customer's balance: 352.39
ABC Manufacturing
47 Industrial Lane
Wild Rose, WI 54984
Customer Name: Ed Gonzales
Customer Balance: 352.39
C:\>
```

Figure 9-2 Output from the Customer Bill program

Exercise 9-1: Writing Functions with No Parameters

In this exercise, you use what you have learned about writing functions with no parameters to answer Questions 1–2.

1. Given the following function calls, write the function's header and function declaration:

a. `printMailingLabel();`

b. `displayOrderNumbers();`

c. `displayTVListing();`

2. Given the following function headers, write a function call:

a. `void printCellPhoneNumbers()`

b. `void displayTeamMembers()`

c. `void showOrderInfo()`

Lab 9-1: Writing Functions with No Parameters

In this lab, you complete a partially prewritten C++ program that includes functions with no parameters. The program asks the user if he or she has preregistered for tickets to an art show. If the user has preregistered, the program should call a function named `discount()` that displays the message, "You are preregistered and qualify for a 5 percent discount." If the user has not preregistered, the program should call a function named `noDiscount()` that displays the message, "Sorry, you did not preregister and do not qualify for a 5 percent discount." The source code file provided for this lab includes the necessary variable declarations and the input statement. Comments are included in the file to help you write the remainder of the program.

1. Open the source code file named `ArtShowDiscount.cpp` using Notepad or the text editor of your choice.
2. Write the C++ statements as indicated by the comments.
3. Save this source code file in a directory of your choice, and then make that directory your working directory.
4. Compile the source code file, `ArtShowDiscount.cpp`.
5. Execute the program.

Writing Functions that Require a Single Parameter

As you learned in *Programming Logic and Design*, some functions require data to accomplish their task. You also learned that designing a program that sends data (which can be different each time the program runs) to a function (which does not change) keeps you from having to write multiple functions to handle similar situations. For example, suppose you are writing a program that has to determine if a number is even or odd. It is certainly better to write a single function that receives a number entered by the user in the main program than to write individual functions for every number.

In Figure 9-3, you see the C++ code for a program that includes a function that can determine if a number is odd or even. The line numbers are not actually part of the program but are included for reference only. The program allows the user to enter a number, and then it

passes that number to a function as an argument. After it receives the argument, the function can determine if the number is an even number or an odd number.

146

```

1 // EvenOrOdd.cpp - This program determines if a number input by the user is an
2 // even number or an odd number.
3 // Input: Interactive
4 // Output: The number entered and whether it is even or odd
5
6 #include <iostream>
7 #include <string>
8 void evenOrOdd(int);
9 using namespace std;
10
11 int main()
12 {
13     int number; ←
14     cout << "Enter a number or 999 to quit: ";
15     cin >> number;
16
17     while(number != 999)
18     {
19         evenOrOdd(number); ←
20         cout << "Enter a number or 999 to quit: ";
21         cin >> number;
22     }
23     return 0;
24 } // End of main function
25
26 void evenOrOdd(int number) ←
27 {
28     if((number % 2) == 0)
29         cout << "Number: " << number << " is even." << endl;
30     else
31         cout << "Number: " << number << " is odd." << endl;
32 } // End of evenOrOdd function

```

The variable named `number` is local to the main function. Its value is stored at one memory location. For example, it may be stored at memory location 2000.

The value of the formal parameter, `number`, is stored at a different memory location and is local to the `evenOrOdd` function. For example, it may be stored at memory location 7800.

Figure 9-3 C++ code for Even or Odd program

On line 14 in this program, the user is asked to enter a number or the sentinel value, 999, when he or she is finished entering numbers and wants to quit the program. (You learned about sentinel values in Chapter 5 of this book.) On line 15, the input value is retrieved and then stored in the variable named `number`. Next, if the user did not enter the sentinel value 999, the `while` loop is entered and the function named `evenOrOdd()` is called (line 19) using the following syntax:

`evenOrOdd(number);`



Notice the function declaration on line 8 specifies that the `evenOrOdd()` function does not return a value (`void`) and expects a single argument of data type `int`.

The `int` variable `number` is placed within the parentheses, which means the value of `number` is passed to the `evenOrOdd()` function. This is referred to as passing an argument by value.

Passing an argument by value means that a copy of the value of the argument is passed to the function. Within the function, the value is stored in the formal parameter at a different memory location, and it is considered local to that function. In this example, as shown on line 26, the value is stored in the formal parameter named `number`.

 It is important that the data types of the formal parameter and the actual parameter are the same.

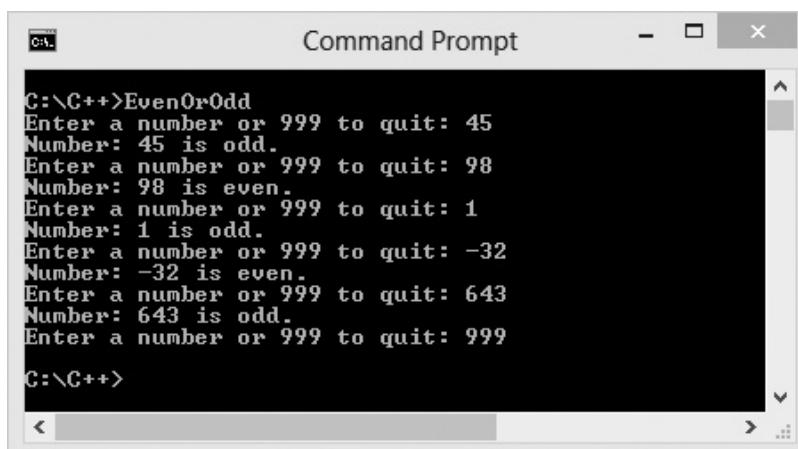
Program control is now transferred to the `evenOrOdd()` function. The header for the `evenOrOdd()` function on line 26 includes the `void` keyword, so you know the function will not return a value. The name of the function follows, and within the parentheses that follow the function name, the parameter `number` is given a local name and declared as the `int` data type.

Remember that even though the parameter `number` has the same name as the local variable `number` in the `main()` function, the values are stored at different memory locations. Figure 9-3 illustrates that the variable `number` that is local to `main()` is stored at one memory location and the parameter `number` in the `evenOrOdd()` function is stored at a different memory location.

Within the function on line 28, the modulus operator (`%`) is used in the test portion of the `if` statement to determine if the value of the local `number` is even or odd. The user is then informed if `number` is even (line 29) or odd (line 31), and program control is transferred back to the statement that follows the call to `evenOrOdd()` in the `main()` function (line 20).

Back in the `main()` function, the user is asked to enter another number on line 20, and the `while` loop continues to execute, calling the `evenOrOdd()` function with a new input value. The loop is exited when the user enters the sentinel value 999 and the program ends. When the input to this program is 45, 98, 1, -32, 643, and 999, the output is shown in Figure 9-4.

In the next section, you will learn how to pass more than one value to a function.



The screenshot shows a Windows Command Prompt window titled "Command Prompt". The window displays the output of a C++ program named "EvenOrOdd". The user inputs several numbers, and the program outputs whether each number is even or odd. The output is as follows:

```
C:\>EvenOrOdd
Enter a number or 999 to quit: 45
Number: 45 is odd.
Enter a number or 999 to quit: 98
Number: 98 is even.
Enter a number or 999 to quit: 1
Number: 1 is odd.
Enter a number or 999 to quit: -32
Number: -32 is even.
Enter a number or 999 to quit: 643
Number: 643 is odd.
Enter a number or 999 to quit: 999
```

Figure 9-4 Output from the Even or Odd program

Exercise 9-2: Writing Functions that Require a Single Parameter

In this exercise, you use what you have learned about writing functions that require a single parameter to answer Questions 1–2.

148

- Given the following variable declarations and function calls, write the function's header and function declaration:

a. `string name;`
`printNameBadge(name);`

b. `double side_length;`
`calculateSquareArea(side_length);`

c. `int hours;`
`displaySecondsInMinutes(minutes);`

- Given the following function headers and variable declarations, write a function call:

a. `string petName = "Mindre";`
`void displayPetName(string petName)`

b. `int currentMonth;`
`void printBirthdays(int month)`

c. `string password;`
`void checkValidPassword(string id)`

Lab 9-2: Writing Functions that Require a Single Parameter

In this lab, you complete a partially written C++ program that includes two functions that require a single parameter. The program continuously prompts the user for an integer until the user enters 0. The program then passes the value to a function that computes the sum of all the whole numbers from 1 up to and including the entered number. Next, the program passes the value to another function that computes the product of all the whole numbers up to and including the entered number. The source code file provided for this lab includes the

necessary variable declarations and the input statement. Comments are included in the file to help you write the remainder of the program.

1. Open the source code file named `SumAndProduct.cpp` using Notepad or the text editor of your choice.
2. Write the C++ statements as indicated by the comments.
3. Save this source code file in a directory of your choice, and then make that directory your working directory.
4. Compile the source code file `SumAndProduct.cpp`.
5. Execute the program.

149

Writing Functions that Require Multiple Parameters

In Chapter 9 of *Programming Logic and Design*, you learned that a function often requires more than one parameter in order to accomplish its task. To specify that a function requires multiple parameters, you include a list of data types and local identifiers separated by commas as part of the function's header. To call a function that expects multiple parameters, you list the actual parameters (separated by commas) in the call to the function.

In Figure 9-5, you see the C++ code for a program that includes a function named `computeTax()` that you designed in *Programming Logic and Design*. The line numbers are not actually part of the program but are included for reference only.

```
1 // ComputeTax.cpp - This program computes tax given a balance
2 // and a rate
3 // Input: Interactive
4 // Output: The balance, tax rate, and computed tax
5
6 #include <iostream>
7 #include <string>
8 void computeTax(double, double);
9 using namespace std;
10
```

Figure 9-5 C++ code for the Compute Tax program (continues)

(continued)

150

```
11 int main()
12 {
13     double balance; --> Memory address 1000
14     double rate; --> Memory address 1008
15
16     cout << "Enter balance: ";
17     cin >> balance;
18     cout << "Enter rate: ";
19     cin >> rate;
20
21     computeTax(balance, rate);
22
23     return 0;
24 } // End of main() function
25
26 void computeTax(double amount, double rate) --> Memory address 9000
27 {
28     double tax; --> Memory address 9008
29
30     tax = amount * rate;
31     cout << "Amount: " << amount << " Rate: " << rate << " Tax: "
32         << tax << endl;
33 } // End of computeTax function
```

Figure 9-5 C++ code for the Compute Tax program



In C++, when you write a function that expects more than one parameter, you must list the parameters separately, even if they have the same data type.



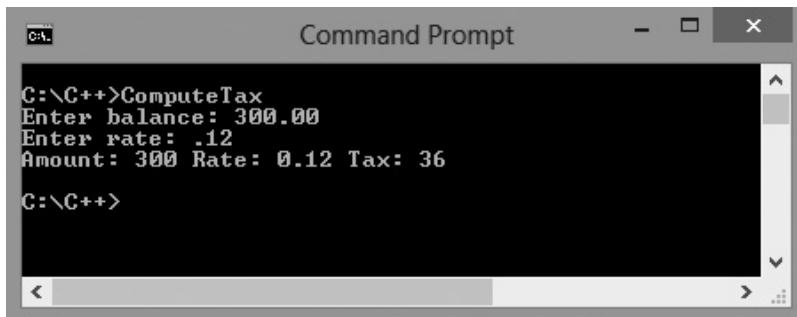
There is no limit to the number of arguments you can pass to a function, but when multiple arguments are passed to a function, the call to the function and the function's header must match. This means that the number of arguments/parameters must be the same, their data types must be the same, and the order in which they are listed must be the same.



Notice the function declaration on line 8 specifies two parameters (arguments) that are both data type `double`. This agrees with the two `doubles` passed to `computeTax()` on line 21 and received as two `doubles` in the function header on line 26.

In the C++ code shown in Figure 9-5, you see that the highlighted call to `computeTax()` on line 21 includes the names of the local variables `balance` and `rate` within the parentheses and that they are separated by commas. These are the arguments (actual parameters) that are passed to the `computeTax()` function. You can also see that the `computeTax()` function header on line 26 is highlighted and includes two formal parameters, `double amount` and `double rate`, listed within parentheses and separated by commas. The value of the variable

named `balance` is passed by value to the `computeTax()` function as an actual parameter and is stored in the formal parameter named `amount`. The value of the variable named `rate` is passed by value to the `computeTax()` function as an actual parameter and is stored in the formal parameter named `rate`. As illustrated in Figure 9-5, it does not matter that one of the parameters being passed, `rate`, has the same name as the parameter received, `rate`, because they occupy different memory locations. When the input to this program is 300.00 (`balance`) and .12 (`rate`), the output is shown in Figure 9-6.



```
C:\>ComputeTax
Enter balance: 300.00
Enter rate: .12
Amount: 300 Rate: 0.12 Tax: 36
C:\>
```

Figure 9-6 Output from the Compute Tax program

Exercise 9-3: Writing Functions that Require Multiple Parameters

In this exercise, you use what you have learned about writing functions that require multiple parameters to answer Questions 1–2.

1. Given the following function calls and variable declarations, write the function's header and function declaration:
 - a. `string name, address;`
`printLabel(name, address);`

 - b. `double side1, side2;`
`calculateRectangleArea(side1, side2);`

 - c. `int day, month, year;`
`birthdayInvitation(day, month, year);`

2. Given the following function headers and variable declarations, write a function call:
 - a. `string customerName = "Smith";`
`double balance = 54000;`
`void printBill(string name, double balance)`

```
b. int val1 = 10, val2 = 20;  
void findProduct(int num1, int num2)
```

```
c. double balance = 37500, interest = .10;  
void newBalance(double bal, double pcnt)
```

152

Lab 9-3: Writing Functions that Require Multiple Parameters

In this lab, you complete a partially written C++ program that includes a function requiring multiple parameters (arguments). The program prompts the user for two numeric values. Both values should be passed to functions named `sum()`, `difference()`, and `product()`. The functions compute the sum of the two values, the difference between the two values, and the product of the two values. Each function should perform the appropriate computation and display the results. The source code file provided for this lab includes the variable declarations and the input statements. Comments are included in the file to help you write the remainder of the program.

1. Open the source code file named `Computation.cpp` using Notepad or the text editor of your choice.
2. Write the C++ statements as indicated by the comments.
3. Save this source code file in a directory of your choice, and then make that directory your working directory.
4. Compile the source code file `Computation.cpp`.
5. Execute the program.

Writing Functions that Return a Value

Thus far in this book, none of the functions you have written (except for the `main()` function) return a value. The header for each of these functions includes the keyword `void`, as in `void nameAndAddress()` indicating that the function does not return a value. However, as a programmer, you will often find that you need to write functions that do return a value. In C++, a function can only return a single value; when you write the code for the function, you must indicate the data type of the value you want returned. This is often referred to as the function's **return type**. The return type can be any of the built-in data types found in C++, as well as a class type, such as `string`. You will learn more about classes in Chapter 10 of this book. For now, you will focus on returning values of the built-in types and `string` objects.

In Chapter 9 of *Programming Logic and Design*, you studied the design for a program that includes a function named `getHoursWorked()`. This function is designed to prompt a user for the number of hours an employee has worked, retrieve the value, and then return that value

```
1 // GrossPay.cpp - This program computes an employee's gross pay.
2 // Input: Interactive
3 // Output: The employee's hours worked and their gross pay
4
5 #include <iostream>
6 #include <string>
7 using namespace std;
8 double getHoursWorked();
9
10 int main()
11 {
12     double hours;
13     const double PAY_RATE = 12.00;
14     double gross;
15
16     hours = getHoursWorked();
17     gross = hours * PAY_RATE;
18
19     cout << "Hours worked: " << hours << endl;
20     cout << "Gross pay is: " << gross << endl;
21
22     return 0;
23 } // End of main() function
24
25 double getHoursWorked()
26 {
27     double workHours;
28
29     cout << "Please enter hours worked: ";
30     cin >> workHours;
31
32     return workHours;
33 } // End of getHoursWorked function
```

Figure 9-7 C++ code for a program that includes the `getHoursWorked()` function

The C++ program shown in Figure 9-7 declares local constants and variables `hours`, `PAY_RATE`, and `gross` on lines 12, 13, and 14 in the `main()` function. The next statement (line 16), shown below, is an assignment statement:

```
hours = getHoursWorked();
```

This assignment statement includes a call to the function named `getHoursWorked()`. As with all assignment statements, the expression on the right side of the assignment operator (`=`) is evaluated, and then the result is assigned to the variable named on the left side of the

assignment operator (`=`). In this example, the expression on the right is a call to the `getHoursWorked()` function.

When the `getHoursWorked()` function is called, program control is transferred to the function. Notice that the header (line 25) for this function is written as follows:

154

```
double getHoursWorked()
```

The keyword `double` is used in the header to specify that a value of data type `double` is returned by this function. Also, notice on line 8 in the program, the function declaration agrees with the function header. That is, the function declaration specifies that the function returns a `double`.

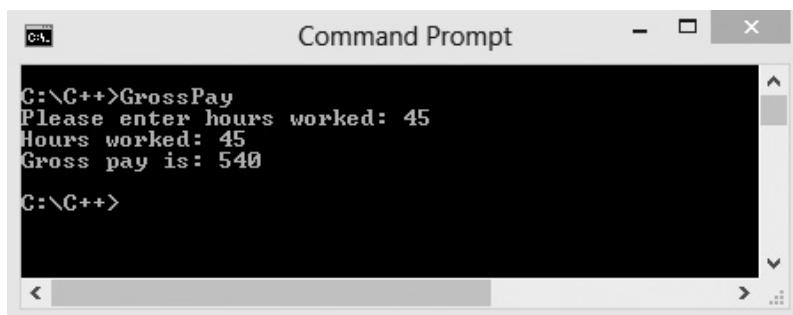
A local, `double` variable named `workHours` is then declared on line 27. On line 29, the user is asked to enter the number of hours worked, and on line 30 the value is retrieved and stored in the local variable named `workHours`. The `return` statement that follows on line 32 returns a copy of the value stored in `workHours` (data type `double`) to the location in the calling function where `getHoursWorked()` is called, which is the right side of the assignment statement on line 16.

The value returned to the right side of the assignment statement is then assigned to the variable named `hours` (data type `double`) in the `main()` function. Next, the gross pay is calculated on line 17, followed by the `cout` statements on lines 19 and 20 that display the value of the local variables, `hours` and `gross`, which contain the number of hours worked and the calculated gross pay.

You can also use a function's return value directly rather than store it in a variable. The two C++ statements that follow make calls to the same `getHoursWorked()` function shown in Figure 9-7, but in these statements, the returned value is used directly in the statement that calculates gross pay and in the statement that prints the returned value:

```
gross = getHoursWorked() * PAY_RATE;  
cout << "Hours worked are " << getHoursWorked() << endl;
```

When the input to this program is 45, the output is shown in Figure 9-8.



```
C:\>GrossPay  
Please enter hours worked: 45  
Hours worked: 45  
Gross pay is: 540  
C:\>
```

Figure 9-8 Output from a program that includes the `getHoursWorked()` function

Exercise 9-4: Writing Functions that Return a Value

In this exercise, you use what you have learned about writing functions that return a value to answer Questions 1–2.

- Given the following variable declarations and function calls, write the function's header:

a. double price, percent, newPrice;
newPrice = calculateNewPrice(price, percent);

b. double area, one_length, two_length;
area = calcArea(one_length, two_length);

c. string lowerCase, upperCase;
upperCase = changeCase(lowerCase);

- Given the following function headers and variable declarations, write a function call:

a. int itemID = 1234;
string itemName;
string findItem(int itemNumber)

b. int val, cube;
int cubed(int num1)

c. int number = 3, exponent = 4, result;
int power(int num, int exp)

Lab 9-4: Writing Functions that Return a Value

In this lab, you complete a partially written C++ program that includes a function that returns a value. The program is a simple calculator that prompts the user for two numbers and an operator (+, -, *, or /). The two numbers and the operator are passed to the function where the appropriate arithmetic operation is performed. The result is then returned to the `main()` function where the arithmetic operation and result are displayed. For example, if the user enters 3, 4, and *, the following is displayed:

3 * 4 = 12

The source code file provided for this lab includes the necessary variable declarations and input and output statements. Comments are included in the file to help you write the remainder of the program.

1. Open the source code file named `Calculator.cpp` using Notepad or the text editor of your choice.
2. Write the C++ statements as indicated by the comments.
3. Save this source code file in a directory of your choice, and then make that directory your working directory.
4. Compile the source code file `Calculator.cpp`.
5. Execute the program.

Passing an Array and an Array Element to a Function

As a C++ programmer, there are times when you will want to write a function to perform a task on all of the elements you have stored in an array. For example, in Chapter 9 of *Programming Logic and Design*, you saw a design for a program that used a function to quadruple all of the values stored in an array. This design is translated into C++ code in Figure 9-9.

```
1 // PassEntireArray.cpp - This program quadruples the values stored in an array.
2 // Input: None
3 // Output: The original values and the quadrupled values
4
5 #include <iostream>
6 #include <string>
7 using namespace std;
8
9 void quadrupleTheValues(int[]);
10 int main()
11 {
12     const int LENGTH = 4;
13     int someNums[] = {10, 12, 22, 35};
14     int x;
15
16     cout << "At beginning of main function. . . " << endl;
17     x = 0;
18     while (x < LENGTH)
19     {
20         cout << someNums[x] << endl;
21         x++;
22     }
```

Figure 9-9 C++ code for Pass Entire Array program (*continues*)

(continued)

```
23     quadrupleTheValues(someNums);
24     cout << "At the end of main function. . . " << endl;
25     x = 0;
26     while (x < LENGTH)
27     {
28         cout << someNums[x] << endl;
29         x++;
30     }
31     return 0;
32 } // End of main() function
33
34 void quadrupleTheValues(int vals[])
35 {
36     const int LENGTH = 4;
37     int x;
38
39     x = 0;
40     while(x < LENGTH)
41     {
42         cout << " In quadrupleTheValues() function, value is " << vals[x] << endl;
43         x++;
44     }
45     x = 0;
46     while(x < LENGTH)
47     {
48         vals[x] = vals[x] * 4;
49         x++;
50     }
51     x = 0;
52     while(x < LENGTH)
53     {
54         cout << " After change, value is " << vals[x] << endl;
55         x++;
56     }
57     return;
58 } // End of quadrupleTheValues function
```

Figure 9-9 C++ code for Pass Entire Array program

The `main()` function begins on line 10 and proceeds with the declaration and initialization of the constant named `LENGTH` (line 12) and the array of `ints` named `someNums` (line 13), followed by the declaration of the variable named `x` (line 14), which is used as a loop control variable. The first `while` loop in the program on lines 18 through 22 is responsible for printing the values stored in the array at the beginning of the program. Next, on line 23, the function named `quadrupleTheValues()` is called. The array named `someNums` is passed as an argument. Notice that when an entire array is passed to a function, the square brackets and the size are not included. Also note that when you pass an entire array to a function, the beginning memory address of the array is passed by value. This means instead of a copy of the array being passed, the memory address of the array is passed. Although you cannot change

the beginning memory address of the array, this does give the function access to that memory location; the function can then change the values stored in the array if necessary.

Program control is then transferred to the `quadrupleTheValues()` function. The header for the function on line 34 includes one parameter, `int vals[]`. The syntax for declaring an array as a formal parameter includes the parameter's data type, followed by a local name for the array, followed by empty square brackets. Note that a size is not included within the square brackets.



On line 9, you see the function declaration for the `quadrupleTheValues()` function. The function declaration includes the return type of the function (`void`), the name of the function (`quadrupleTheValues()`), the data type of the parameter(s) (`int`), and empty square brackets (`[]`) that specify the parameter is an array.

In the `quadrupleTheValues()` function, the first `while` loop on lines 40 through 44 prints the values stored in the array, and the second `while` loop on lines 46 through 50 accesses each element in the array, quadruples the value, and then stores the quadrupled values in the array at their same location. The third `while` loop on lines 52 through 56 prints the changed values now stored in the array. Program control is then returned to the location in the `main()` function where the function was called.

When program control returns to the `main()` function, the next statements to execute (lines 24 through 30) are responsible for printing out the values stored in the array once more. The output from this program is displayed in Figure 9-10.

The screenshot shows a Windows Command Prompt window titled "Command Prompt". The window displays the following text:

```
C:\>PassEntireArray
At beginning of main function. . .
10
12
22
35
In quadrupleTheValues() function, value is 10
In quadrupleTheValues() function, value is 12
In quadrupleTheValues() function, value is 22
In quadrupleTheValues() function, value is 35
After change, value is 40
After change, value is 48
After change, value is 88
After change, value is 140
At the end of main function. . .
40
48
88
140
C:\>
```

Figure 9-10 Output from Pass Entire Array program

As shown in Figure 9-10, the array values printed at the beginning of the `main()` function (lines 18 through 22) are the values with which the array was initialized. Next, the

The `quadrupleTheValues()` function prints the array values (lines 40 through 44) again before they are changed. As shown, the values remain the same as the initialized values. The `quadrupleTheValues()` function then prints the array values again after the values are quadrupled (lines 52 through 56). Finally, after the call to `quadrupleTheValues()`, the `main()` function prints the array values one last time (lines 26 through 30). These are the quadrupled values, indicating that the `quadrupleTheValues()` function has access to the memory location where the array is stored and is able to permanently change the values stored there.

You can also pass a single array element to a function, just as you pass a variable or constant. The following C++ code initializes an array named `someNums`, declares a variable named `newNum`, and passes one element of the array to a function named `tripleTheNumber()`. The array element is passed by value in this example, which means a copy of the value stored in the array is passed to the `tripleTheNumber()` function:

```
int someNums[] = {10, 12, 22, 35};  
int newNum;  
newNum = tripleTheNumber(someNums[1]);
```

The following C++ code includes the header for the function named `tripleTheNumber()` along with the code that triples the value passed to it:

```
int tripleTheNumber(int num)  
{  
    int result;  
    result = num * 3;  
    return result;  
}
```

Exercise 9-5: Passing Arrays to Functions

In this exercise, you use what you have learned about passing arrays and array elements to functions to answer Questions 1–3.

- Given the following function calls, write the function's header:

a. `int marchBirthdays [] = {17, 24, 21, 14, 28, 9};
printBirthdays(marchBirthdays);`

b. `double octoberInvoices [] = {100.00, 200.00, 55.00, 230.00};
double total;
total = monthlyIncome(octoberInvoices);`

c. `double balance[] = {34.56, 33.22, 65.77, 89.99};
printBill(balance[1]);`

2. Given the following function headers and variable declarations, write a function call:

a. `string names[] = {"Jones", "Smith", "Brown", "Perez"};`
`double grades[] = {95, 76, 88, 72};`
`void midtermGrades(string names[], double grades[])`

160

b. `int numbers[] = {1, 4, 6, 8, 3, 7};`
`int result;`
`int printAverage(int nums[])`

3. Given the following function header (in which `sal` is one element of an array of `doubles`), write a function call that passes the last value in the array named `salaries`:

a. `double salaries[] = {45000, 23000, 35000};`
`void bonus(double sal)`

Lab 9-5: Passing Arrays to Functions

In this lab, you complete a partially written C++ program that reverses the order of five numbers stored in an array. The program should first print the five numbers stored in the array. Next, the program passes the array to a function where the numbers are reversed. The program then prints the reversed numbers in the main program.

The source code file provided for this lab includes the necessary variable declarations. Comments are included in the file to help you write the remainder of the program.

1. Open the source code file named `Reverse.cpp` using Notepad or the text editor of your choice.
2. Write the C++ statements as indicated by the comments.
3. Save this source code file in a directory of your choice, and then make that directory your working directory.
4. Compile the source code file `Reverse.cpp`.
5. Execute the program.

Passing Arguments by Reference and by Address

As a C++ programmer, you will sometimes want to pass one or more arguments to a function and allow the function to permanently change the value of that argument. In the previous section, you learned when you pass an array to a function, it is actually the beginning memory address of the array that is passed. Because the array's address is passed to a function, the function has access to that memory address and can therefore change the value (or values) that are stored at the memory location.

In C++, there are two techniques you can use to pass the memory address of variables that are not arrays to functions. The two techniques are referred to as **pass by reference** and **pass by address**. You begin with a discussion of pass by reference.

Pass by Reference

Passing an argument to a function by reference allows the called function to access a variable in the calling function. In other words, the called function gains the ability to change the value of a local variable in the calling function. The argument in the called function becomes an **alias** (another name) for the variable in the calling function.

To specify that an argument is passed by reference, you include the ampersand (&) operator in the function declaration and in the function header. You do not use the & operator in the function call. Figure 9-11 shows a C++ program named `IncreaseSalary.cpp` that includes a function named `increase()`. The `increase()` function calculates an employee's new salary based on a current salary and a percentage that represents an employee's raise.

```
1 // IncreaseSalary.cpp - This program demonstrates pass by reference.
2 // Input: None
3 // Output: Beginning salary along with new salary
4 #include <iostream>
5 using namespace std;
6
7 void increase(double, double, double&);
8 int main()
9 {
10    double salary = 50000.00;
11    double raise = .15;
12    double new_salary;
13
14    cout << "Salary is: " << salary << endl;
15
16    // Call the increase function; the variables salary and raise
17    // are passed by value, the variable new_salary is passed by
18    // reference.
19    increase(salary, raise, new_salary);
20
21    cout << "New salary is: " << new_salary << endl;
22    return 0;
23 } // End of main() function
24
25 void increase(double amt, double pcnt, double& new_sal)
26 {
27    // Changes the value of new_salary in the main function
28    new_sal = amt * (1 + pcnt);
29    return;
30 }
```

Figure 9-11 Pass by reference example

On line 7 in Figure 9-11, you see the function declaration for the function named `increase()`. The function declaration includes the data types for three arguments that are passed to the `increase()` function. The first two arguments are `doubles`, and the third argument is a reference to a `double` (`double&`). The function declaration on line 7 has been shaded in Figure 9-11.

162



In Chapter 3 of this book, you learned that local variables are not available to other functions in the program.

On line 19, you see the function call to the `increase()` function in the `main()` function. The `increase()` function is also shaded so you can easily see that it passes three arguments (`salary`, `raise`, and `new_salary`) that are all local variables declared in the `main()` function.

On line 25, also shaded, you see the header for the `increase()` function that includes three parameters, `double amt`, `double pcnt`, and `double& new_sal`. The first two arguments are passed by value, which means that `amt` contains a copy of the value stored in `salary`, and `pcnt` contains a copy of the value stored in `raise`. The third argument is passed by reference, which means that `new_sal` is actually an alias for the variable `new_salary`. Because `new_sal` is an alias for `new_salary`, when the value of `new_sal` is changed on line 28, the value of `new_salary`, which is declared in the `main()` function, is also changed.

You can compile and execute this program to verify that the `increase()` function changes the value of the variable named `new_salary`. The output from this program is shown in Figure 9-12. The program, named `IncreaseSalary.cpp`, is included with the data files provided with this book.

The screenshot shows a command prompt window titled "Developer Command Prompt for VS2012". The output of the program is displayed:

```
C:\>IncreaseSalary
Salary is: 50000
New salary is: 57500
C:\>
```

Figure 9-12 Output from pass by reference example

Pass by Address

The second technique you can use to pass the memory address of variables (that are not arrays) to functions is called *pass by address*. Passing an argument by address allows the called function to change the value of an argument.

Before you can understand passing arguments by address, you must learn about pointers. A **pointer** is a variable that stores a memory address as its value. You are accustomed to using a variable name to directly access a value stored at a memory address associated with a variable.

You can also use a pointer variable and the memory address stored in the pointer variable to indirectly access a value stored at a memory address.

The first thing you need to learn in order to use pointers is how to use the address operator. You use the C++ address operator (&) to obtain the memory address associated with a variable. The following C++ code prints the value of a variable named num and also uses the address operator (&) to print the memory address associated with the variable named num:

```
int num = 5;
cout << "Value of num is: " << num << endl;
cout << "Address of num is: " << &num << endl;
```

As you can see in the preceding, highlighted line of code, the address operator (&) is placed before the name of the variable (num) to obtain the address of num.

Next, you must learn to use the asterisk (*) symbol to declare a pointer variable. When you declare a pointer, a specific data type must be explicitly associated with the pointer. The C++ code that follows declares an int variable named age and ptr_age as a pointer to an int:

```
int age = 21;
int* ptr_age = &age;
```

Figure 9-13 shows memory after the preceding code executes. The variable named age is associated with memory location 2000, where the value 21 is stored. The variable named ptr_age is associated with memory location 2004, where the value 2000 (the address of age) is stored. Because ptr_age contains the address of age, you can say that ptr_age is a pointer to an int and that it points to age.

The (*) symbol also serves as the **indirection operator**, which means that it can be used to indirectly access the value a pointer points to. Because the (*) symbol is used in many ways in C++, the C++ compiler must use context to determine the appropriate meaning. The C++ code that follows shows the use of the (*) symbol to carry out multiplication, to declare a pointer variable, and to function as the indirection operator:

```
int y = 5;
int x;
int* ptr_x; // Pointer variable
x = y * 10; // Multiplication; x is assigned the value 50
ptr_x = &x; // ptr_x is assigned the memory address of x
cout << *ptr_x; // Indirection operator; prints 50
```

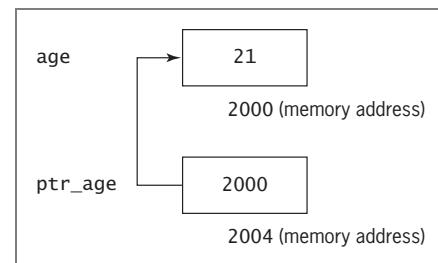


Figure 9-13 Memory contents for age and ptr_age

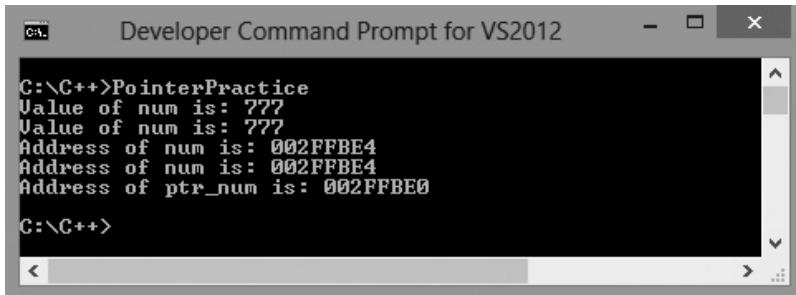
The C++ program shown in Figure 9-14 shows additional examples when the (*) symbol is used to declare a pointer variable and also when it is used as the indirection operator.

```
1 // PointerPractice.cpp - This program demonstrates pointers.
2 // Input: None
3 // Output: Values of variables
4 #include <iostream>
5 using namespace std;
6
7 int main()
8 {
9     int num = 777;
10    int* ptr_num;
11
12    ptr_num = &num; // Assigns address of num to ptr_num
13
14    cout << "Value of num is: " << num << endl;
15
16    cout << "Value of num is: " << *ptr_num << endl;
17
18    cout << "Address of num is: " << &num << endl;
19
20    cout << "Address of num is: " << ptr_num << endl;
21
22    cout << "Address of ptr_num is: " << &ptr_num << endl;
23
24    return 0;
25 } // End of main() function
```

Figure 9-14 C++ program that uses pointer variables and indirection operator

On line 9, the variable `num` is declared as an `int` and initialized with the value `777`, and on line 10, the variable named `ptr_num` is declared as a pointer to an `int`. Line 12 assigns the address of the variable `num` to `ptr_num`, and then on line 14, the `cout` statement prints the value of `num` by using the name of the variable to directly access its value. The `cout` statement on line 16 also prints the value of `num`, but this time the indirection operator is used along with the pointer variable (`*ptr_num`) to indirectly access the value of `num`. Next, the address of `num` is printed on line 18 using the address operator and the name of the variable (`&num`) to obtain the address of the variable. On line 20, the address of `num` is printed again, this time using the value stored in the pointer variable `ptr_num`. The `cout` statement on line 22 prints the address of `ptr_num` using the address operator (`&ptr_num`).

The output of this program is shown in Figure 9-15. Keep in mind that actual memory addresses will vary from one system to another and even from one run of the program to another.

A screenshot of the Developer Command Prompt for VS2012. The window title is "Developer Command Prompt for VS2012". The command line shows the path "C:\C++>PointerPractice". The output window displays the following text:

```
C:\C++>PointerPractice
Value of num is: 777
Value of num is: 777
Address of num is: 002FFBE4
Address of num is: 002FFBE4
Address of ptr_num is: 002FFBE0
```

Figure 9-15 Output of Pointer Practice program

You can compile and execute this program to see the output values on your system. The program file, named `PointerPractice.cpp`, is included with the data files provided with this book.

To pass an argument(s) to a function using pass by address, the memory address of a local variable(s) in the calling function is passed to the called function. You use the (&) operator (address operator) when you call the function. The parameter(s) must be declared as a pointer in the function declaration and in the function's header. You use the (*) symbol to declare the pointer parameter(s). You also must use the * symbol (indirection operator) to dereference the pointer(s) in the body of the function. **Dereferencing** a pointer means that you use the indirection operator to indirectly gain access to the value of a variable.

The C++ program shown in Figure 9-16 illustrates the use of pass by address.

```
1 // PassByAddress.cpp - This program demonstrates pass by address.
2 // Input: None
3 // Output: Values of variables
4 #include <iostream>
5 using namespace std;
6
7 void swap(int*, int*);
8
9 int main()
10 {
11     int num1 = 5;
12     int num2 = 10;
13
14     cout << "Value of num1 is: " << num1 << endl;
15
16     cout << "Value of num2 is: " << num2 << endl;
17
18     swap(&num1, &num2);
19
20     cout << "Value of num1 is: " << num1 << endl;
21
22     cout << "Value of num2 is: " << num2 << endl;
23 }
```

Figure 9-16 C++ program that uses pass by address (continues)

(continued)

166

```
24     return 0;
25 } // End of main() function
26
27 void swap(int* val1, int* val2)
28 {
29     int temp;
30     temp = *val1;
31     *val1 = *val2;
32     *val2 = temp;
33 }
```

Figure 9-16 C++ program that uses pass by address

The function declaration for the `swap()` function is shown on line 7. Notice the `swap()` function does not return a value (`void`); instead, it specifies two parameters, both pointers to `ints` (`int*`, `int*`).

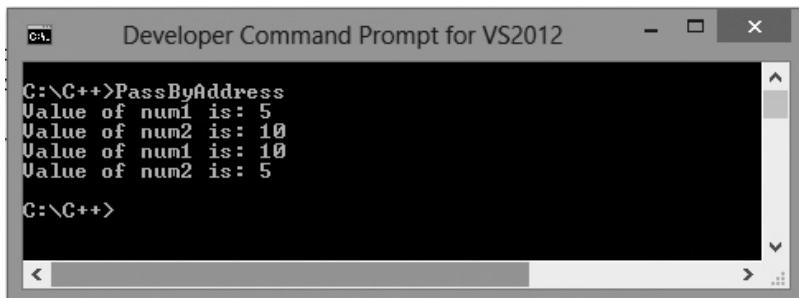
In the `main()` function on lines 11 and 12, two `int` variables, `num1` and `num2`, are declared and initialized to the values 5 and 10, respectively. Both of these variables, `num1` and `num2`, are local to the `main()` function, which means that the `swap()` function does not have direct access to them. On lines 14 and 16, the values of `num1` (5) and `num2` (10) are printed.

Next, on line 18, the `swap()` function is called, passing the address of `num1` (`&num1`) and the address of `num2` (`&num2`) to the function. Because the memory addresses of `num1` and `num2` are passed, the `swap()` function has access to the values stored at those memory addresses and will be able to change the values stored there.

Program control is now passed to the `swap()` function. The header for the `swap()` function on line 27 specifies that the function does not return a value and accepts two parameters, `val1` and `val2`. The two parameters are declared as pointers to `ints` (`int* val1, int* val2`). The parameter `val1` contains the address of `num1`, and the parameter `val2` contains the address of `num2`. On line 29, the `int` variable named `temp` is declared. Because `val1` and `val2` are both pointers, the indirection operator (`*`) is used on lines 30, 31, and 32 to gain access to the values stored in memory that are pointed to by `val1` and `val2`. On line 30, the value pointed to by `val1` (the value of `num1`) is accessed indirectly and then assigned to `temp`. On line 31, the value pointed to by `val2` (the value of `num2`) is accessed indirectly then assigned to what `val1` points to, which is `num1`. This assignment statement changes the value of `num1` in the `main()` function. Line 32 assigns the value stored in `temp` to what `val2` points to, which is `num2`. This assignment changes the value of `num2` in the `main()` function.

When the `swap()` function returns control to the `main()` function, the values of `num1` and `num2` are displayed again on lines 20 and 22. This time, the value of `num1` is 10 and the value of `num2` is 5, illustrating that the local variables `num1` and `num2` have been changed by the `swap()` function.

The output from this program is shown in Figure 9-17. You can compile and execute this program to see the output values on your system. The program, named `PassByAddress.cpp`, is included with the data files provided with this book.



```
C:\>PassByAddress
Value of num1 is: 5
Value of num2 is: 10
Value of num1 is: 10
Value of num2 is: 5
C:\>
```

Figure 9-17 Output from C++ program that uses pass by address

Exercise 9-6: Pass by Reference and Pass by Address

In this exercise, you use what you have learned about passing arguments by reference and by address to functions to answer Questions 1–2.

- Given the following variable and function declarations, write the function call and the function's header:

a. `double price = 22.95, increase = .10;`
`void changePrice(double&, double);`

b. `double price = 22.95, increase = .10;`
`void changePrice(double* , double);`

c. `int age = 23;`
`void changeAge(int&);`

d. `int age = 23;`
`void changeAge(int*);`

2. Given the following function headers and variable declarations, write a function call:

a. `custNames[] = {"Perez", "Smith", "Patel", "Shaw"};`
`balances[] = {34.00, 21.00, 45.50, 67.00};`
`void cust(string name[], double bal[])`

b. `int values[] = {1, 77, 89, 321, -2, 34};`
`void printSum(int nums[])`

168

Lab 9-6: Pass by Reference and Pass by Address

In this lab, you complete a partially written C++ program that includes a function named `multiplyNumbers()` that multiplies two `int` values to find their product. Three `ints` should be passed to the `multiplyNumbers()` function, the two numbers to be multiplied (`num1` and `num2`) should be passed by value, and another `int` (`product`) to hold the product of the two numbers should be passed by reference, enabling the `multiplyNumbers()` function to change its value.

The source code file provided for this lab includes the necessary variable declarations and input and output statements. Comments are included in the file to help you write the remainder of the program.

1. Open the source code file named `MultiplyTwo.cpp` using Notepad or the text editor of your choice.
2. Write the `multiplyNumbers()` function, the function declaration, and the function call as indicated by the comments.
3. Save this source code file in a directory of your choice, and then make that directory your working directory.
4. Compile the source code file `MultiplyTwo.cpp`.
5. Execute the program.
6. Rewrite the `multiplyNumbers()` function to pass the two numbers (`num1` and `num2`) by value and to pass `product` by address.
7. Save this program as `MultiplyTwo2.cpp`.
8. Compile `MultiplyTwo2.cpp`.
9. Execute the program. Both programs (`MultiplyTwo.cpp` and `MultiplyTwo2.cpp`) should generate the same output.

Overloading Functions

You can **overload** functions by giving the same name to more than one function. Overloading functions is useful when you need to perform the same action on different types of inputs. For example, you may want to write multiple versions of an `add()` function—one that can add two integers, another that can add two doubles, another that can add three integers, and

another that can add two integers and a double. Overloaded functions have the same name, but they must either have a different number of arguments or the arguments must be of a different data type. C++ figures out which function to call based on the function's name and its arguments, the combination of which is known as the function's **signature**. The signature of an overloaded function consists of the function's name and its argument list; it does not include the function's return type.

Overloading functions allows a C++ programmer to choose meaningful names for functions, and it also permits the use of polymorphic code. **Polymorphic** code is code that acts appropriately depending on the context. (The word *polymorphic* is derived from the Greek words *poly*, meaning “many,” and *morph*, meaning “form.”) Polymorphic functions in C++ can take many forms. You will learn more about polymorphism in other C++ courses, when you learn more about object-oriented programming. For now, you can use overloading to write functions that perform the same task but with different data types.

In Chapter 9 of *Programming Logic and Design*, you studied the design for an overloaded function named `printBill()`. One version of the function includes a numeric parameter, a second version includes two numeric parameters, a third version includes a numeric parameter and a `string` parameter, and a fourth version includes two numeric parameters and a `string` parameter. All versions of the `printBill()` function have the same name with a different signature; therefore, it is an overloaded function. In Figure 9-18, you see a C++ program that includes the four versions of the `printBill()` function.

```
1 // Overloaded.cpp - This program illustrates overloaded functions.
2 // Input:  None
3 // Output: Bill printed in various ways
4
5 #include <iostream>
6 #include <string>
7 using namespace std;
8
9 void printBill(double);
10 void printBill(double, double);
11 void printBill(double, string);
12 void printBill(double, double, string);
13
14 int main()
15 {
16     double bal = 250.00, discountRate = .05;
17     string msg = "Due in 10 days.";
18 }
```

Figure 9-18 Program that uses overloaded `printBill()` functions (continues)

(continued)

170

```
19     printBill(bal);                      // Call version #1
20     printBill(bal, discountRate);        // Call version #2
21     printBill(bal, msg);                // Call version #3
22     printBill(bal, discountRate, msg);  // Call version #4
23
24     return 0;
25 } // End of main() function
26
27 // printBill() function #1
28 void printBill(double balance)
29 {
30     cout << "Thank you for your order." << endl;
31     cout << "Please remit " << balance << endl;
32 } // End of printBill #1 function
33
34 // printBill() function #2
35 void printBill(double balance, double discount)
36 {
37     double newBal;
38     newBal = balance - (balance * discount);
39     cout << "Thank you for your order." << endl;
40     cout << "Please remit " << newBal << endl;
41 } // End of printBill #2 function
42
43 // printBill() function #3
44 void printBill(double balance, string message)
45 {
46     cout << "Thank you for your order." << endl;
47     cout << message << endl;
48     cout << "Please remit " << balance << endl;
49 } // End of printBill #3 function
50
51 // printBill() function #4
52 void printBill(double balance, double discount, string message)
53 {
54     double newBal;
55     newBal = balance - (balance * discount);
56     cout << "Thank you for your order." << endl;
57     cout << message << endl;
58     cout << "Please remit " << newBal << endl;
59 } // End of printBill #4 function
```

Figure 9-18 Program that uses overloaded printBill() functions

On line 19, the first call to the `printBill()` function passes one argument, the variable named `bal`, which is declared as a `double`. This causes the runtime system to find and execute the `printBill()` function that is written to accept one `double` as an argument (line 28). The third call to the `printBill()` function (line 21) passes two arguments, a `double` and a `string`. This causes the runtime system to find and execute the `printBill()` function that is written to accept a `double` and a `string` as arguments (line 44). You can compile and execute this program if you would like to verify that a different version of the `printBill()` function is

A screenshot of a Windows command prompt window titled "Developer Command Prompt for VS2012". The window shows the output of a C++ program named "Overloaded". The output consists of several lines of text: "C:\>Overloaded", "Thank you for your order.", "Please remit 250", "Thank you for your order.", "Please remit 237.5", "Thank you for your order.", "Due in 10 days.", "Please remit 250", "Thank you for your order.", "Due in 10 days.", and "Please remit 237.5". The window has a standard title bar with minimize, maximize, and close buttons. The text area is black with white text, and there is a vertical scroll bar on the right side.

```
C:\>Overloaded
Thank you for your order.
Please remit 250
Thank you for your order.
Please remit 237.5
Thank you for your order.
Due in 10 days.
Please remit 250
Thank you for your order.
Due in 10 days.
Please remit 237.5
C:\>
```

Figure 9-19 Output from program that uses overloaded printBill() functions

Exercise 9-7: Overloading Functions

In this exercise, you use what you have learned about overloading functions to answer Question 1.

1. In Figure 9-20, which function header would the following function calls match? Use a line number as your answer.

```
1 // Function headers
2 int sum(int num1, int num2)
3 int sum(int num2, int num2, int num3)
4 double sum(double num1, double num2)
5 // Variable declarations
6 double number1 = 1.0, ans1;
7 int number2 = 5, ans2;
```

Figure 9-20 Function headers and variable declarations

a. `ans2 = sum(2, 7, 9);`

b. `ans1 = sum(number2, number2);`

c. `ans1 = sum(10.0, 7.0);`

d. `ans2 = sum(2, 8, number2);`

e. `ans2 = sum(3, 5);`

172

Lab 9-7: Overloading Functions

In this lab, you complete a partially written C++ program that computes hotel guest rates at Cornwall's Country Inn. The program is described in Chapter 9, Exercise 11, in *Programming Logic and Design*. In this program, you should include two overloaded functions named `computeRate()`. One version accepts a number of days and calculates the rate at \$99.99 per day. The other accepts a number of days and a code for a meal plan. If the code is *A*, three meals per day are included, and the price is \$169.00 per day. If the code is *C*, breakfast is included, and the price is \$112.00 per day. Each function returns the rate to the calling program where it is displayed. The main program asks the user for the number of days in a stay and whether meals should be included; then, based on the user's response, the program either calls the first function or prompts for a meal plan code and calls the second function. Comments are included in the file to help you write the remainder of the program.

1. Open the source code file named `Cornwall.cpp` using Notepad or the text editor of your choice.
2. Write the C++ statements as indicated by the comments.
3. Save this source code file in a directory of your choice, and then make that directory your working directory.
4. Compile the source code file `Cornwall.cpp`.
5. Execute the program.

Using C++ Built-in Functions

Throughout this book, you have been told that you do not know enough about the C++ programming language to be able to control the number of places displayed after the decimal point when you print a value of data type `double`. There are actually several ways to control the number of places displayed after a decimal point. In this section, you look at one approach: using a manipulator named `setprecision()`. A **manipulator** is a special C++ function that can control (manipulate) how data is displayed. Throughout this book, you have been using the `endl` manipulator to display a new line character.

In the C++ program that follows, you see how the `setprecision()` manipulator is used to control the number of digits displayed.

```
#include <iostream>
#include <iomanip>
using namespace std;
int main()
```

```
{  
    double value = 3.14159;  
    cout << value << endl;  
    cout << setprecision(1) << value << endl;  
    cout << setprecision(2) << value << endl;  
    cout << setprecision(3) << value << endl;  
    cout << setprecision(4) << value << endl;  
    cout << setprecision(5) << value << endl;  
    cout << setprecision(6) << value << endl;  
    return 0;  
}
```

In the preceding code sample, `setprecision()` is a manipulator that controls how the floating point value should be formatted. The `setprecision()` manipulator requires one integer argument that specifies the maximum number of digits to be displayed including both those before and those after the decimal point. For example, the C++ code

```
double value = 3.14159;  
cout << setprecision(3) << value << endl;
```

produces the following output: 3.14.

The output from this program is shown in Figure 9-21. You can compile and execute this program to see how the number of digits changes with different integer arguments provided to `setprecision`. The program, named `Precision.cpp`, is included with the data files provided with this book.

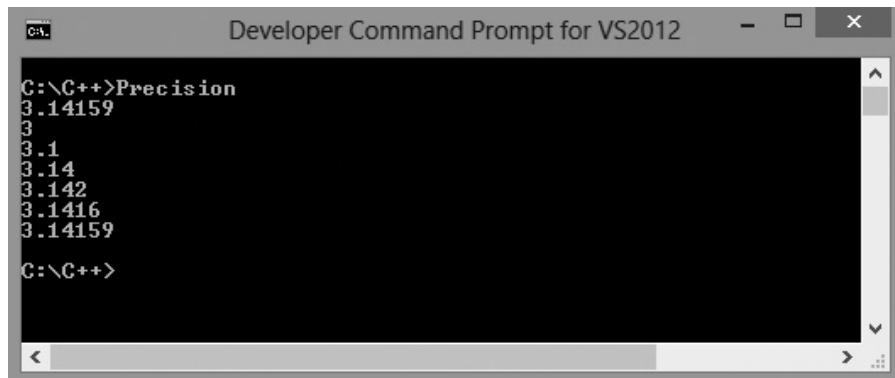
A screenshot of a Windows command prompt window titled "Developer Command Prompt for VS2012". The window shows the output of a C++ program named "Precision". The output consists of several lines of text, each showing the value 3.14159 with a different number of digits displayed. The first line shows all six digits. Subsequent lines show fewer digits, starting with three digits (3.14) and then decreasing to one digit (3). The command prompt window has a standard Windows interface with a title bar, minimize, maximize, and close buttons. The text area contains the program's output and ends with the command prompt ">".

Figure 9-21 Output from Precision program

Exercise 9-8: Using C++ Built-in Functions

In this exercise, you use a browser, such as Google, to find information about a built-in function that belongs to the C++ function library to answer Questions 1–7.

1. What does the `pow()` function do?

2. What data type does the `pow()` function return?
-

3. Is the `pow()` function overloaded? How do you know?
-

4. How many arguments does the `pow()` function require?
-

5. What is the data type of the argument(s) that the `pow()` function requires?
-

6. What is the value of the variable named `result`?

```
result = pow(2,4);
```

7. What is the value of the variable named `result`?

```
result = pow(10, 2);
```

Lab 9-8: Using C++ Built-in Functions

In this lab, you complete a partially written C++ program that includes built-in functions that convert characters stored in a character array to all uppercase or all lowercase. The program prompts the user to enter nine characters. For each character in the array, you call the built-in functions `tolower()` and `toupper()`. Both of these functions return a character with the character changed to uppercase or lowercase. Here is an example:

```
char sample1 = 'A';
char sample2 = 'a';
char result1, result2;
result1 = tolower(sample1);
result2 = toupper(sample2);
```

The source code file provided for this lab includes the necessary variable declarations and the necessary input and output statements. Comments are included in the file to help you write the remainder of the program.

1. Open the source code file named `ChangeCase.cpp` using Notepad or the text editor of your choice.
2. Write the C++ statements as indicated by the comments.
3. Save this source code file in a directory of your choice, and then make that directory your working directory.
4. Compile the source code file `ChangeCase.cpp`.
5. Execute the program with the following data:

```
uppercase
LOWERCASE
```

10

CHAPTER

Object-Oriented C++

After studying this chapter, you will be able to:

- ◎ Create a simple programmer-defined class
- ◎ Use inheritance to create a derived C++ class

This chapter covers topics that include the material covered in Chapters 10 and 11 in *Programming Logic and Design, Eighth Edition*, by Joyce Farrell.

A Programmer-Defined Class

176

Before you continue with this chapter, you should take a moment to review the object-oriented terminology (class, attribute, and method) presented in Chapter 1 of this book and in Chapters 10 and 11 of *Programming Logic and Design, Eighth Edition*.

You have been using prewritten classes, objects, and methods throughout this book. For example, you have used the `open` method that is a member of the `ifstream` and `ofstream` class to open input and output files. In this section, you learn how to create your own class that includes attributes and methods of your choice. In programming terminology, a class created by the programmer is referred to as a **programmer-defined class** or a **custom class**.

To review, procedural programming focuses on declaring data and defining functions separate from the data and then calling those functions to operate on the data. This is the style of programming you have been using in Chapters 1 through 9 of this book.

Object-oriented programming is different from procedural programming. Object-oriented programming focuses on an application's data and the functions you need to manipulate that data. The data and functions (which are called "methods" in object-oriented programming) are **encapsulated**, or contained within, a class. Objects are created as an instance of a class. The program tells an object to perform tasks by passing messages to it. Such a message consists of an instruction to execute one of the class's methods. The class method then manipulates the data (which is part of the object itself).

Creating a Programmer-Defined Class

In Chapter 10 of *Programming Logic and Design*, you studied pseudocode for the `Employee` class. This pseudocode is shown in Figure 10-1. The C++ code that implements the `Employee` class is shown in Figure 10-2.

```
1  class Employee
2      string lastName
3      num hourlyWage
4      num weeklyPay
5
6      void setLastName(string name)
7          lastName = name
8      return
9
10     void setHourlyWage(num wage)
11         hourlyWage = wage
12         calculateWeeklyPay()
13     return
```

Figure 10-1 Pseudocode for `Employee` class (continues)

(continued)

```

14     string getLastName()
15         return lastName
16
17     num getHourlyWage()
18         return hourlyWage
19
20     num getWeeklyPay()
21         return weeklyPay
22
23
24     void calculateWeeklyPay()
25         num WORK_WEEK_HOURS = 40
26         weeklyPay = hourlyWage * WORK_WEEK_HOURS
27
28     return
29 endClass

```

177

Figure 10-1 Pseudocode for Employee class

```

1 // Employee.cpp
2 #include <string>
3 using namespace std;
4 class Employee
5 {
6     public:
7         void setLastName(string);
8         void setHourlyWage(double);
9         double getHourlyWage();
10        double getWeeklyPay();
11        string getLastname();
12    private:
13        string lastName;
14        double hourlyWage;
15        double weeklyPay;
16        void calculateWeeklyPay();
17    }; // You end a class definition with a semicolon
18
19 void Employee::setLastName(string name)
20 {
21     lastName = name;
22     return;
23 }
24 void Employee::setHourlyWage(double wage)
25 {
26     hourlyWage = wage;
27     calculateWeeklyPay();
28     return;
29 }

```

Figure 10-2 Employee class implemented in C++ (continues)

(continued)

178

```
30     string Employee::getLastName()
31     {
32         return lastName;
33     }
34     double Employee::getHourlyWage()
35     {
36         return hourlyWage;
37     }
38     double Employee::getWeeklyPay()
39     {
40         return weeklyPay;
41     }
42     void Employee::calculateWeeklyPay()
43     {
44         const int WORK_WEEK_HOURS = 40;
45         weeklyPay = hourlyWage * WORK_WEEK_HOURS;
46         return;
47     }
```

Figure 10-2 Employee class implemented in C++

Looking at the pseudocode in Figure 10-1, you see that you begin creating a class by specifying that it is a class. In the C++ code in Figure 10-2, line 1 is a comment, line 2 is a preprocessor directive, `#include <string>`, and line 3 is a `using` statement, `using namespace std;`. (You learned about preprocessor directives and `using` statements in Chapter 1 of this book.) This is followed by the class declaration for the `Employee` class on line 4. The class declaration begins with the keyword `class`, which specifies that what follows is a C++ class. The opening curly brace on line 5 and the closing curly brace on line 17 mark the beginning and the end of the class.

Notice the closing curly brace on line 17 is followed by a semicolon. This closing curly brace and semicolon is *absolutely necessary* because it is part of the C++ syntax for creating a class. Omitting the semicolon is a common error.

Adding Attributes to a Class

The next step is to define the attributes (data) that are included in the `Employee` class. As shown on lines 2, 3, and 4 of the pseudocode in Figure 10-1, there are three attributes in this class: `string lastName`, `num hourlyWage`, and `num weeklyPay`.

Lines 13, 14, and 15 in Figure 10-2 include these attributes in the C++ version of the `Employee` class. Notice in the C++ code that `hourlyWage` and `weeklyPay` are defined using the `double` data type, and `lastName` is defined as a `string`. Also, notice that all three attributes are included in the **private** section of the class. The private section is defined by including the keyword `private` followed by a colon on line 12. As explained in *Programming Logic and Design*, this means the data cannot be accessed by any method (function) that is not part of the class. Programs that use the `Employee` class must use the methods that are part of the class to access private data.

Adding Methods to a Class

The next step is to add methods to the `Employee` class. The pseudocode versions of these methods, shown on lines 6 through 27 in Figure 10-1, are nonstatic methods. As you learned in Chapter 10 of *Programming Logic and Design*, **nonstatic methods** are methods that are meant to be used with an object created from a class. In other words, to use these methods, you must create an object of the `Employee` class first and then use that object to **invoke** (or call) the method.

The code shown in Figure 10-2 shows how to include methods in the `Employee` class using C++. The discussion starts with the set methods. You learned in *Programming Logic and Design* that **set methods** are methods that set the values of attributes (data fields) within the class. There are three data fields in the `Employee` class, but you will only add two set methods, `setLastName()` and `setHourlyWage()`. You do not add a `setWeeklyPay()` method because the `weeklyPay` data field is set by the `setHourlyWage()` method. The `setHourlyWage()` method uses another method, `calculateWeeklyPay()`, to accomplish this.

The two set methods, `setLastName()`, shown on lines 19 through 23 in Figure 10-2, and `setHourlyWage()`, shown on lines 24 through 29, are declared in the `public` section of the class on lines 7 and 8. The public section of the class is specified by using the keyword `public` followed by a colon on line 6. Including methods in the `public` section means that programs may use these methods to gain access to the private data. The `calculateWeeklyPay()` method, shown on lines 42 through 47 in Figure 10-2, is **private**, which means it can only be called from within another method that already belongs to the class. Notice the `calculateWeeklyPay()` method is declared in the `private` section in the `Employee` class on line 16. The `calculateWeeklyPay()` method is called from the `setHourlyWage()` method (line 27) and ensures that the class retains full control over when and how the `calculateWeeklyPay()` method is used.

The `setLastName()` method (lines 19 through 23) begins with the keyword `void`, followed by the name of the class, `Employee`, followed by the scope resolution operator (`::`), followed by the name of the method, `setLastName()`. The **scope resolution operator** is used to associate the method with the class. In this case, the `setLastName()` method is associated with the `Employee` class. The `setLastName()` method accepts one argument, `string name`, that is assigned to the private attribute `lastName`. This sets the value of `lastName`. The `setLastName()` method is a `void` method—that is, it returns nothing. Notice the declaration of the `setLastName()` method on line 7 agrees with the method definition; that is, both specify a return value of `void` and a single `string` argument.

The `setHourlyWage()` method (lines 24 through 29) accepts one argument, `double wage`, that is assigned to the private attribute `hourlyWage`. This sets the value of `hourlyWage`. Next, it calls the `private` method `calculateWeeklyPay()` on line 27. The `calculateWeeklyPay()` method does not accept arguments. Within the method, on line 44, a constant, `const int WORK_WEEK_HOURS`, is declared and initialized with the value 40. The `calculateWeeklyPay()` method then calculates the weekly pay (line 45) by multiplying the private attribute `hourlyWage` by `WORK_WEEK_HOURS`. The result is assigned to the private attribute `weeklyPay`. The `setHourlyWage()` method and the `calculateWeeklyPay()` method are `void` methods, which means they return nothing.

The final step in creating the `Employee` class is adding the get methods. **Get methods** are methods that return a value to the program using the class. The pseudocode in Figure 10-1 includes three get methods: `getLastName()` on lines 15 and 16, `getHourlyWage()` on lines 18 and 19, and `getWeeklyPay()` on lines 21 and 22. Lines 30 through 41 in Figure 10-2 illustrate the C++ version of the get methods in the `Employee` class.



It is not a requirement to use the prefix `get` when naming get methods, but by naming them using the word `get`, their intended purpose is clearer.

The three get methods are `public` methods and accept no arguments. The `getLastName()` method, shown on lines 30 through 33, returns a `string`, which is the value of the private attribute `lastName`. The `getHourlyWage()` method, shown on lines 34 through 37, returns a `double`, which is the value of the private attribute `hourlyWage`. The `getWeeklyPay()` method, shown on lines 38 through 41, also returns a `double`, which is the value of the private attribute `weeklyPay`. The three get methods are declared in the `Employee` class on lines 9, 10, and 11.

The `Employee` class is now complete and may be used in a C++ program. The `Employee` class does not contain a `main()` method (function) because it is not an application but rather a class that an application may now use to instantiate objects. The completed `Employee` class is included in the student files provided for this book, in a file named `Employee.cpp`.

Figure 10-3 illustrates a program named `MyEmployeeClassProgram` that uses the `Employee` class.

```
1 // This program uses the programmer-defined Employee class.
2 #include "Employee.cpp"
3 #include <iostream>
4 using namespace std;
5
6 int main()
7 {
8     const double LOW = 9.00;
9     const double HIGH = 14.65;
10    Employee myGardener;
11
12    myGardener.setLastName("Greene");
13    myGardener.setHourlyWage(LOW);
14    cout << "My gardener makes " << myGardener.getWeeklyPay()
15        << " per week." << endl;
16
17    myGardener.setHourlyWage(HIGH);
18    cout << "My gardener makes " << myGardener.getWeeklyPay()
19        << " per week." << endl;
20    return 0;
21 }
```

Figure 10-3 `MyEmployeeClassProgram` that uses the `Employee` class

As shown in Figure 10-3, the `MyEmployeeClassProgram` begins with a comment on line 1, followed by the preprocessor directive `#include "Employee.cpp"` on line 2. This preprocessor directive instructs the C++ compiler to include the file named `Employee.cpp`. The `Employee.cpp` file contains the `Employee` class declaration and the implementation of the methods defined in the `Employee` class. Notice the file named `Employee.cpp` is enclosed in double quotes. The double quotes specify that this is a user-created file rather than a C++ header file. C++ header file names are surrounded by angle brackets as shown on line 3, which instructs the C++ compiler to include the C++ `iostream` header file. Line 4 includes a `using` statement that specifies the `std` namespace.

Since this is a C++ program, it must contain a `main()` function. The `main()` function begins on line 6. On lines 7 and 19, you see the opening and closing curly braces that define the beginning and end of the `main()` function. Within the `main()` function, two constants, `LOW` and `HIGH`, are declared and initialized on lines 8 and 9, respectively. Next, on line 10, an `Employee` object (an instance of the `Employee` class) is created with the following statement:

```
Employee myGardener;
```

In C++, a statement that creates a new object consists of the class name followed by the object's name. In the preceding example, the class is `Employee` and the name of the object is `myGardener`.

As you learned in *Programming Logic and Design*, a **constructor** is a method that creates an object. You also learned that you can use a **default constructor**, which is a constructor that expects no arguments and is created automatically by the compiler for every class you write. The `Employee` constructor used in the `MyEmployeeClassProgram` is an example of a prewritten default constructor.



You can also write your own constructors if you wish. You will learn more about additional constructors in C++ courses you take after this course.

Once the `myGardener` object is created, you can use `myGardener` to invoke the set methods to set the value of `lastName` to "Greene" and the `hourlyWage` to `LOW`. The syntax used is shown in the following code sample:

```
myGardener.setLastName("Greene");  
myGardener.setHourlyWage(LOW);
```

This is the syntax used to invoke a method with an instance (an object) of a class.



Notice the syntax `objectName.methodName`; the name of the object is separated from the name of the method by a dot. A dot is actually a period.

On line 14 in Figure 10-3, the program then prints "My gardener makes " (a string constant) followed by the return value of `myGardener.getWeeklyPay()`, followed by the string constant " per week.", followed by a newline. Here, the `myGardener` object is used again—this time to invoke the `getWeeklyPay()` method.

On line 16, `myGardener` invokes the set method `setHourlyWage()` to specify a new value for `hourlyWage`. This time, `hourlyWage` is set to `HIGH`. The program then prints (line 17) "My gardener makes " (a string constant) followed by the return value of `myGardener.getWeeklyPay()`, followed by the string constant " per week.", followed by a newline. The `return 0;` statement on line 18 ends the program. The output from this program is shown in Figure 10-4.

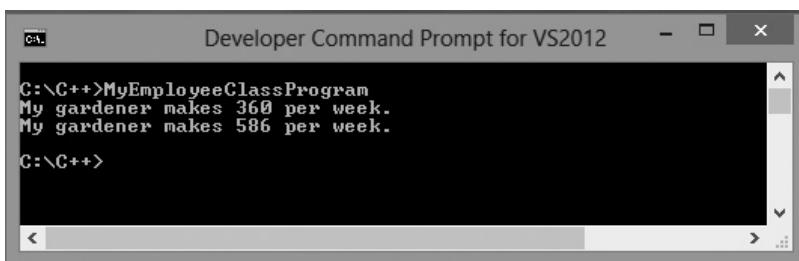
A screenshot of a Windows command prompt window titled "Developer Command Prompt for VS2012". The window shows the following text:
C:\>MyEmployeeClassProgram
My gardener makes 360 per week.
My gardener makes 586 per week.
C:\>
The window has standard Windows-style scroll bars on the right side.

Figure 10-4 Output from the `MyEmployeeClassProgram`

You will find the completed program in a file named `MyEmployeeClassProgram.cpp` included with the student files for this book.

Exercise 10-1: Creating a Class in C++

In this exercise, you use what you have learned about creating and using a programmer-defined class. Study the following code, and then answer Questions 1–4.

```
class Circle
{
public:
    void setRadius(double);
    double getRadius();
    double calculateCircumference();
    double calculateArea();
private:
    double radius; // Radius of this circle
    const double PI = 3.14159;
};
void Circle::setRadius(double rad)
{
    radius = rad;
}
double Circle::getRadius()
{
    return radius;
}
```

```
double Circle::calculateCircumference()
{
    return (2 * PI * radius)
}
double Circle::calculateArea()
{
    return(PI * radius * radius)
```

In the following exercise, assume that a `Circle` object named `oneCircle` has been created in a program that uses the `Circle` class, and `radius` is given a value as shown in the following code:

```
Circle oneCircle;
oneCircle.setRadius(4.5);
```

1. What is the output when the following line of C++ code executes?

```
cout << "The circumference is : " << oneCircle.calculateCircumference();
```

-
2. Is the following a legal C++ statement? Why or why not?

```
cout << "The area is : " << calculateArea();
```

-
3. Consider the following C++ code. What is the value stored in the `oneCircle` object's attribute named `radius`?

```
oneCircle.setRadius(10.0);
```

-
4. Write the C++ code that will assign the circumference of `oneCircle` to a double variable named `circumference1`.
-

Lab 10-1: Creating a Class in C++

In this lab, you create a programmer-defined class and then use it in a C++ program. The program should create two `Rectangle` objects and find their area and perimeter. Use the `Circle` class that you worked with in Exercise 10-1 as a guide.

1. Open the class file named `Rectangle.cpp` using Notepad or the text editor of your choice.
2. In the `Rectangle` class, create two private attributes named `length` and `width`. Both `length` and `width` should be data type `double`.
3. Write `public` set methods to set the values for `length` and `width`.
4. Write `public` get methods to retrieve the values for `length` and `width`.

5. Write a `public calculateArea()` method and a `public calculatePerimeter()` method to calculate and return the area of the rectangle and the perimeter of the rectangle.
 6. Save this class file, `Rectangle.cpp`, in a directory of your choice, and then open the file named `MyRectangleClassProgram.cpp`.
 7. In the `MyRectangleClassProgram`, create two `Rectangle` objects named `rectangle1` and `rectangle2` using the default constructor as you saw in `MyEmployeeClassProgram.cpp`.
 8. Set the length of `rectangle1` to 10.0 and the width to 5.0. Set the length of `rectangle2` to 7.0 and the width to 3.0.
 9. Print the value of `rectangle1`'s perimeter and area, and then print the value of `rectangle2`'s perimeter and area.
 10. Save `MyRectangleClassProgram.cpp` in the same directory as `Rectangle.cpp`.
 11. Compile the source code file `MyRectangleClassProgram.cpp`.
 12. Execute the program.
 13. Record the output below.
-
-
-
-

Reusing C++ Classes

You can reuse classes and reuse the code that is written for the members of a class by taking advantage of inheritance. **Inheritance** allows you to create a new class that is based on an existing class. To use inheritance, you create a new class (called the **derived class**) that contains the members of the original class (called the **base class**). You can then modify the derived class by adding new members that allow it to behave in new ways. You can also redefine methods that are inherited from the base class if they do not meet your exact needs in the derived class.

For example, if you have an existing `Vehicle` class, you could derive a new class named `Automobile` from the `Vehicle` class. You could also derive a `TaxiCab` class, a `Motorcycle` class, or a `GoKart` class from the `Vehicle` class. In C++, any class can serve as a base class.

Defining a Derived Class

You begin by taking a look at the base class, `Vehicle`. Figure 10-5 contains the C++ code that implements the `Vehicle` class.

```
1 // Vehicle.cpp
2 #include <iostream>
3 using namespace std;
4 class Vehicle
5 {
6     public:
7         void setSpeed(double);
8         double getSpeed();
9         void accelerate(double);
10        void setFuelCapacity(double);
11        double getFuelCapacity();
12        void setMaxSpeed(double);
13        double getMaxSpeed();
14    private:
15        double fuelCapacity;
16        double maxSpeed;
17        double currentSpeed;
18    };
19
20 void Vehicle::setSpeed(double speed)
21 {
22     currentSpeed = speed;
23     return;
24 }
25
26 double Vehicle::getSpeed()
27 {
28     return currentSpeed;
29 }
30
```

Figure 10-5 `Vehicle` class implemented in C++ (continues)

(continued)

186

```
31     void Vehicle::accelerate(double mph)
32     {
33         if(currentSpeed + mph < maxSpeed)
34             currentSpeed = currentSpeed + mph;
35         else
36             cout << "This vehicle cannot go that fast." << endl;
37     }
38
39     void Vehicle::setFuelCapacity(double fuel)
40     {
41         fuelCapacity = fuel;
42     }
43
44     double Vehicle::getFuelCapacity()
45     {
46         return fuelCapacity;
47     }
48
49     void Vehicle::setMaxSpeed(double max)
50     {
51         maxSpeed = max;
52     }
53
54     double Vehicle::getMaxSpeed()
55     {
56         return maxSpeed;
57     }
```

Figure 10-5 Vehicle class implemented in C++

The `Vehicle` class contains three private attributes: `fuelCapacity`, `maxSpeed`, and `currentSpeed` (lines 15 through 17). The `Vehicle` class also contains seven methods that are declared in the `public` section (lines 7 through 13). The methods are `setSpeed()`, `getSpeed()`, `accelerate()`, `setFuelCapacity()`, `getFuelCapacity()`, `setMaxSpeed()`, and `getMaxSpeed()`. The methods are implemented on lines 20 through 57. Read over the C++ code written for these methods to be sure you understand them. Three of the methods are set methods and are written in a similar manner as the set methods that belong to the `Employee` class discussed in the previous section. Three of the methods are get methods and are written in a similar manner as the get methods that belong to the `Employee` class discussed in the previous section. Note that the `accelerate()` method contains something you have not yet seen. The `if` statement that is part of this method (line 33) tests to prohibit this `Vehicle` from exceeding its maximum speed.

The C++ program `MyVehicleClassProgram.cpp` shown in Figure 10-6 uses the `Vehicle` class, and the output generated by this program is shown in Figure 10-7.

```
1 // This program uses the programmer-defined Vehicle class.
2 #include "Vehicle.cpp"
3 #include <iostream>
4 using namespace std;
5 int main()
6 {
7     Vehicle vehicleOne;
8
9     vehicleOne.setMaxSpeed(100.0);
10    vehicleOne.setSpeed(35.0);
11    vehicleOne.accelerate(10.0);
12    cout << "The current speed is " << vehicleOne.getSpeed() << endl;
13
14    vehicleOne.accelerate(60.0);
15    cout << "The current speed is " << vehicleOne.getSpeed() << endl;
16
17    return 0;
18 }
```

Figure 10-6 MyVehicleClassProgram that uses the Vehicle class

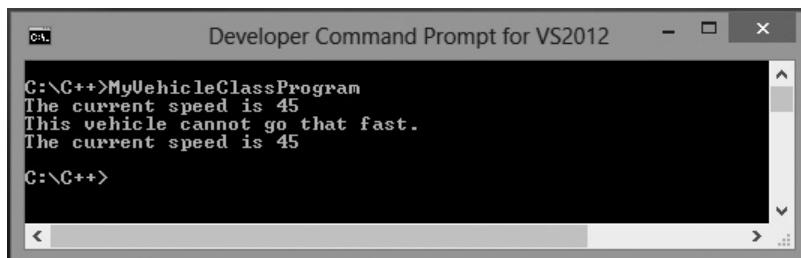
A screenshot of a Windows command prompt window titled "Developer Command Prompt for VS2012". The window shows the command "C:\>C++>MyVehicleClassProgram" followed by the program's output:
The current speed is 45
This vehicle cannot go that fast.
The current speed is 45
The output is displayed in a black terminal window with white text, and the window has a standard Windows title bar and scroll bars.

Figure 10-7 Output generated by the MyVehicleClassProgram

Looking at Figure 10-6, you see that line 1 is a comment that describes what the program does. On line 2, you see `#include "Vehicle.cpp"`. This preprocessor directive instructs C++ to include the file that contains the definition of the `Vehicle` class. Next, on lines 3 and 4, you see the `#include` directive that allows you to use `cout` in this program and the `using` statement you learned about in Chapter 1 of this book. The `main()` function begins on line 5.

On line 7, a `Vehicle` object is created named `vehicleOne`. Next, on line 9, the maximum speed for `vehicleOne` is set to 100.0 miles per hour (mph) using the `setMaxSpeed()` method. Line 10 uses the `setSpeed()` method to set the current speed for `vehicleOne` at 35.0 mph, and on line 11 `vehicleOne` accelerates by 10.0 mph using the `accelerate()` method. The

`cout` statement on line 12 produces the following output: "The current speed is 45". Notice that the `getSpeed()` method is used in the `cout` statement to retrieve `vehicleOne`'s current speed. On line 14 `vehicleOne` invokes the `accelerate()` method again; this time to accelerate by 60.0 mph. Accelerating by 60.0 mph would cause `vehicleOne` to travel faster than its maximum speed and causes the `accelerate()` method to produce the following output: "This vehicle cannot go that fast." The `cout` statement on line 15 produces the following output: "The current speed is 45". This output shows that the `accelerate` method will not allow this `Vehicle` object to travel faster than its maximum speed.

Now that the `Vehicle` class is complete, it can be used as a base class. You would like to use the `Vehicle` class to create a new `Automobile` class. You realize that an automobile is a vehicle just like the vehicle defined by the `Vehicle` class. You would like to be able to set a maximum speed, a current speed, and the fuel capacity for the automobile. You would also like to get its maximum speed, its current speed, and its fuel capacity. If you use inheritance, you will be able to do all of these actions without having to write any new code. Additionally, you would like to accelerate the automobile, but if you want the automobile to accelerate beyond its maximum speed, you would like to change the message generated to say "This automobile cannot go that fast". You also want to be able to indicate if the automobile has a convertible top. You realize you will have to modify the inherited `accelerate()` method and add new data to the `Automobile` class to indicate whether or not a particular `Automobile` object is a convertible, along with get and set methods to set the convertible top status and to get the convertible top status.

Figure 10-8 contains the C++ code that uses inheritance to create a derived class named `Automobile` using the `Vehicle` class as its base class.

```
1 // Automobile.cpp
2 #include "Vehicle.cpp"
3 #include <iostream>
4 using namespace std;
5 class Automobile : public Vehicle
6 {
7     public:
8         void accelerate(double);
9         void setConvertibleStatus(bool);
10        bool getConvertibleStatus();
11    private:
12        bool convertibleStatus;
13    };
14
15 void Automobile::accelerate(double mph)
16 {
17     if(getSpeed() + mph > getMaxSpeed())
18         cout << "This automobile cannot go that fast" << endl;
19     else
20         setSpeed(getSpeed() + mph);
21 }
22
23 void Automobile::setConvertibleStatus(bool status)
24 {
25     convertibleStatus = status;
26     return;
27 }
28
29 bool Automobile::getConvertibleStatus()
30 {
31     return convertibleStatus;
32 }
```

Figure 10-8 Automobile class implemented in C++

Looking at Figure 10-8, you see that line 1 is a comment, and on line 2 you see `#include "Vehicle.cpp"`. This preprocessor directive instructs C++ to include the file that contains the definition of the `Vehicle` class. C++ needs to know about the `Vehicle` class in order to create the `Automobile` class. Next, on lines 3 and 4, you see the `#include` directive that allows you to use `cout` in this class's methods and the `using` statement you learned about in Chapter 1 of this book. The `Automobile` class begins on line 5.

As shown on line 5, when you write the declaration for a derived class, you begin with the keyword `class` followed by the name of the derived class (`Automobile`). Next, you include a colon (:) followed by the keyword `public` and then the name of the class from which you are deriving the new class (`Vehicle`).

The keyword `public` specifies a derivation type. A **public** derivation means that all of the public members of the base class will be public in the derived class. Therefore, you do not have to repeat

these members in the derived class; you simply use the inherited members. This means that the new `Automobile` class inherits, and will be able to use, the `public` methods from the `Vehicle` class. These methods are `setSpeed()`, `getSpeed()`, `accelerate()`, `setFuelCapacity()`, `getFuelCapacity()`, `setMaxSpeed()`, and `getMaxSpeed()`. Now, in the `Automobile` class, you are able to use these `public` methods to gain access to the `private` members (`fuelCapacity`, `maxSpeed`, and `currentSpeed`) in the `Vehicle` class. In the derived `Automobile` class, you can also add new members or modify the inherited members. There are additional derivation types in C++, but in this book, you will always use the `public` derivation type. You will learn more about derivation types when you take additional courses in C++.

On line 12 in Figure 10-8, you see that one `private` data member, `convertibleStatus`, is added to the `Automobile` class. This `bool` data member is used to store a `true` or `false` value and specifies whether or not the `Automobile` has a convertible top. To provide the ability to set a value and get a value for the new `convertibleStatus` data member, two `public` methods, `setConvertibleStatus()` and `getConvertibleStatus()`, are added to the `Automobile` class on lines 9 and 10.

On line 8, you see the `accelerate()` method. This method has the same signature as the `accelerate()` method in the `Vehicle` class.

When you declare a method in a derived class (`Automobile`) with the same signature as a method in the base class (`Vehicle`), the derived class method **overrides** the inherited method. This means the method must be rewritten in the derived class and will be used with `Automobile` objects. The `accelerate()` method needs to be rewritten for the `Automobile` class because you want it to generate the message "This automobile cannot go that fast" rather than the message "This vehicle cannot go that fast."



In Chapter 9 of this book, you learned that a signature is made up of the method (function) name, the number of arguments it receives, and the data type of the arguments.

Now, the methods declared in the derived `Automobile` class must be written. The `accelerate()` method is written on lines 15 through 21. On line 17, you see the following `if` statement:

```
if(getSpeed() + mph > getMaxSpeed())
```

Notice the test portion of the `if` statement uses (1), the inherited `getSpeed()` method, to retrieve the value stored in the `private` data member `currentSpeed`, and (2), the inherited `getMaxSpeed()` method, to retrieve the value stored in the `private` data member `maxSpeed`. You do not have direct access to the `private` members of the base class (`Vehicle`); therefore, the inherited `public` methods must be used to gain access to the `private` data members `currentSpeed` and `maxSpeed`.

On line 18, you see the `cout` statement that generates the modified output: "This automobile cannot go that fast". This statement executes if you try to accelerate the `Automobile` beyond its maximum speed. If the acceleration does not cause the `Automobile` to travel at a speed that is beyond its maximum speed, line 20 executes. Line 20 invokes the

inherited `setSpeed()` method and passes the acceleration amount (`mph`) added to the current speed. The current speed is retrieved by using the inherited `getSpeed()` method.

On lines 23 through 27, you see the `setConvertibleStatus()` method, which sets the value of the `private bool` data member `convertibleStatus`. And, on lines 29 through 32, you see the `getConvertibleStatus()` method that retrieves the value of the `private bool` data member `convertibleStatus`.

Using a Derived Class in a C++ Program

Now that the derived `Automobile` class is defined, you can use it in a C++ program. The C++ program, `MyAutomobileClassProgram.cpp`, shown in Figure 10-9 includes a comment on line 1. Line 2 includes the file, `Automobile.cpp`, that contains the `Automobile` class definition; line 3 includes `<iostream>`; and the `using` statement is on line 4. An `Automobile` object (`automobileOne`) is then created on line 7, and a `bool` variable (`convertible`) is declared on line 8.

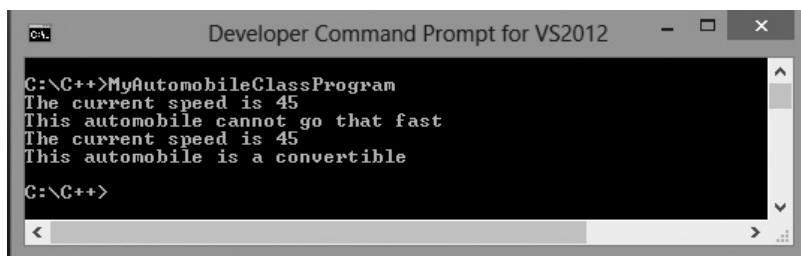
```
1 // This program uses the programmer-defined Automobile class.
2 #include "Automobile.cpp"
3 #include <iostream>
4 using namespace std;
5 int main()
6 {
7     Automobile automobileOne;
8     bool convertible;
9
10    automobileOne.setMaxSpeed(100.0);
11    automobileOne.setSpeed(35.0);
12    automobileOne.accelerate(10.0);
13    cout << "The current speed is " << automobileOne.getSpeed() << endl;
14
15    automobileOne.accelerate(60.0);
16    cout << "The current speed is " << automobileOne.getSpeed() << endl;
17
18    automobileOne.setConvertibleStatus(true);
19    convertible = automobileOne.getConvertibleStatus();
20
21    if(convertible == true)
22        cout << "This automobile is a convertible" << endl;
23    else
24        cout << "This automobile is not a convertible" << endl;
25
26    return 0;
27 }
```

Figure 10-9 MyAutomobileClassProgram that uses the Automobile class

Next, on line 10, the maximum speed for automobileOne is set to 100.0 mph using the inherited `setMaxSpeed()` method. Line 11 uses the inherited `setSpeed()` method to set the current speed for automobileOne at 35.0 mph, and on line 12 automobileOne accelerates by 10.0 mph using the `accelerate()` method that was overridden in the `Automobile` class. The `cout` statement on line 13 produces the following output: "The current speed is 45". Notice that the inherited `getSpeed()` method is used in the `cout` statement to retrieve automobileOne's current speed. On line 15, automobileOne invokes the overridden `accelerate()` method again, this time to accelerate by 60.0 mph. Accelerating by 60.0 mph would cause automobileOne to travel faster than its maximum speed, and so it causes the `accelerate()` method that was overridden in the `Automobile` class to produce the following output: "This automobile cannot go that fast". This output shows that the overridden `accelerate()` method is used with an `Automobile` object. The `cout` statement on line 16 produces the following output: "The current speed is 45". This output shows that the `accelerate()` method will not allow this `Automobile` object to travel faster than its maximum speed.

Next, on line 18, automobileOne invokes the `setConvertibleStatus()` method and passes a `true` value that indicates this `Automobile` object is a convertible. The `getConvertibleStatus()` method is invoked on line 19, where its return value is assigned to the `bool` variable `convertible`. The value of `convertible` is tested in the `if` statement on line 21. If the value of `convertible` is `true`, the `cout` statement on line 22 executes, generating the following output: "This automobile is a convertible". If the value of `convertible` is `false`, the `cout` statement on line 24 executes, generating the following output: "This automobile is not a convertible".

The output from the `MyAutomobileClassProgram` is shown in Figure 10-10.

A screenshot of a Windows command prompt window titled "Developer Command Prompt for VS2012". The window contains the following text:

```
C:\>MyAutomobileClassProgram
The current speed is 45
This automobile cannot go that fast
The current speed is 45
This automobile is a convertible
C:\>
```

The text is displayed in white on a black background, with the command prompt line starting with "C:\>".

Figure 10-10 MyAutomobileClassProgram output

Exercise 10-2: Using Inheritance to Create a Derived Class in C++

In this exercise, you use what you have learned about using inheritance to create a derived class to answer Questions 1–4.

1. Which line of code is used to create a derived class named `SubWidget` from a base class named `Widget`?
 - a. `class Widget : public SubWidget`
 - b. `class SubWidget : base public Widget`
 - c. `class Widget : derived public SubWidget`
 - d. `class SubWidget : public Widget`

2. An advantage of using inheritance is:
 - a. It maximizes the number of functions.
 - b. It allows reuse of code.
 - c. It requires no coding.

3. True or False: The methods in a derived class have direct access to the base class private data members.

4. True or False: A derived class may add new methods or override existing methods when inheriting from a base class.

193

Lab 10-2: Using Inheritance to Create a Derived Class in C++

In this lab, you create a derived class from a base class, and then use the derived class in a C++ program. The program should create two `Motorcycle` objects, and then set the `Motorcycle`'s speed, accelerate the `Motorcycle` object, and check its sidecar status. Use the `Vehicle` and `Automobile` classes that you worked with earlier in this chapter as a guide.

1. Open the file named `Motorcycle.cpp` using Notepad or the text editor of your choice.
2. Create the `Motorcycle` class by deriving it from the `Vehicle` class. Use a `public` derivation.
3. In the `Motorcycle` class, create a `private` attribute named `sidecar`. The `sidecar` attribute should be data type `bool`.
4. Write a `public` `set` method to set the value for `sidecar`.
5. Write a `public` `get` method to retrieve the value of `sidecar`.
6. Write a `public` `accelerate()` method. This method overrides the `accelerate()` method inherited from the `Vehicle` class. Change the message in the `accelerate()`

method so the following is displayed when the `Motorcycle` tries to accelerate beyond its maximum speed: "This motorcycle cannot go that fast".

7. Save this class file, `Motorcycle.cpp`, in a directory of your choice, and then open the file named `MyMotorcycleClassProgram.cpp`.
 8. In the `MyMotorcycleClassProgram`, create two `Motorcycle` objects named `motorcycleOne` and `motorcycleTwo`.
 9. Set the `sidecar` value of `motorcycleOne` to `true` and the `sidecar` value of `motorcycleTwo` to `false`.
 10. Set `motorcycleOne`'s maximum speed to 90 and `motorcycleTwo`'s maximum speed to 85.
 11. Set `motorcycleOne`'s current speed to 65 and `motorcycleTwo`'s current speed to 60.
 12. Accelerate `motorcycleOne` by 30 mph, and accelerate `motorcycleTwo` by 20 mph.
 13. Print the current speed of `motorcycleOne` and `motorcycleTwo`.
 14. Determine if `motorcycleOne` and `motorcycleTwo` have sidecars. If yes, display the following: "This motorcycle has a sidecar". If not, display the following: "This motorcycle does not have a sidecar".
 15. Save `MyMotorcycleClassProgram.cpp` in the same directory as `Motorcycle.cpp`.
 16. Compile the source code file `MyMotorcycleClassProgram.cpp`.
 17. Execute the program.
 18. Record the output below.
-
-
-
-

Index

Note: Page numbers in **boldface** indicate where key terms are defined.

A

`accelerate()` method, 186–188, 190, 192
accumulator, 87
actual parameters, **143**
`add()` function, 168
addition operator (+), 19, 21
`ADDRESS_LINE1` constant, 144
`ADDRESS_LINE2` constant, 144
`ADDRESS_LINE3` constant, 144
address operator (&), 163, 165
`age` variable, 163
algorithms, **126**
aliases, **161**
All apps button, 7
All Programs, Accessories,
 Notepad command, 7
All Programs, Accessories, Visual
 Studio 2012 Express,
 Developer Command
 Prompt for VS2012
 command, 7
`amount` formal parameter, 151
ampersand operator (&), 161
AND logic, 63
AND operator (&&), 45–46, 47, 63
`AnswerSheet.cpp` file, 85
Answer Sheet program, 85
`answer` variable, 20, 90
arguments, **143**, 159–167
 changing value of, 162
 pass by address, 162–167
 pass by reference, 161–162
arithmetic operators, **18–19**
arrays, **94**
 accessing elements, 97

assigning values to elements, 96–97
bubble sort, 128–134
constants, 128
counting elements in, 132
data types, 94
declaring, **94–96**
`double` data type, 95
elements, 94
as formal parameter, 158
garbage values, **96**, 98
initializing, 96–97
`int` data type, 94
memory allocation, 94–95
multidimensional, 135–138
named constants, 98
naming, 94
one-dimensional, **135**
parallel, 103–106
passing array and array
 element to, 156–159
passing by value, 157
passing to functions, 156–159
primitive data types, 95
printing out values, 158
relationship between elements, 103–106
searching for exact match, 100–102
single-dimensional, **135**
sorting records, 126
square brackets ([]), 94
staying within bounds, 98
storing data, 132
`string` objects, 94, 95
subscripts, **94**, 97–98, 133
two-dimensional, **135–138**

B

`balance` variable, 143, 150–151
`bal` variable, 170
base class, **184**
`BEDROOMS` constant, 136
behaviors, 3
block statements, 51, 72–73, 80
`bool` data type, 13, 101–102
Boolean comparison, **72**
Boolean operators, **44**
 logical operators, **45–46**
 relational operators, **44–45**
true or false values, 44

branching statements, 50
See also decision statements
break statement, 60–61
bubble sort, 128–134
built-in functions, 172–173

C

calculatedAnswer variable, 31–32
calculateWeeklyPay() method, 179
calling methods, 3
camel case, **13**
case keyword, 60
case sensitivity, 12, **49**
cd command, 7
C++ development cycle
compiling C++ programs, 7–8
executing C++ program, 8–9
writing source code, 6–7
Celsius temperature: string, 25
celsius variable, 24–25
characters, 13–14
char data type, 13
cin object, 25
cin statement, 88
cities array, 94, 97
memory allocation, 95–96
cityPopulations array, 94, 97
memory allocation, 95
classes, **3–4**
associating methods with, 179
base class, **184**
custom, **176**
derived classes, **184**
inheritance, **184**, 188
instances, **3**
programmer-defined, **176–182**
reusing, 184–192
class keyword, 178, 189
classRank variable, 33
cl compiler, 7–9, 25
Client By State program,
117–121
City field, 117
comments, 119
main() function, 119–120
Name field, 117
State field, 117
client.dat file, 120
close() method, 113

code
polymorphic, **169**
reusing, 4
COLUMN_HEADING constant, 36–37
Command Prompt window, 7–8
comments, **23**
multiline (*/* */*), 24
single-line (*//*), 24
comparing strings, 48–50
comparisons variable, 132–133
compiling C++ programs, 7–8
computeTax() function,
149–151
Compute Tax program, 149–151
constants, **17**
array size, 128
global, 35
local, 35, 37–38
named, **17**
naming, 17, 37, 98
numeric, **17**
string, **17**
unnamed, **17**
const keyword, 17
constructors, **181**
controlBreak() function, 120
control break logic, 116–121
convertibleStatus data
member, 190–191
convertible variable, 191–192
counters, 87
controlling loops, 74–75
counter variable, 15, 83
count variable, 75, 80–81,
120–121
cout statement, 5, 25, 88, 154,
164, 187–189, 192
.cpp extension, 7
C programming language, 2
C++ programming language, 2
case sensitivity, 12, **49**
data types, **13–14**
C++ programs
See also programs
compiling, 7–8
executing, 8–9
structure, 4–6
currentSpeed attribute, 186, 190
custom classes, **176**
customerAge variable, 51
Customer Bill program, 142–144

D

data
encapsulated, 3, **176**
fields, **110**
hierarchy, 110
records, **110**
user input, 24–25, 31
data_in object, 110–112
data_out object, 112–113
data types, **13**
arrays, 94
pointer variables, 163
variables, 13–14
decimal numbers, number of
places after decimal point,
172–173
decision statements, 50
dual-alternative, 54–55
dual-path, 54–55
if else statements, 54–55
if statements, 50–52
AND logic, 63
multipath, 57–58, 60–61
multiple comparisons, 63–64
nesting if statements, 57–58
OR logic, 64
single-path, **50–52**
switch statements, 60–61
declaring
arrays, **94–96**
functions, 143
pointer variables, 163
variables, 12, **14–15**
decrement operator (--)
postfix form, **70–71**, 73
prefix form, **70–71**, 73
deduct variable, 36, 78
default constructor, **181**
default keyword, 60
definite loops, **80**
dentalPlan variable, 63–64
dentPlan object, 52
deptName variable, 61
deptNum variable, 61
dereferencing pointers, **165**
derived classes, **184**
declaring methods, 190
defining, 185–191
methods overriding inherited
methods, **190**
in programs, 191–192

- d**
- `detailLoop()` module, 38–39
 - detail loop tasks, **35**
 - Developer Command Prompt for VS2012 command, 7
 - `didSwap` variable, 132–133
 - `dir` command, **8**, 9
 - directory listing, 8
 - `displayArray()` function, 133–134
 - division, floating-point, 25
 - division operator (/), 19, 21
 - documenting source code, 23–24
 - `done` variable, 120
 - `double` amount formal parameter, 150
 - `double amt` parameter, 162
 - `double` data type, 13–15, 17, 25, 97, 154, 172
 - arrays, 95
 - `double` keyword, 154
 - `double& new_sal` parameter, 162
 - `do until` loops, 82
 - `do while` loops, 82–83
 - dual-alternative decision statements, 54–55
 - dual-path decision statements, 54–55
 - dynamic memory allocation, **95**–**96**
- E**
- elements
 - accessing, 97
 - arrays, 94
 - counting, 132
 - garbage values, **96**
 - passing to functions, 156–159
 - `else` keyword, 54–55, 57
 - `empDept` variable, 57–58
 - `Employee` class
 - adding attributes, 178
 - creation, 176–178
 - get methods, **180**
 - methods, 179–182
 - private section, 179
 - public section, **179**
 - set methods, **179**
 - `Employee.cpp` file, 181
 - empty string (""), 96
 - encapsulated, **176**
 - encapsulating data and tasks, **3**
- F**
- `Fahrenheit` temperature:
 - string, 25
 - `fahrenheit` variable, 24–25
 - fields, **110**
 - `fileIOTest.cpp` file, 113
 - files
 - input, 110
 - naming, 7
 - object code, 7
 - opening for reading, 110–111
 - opening for writing, 112–113
 - output, 110
 - reading data, 111–112
 - saving, 7
 - sequential, **116**–**121**
 - writing data to, 113–114
 - `fillArray()` function, 132
 - `finishUp()` function, 121
 - `firstName` variable, 13, 111, 113, 143
 - `flag` variable, 128
 - `float` data type, 13
 - floating-point division, 25
 - floating point numbers, number of places after decimal point, 172–173
 - floating-point values, 13
 - `FLOORS` constant, 136
 - flowcharts, 30–34
 - flow of control, **50**
 - `fn` variable, 13
 - `for` loops, 80–81, 88
 - formal parameters, **144**
 - arrays as, 158
 - `foundIt` variable, 101–102
 - `freshmanStudentFirstName` variable, 13
 - `fstream` class, 110
 - `fuelCapacity` attribute, 186, 190
 - function declaration, **143**
 - function prototype, **143**
 - functions, **5**
 - actual parameters, **143**
 - arguments, **143**, 159–167
 - braces ({}) and, 5
 - built-in, 172–173
 - declaring, 143
 - encapsulated, **176**
 - formal parameters, **144**
 - headers, **5**, **143**–**144**
 - local constants, 35
 - local variables, **35**
 - manipulators, **172**–**173**
 - multiple-parameter, 149–151
 - multiple versions with same name, 168–171
 - overloading, **168**–**171**
 - pass by address, **161**
 - pass by reference, **161**
 - passing argument by value, **147**
 - polymorphic, 169
 - return type, **152**
 - signature, **169**
 - single-parameter, 145–147
 - value-returning, 5, 152–154
 - without parameters, 142–144

G

 - garbage values, 87, **96**, 98
 - `getConvertibleStatus()` method, 190–192
 - `getFuelCapacity()` method, 186, 190
 - `getHourlyWage()` method, 180

`getHoursWorked()` function, 152–154
`getLastname()` method, 180
`getMaxSpeed()` method, 186, 190
`get` methods, 180, 186
`getReady()` function, 120
`getSpeed()` method, 186, 188, 190–192
`getWeeklyPay()` method, 180–181
global constants, 35
global variables, 35, 37
`GoKart` class, 184
`grades` array, 103–104
greater than operator (`>`), 44–45
greater than or equal to operator (`>=`), 44–45
`grossPay` variable, 52, 55
`gross` variable, 36, 78, 153–154

H

header files, 4, 181
headers, 5, 143–144
HelloWorld command, 8
`HelloWorld.cpp` file, 7–9
`HelloWorld.exe` file, 8–9
`HelloWorld.obj` file, 8–9
Hello World program, 4–6, 8
HIGH constant, 181–182
high-level programming languages, 2
`hourlyWage` attribute, 178–182
`HOURS_IN_WEEK` constant, 54–55
`hours` variable, 153–154
`hoursWorked` variable, 55
`housekeeping()` function, 37–38
`housekeeping()` method, 77
`housekeeping()` module, 37–38
housekeeping tasks, 35

I

`if else` statements, 54–55
`if` keyword, 50, 54
`if` statements, 50–52, 101, 120, 147, 186, 190
less than operator (`<`), 51
nesting, 57–58
parentheses () and, 50–51
syntax, 50
testing string objects for equality, 52

`ifstream` class, 110–111, 111
`#include` preprocessor directive, 4, 33, 37, 143, 178, 181, 187, 189
`increase()` function, 161–162
`IncreaseSalary.cpp` file, 161–162
increment operator (`++`)
postfix form, 70–71, 73
prefix form, 70–71
indirection operator (`*`), 163–166
infinite loops, 72, 75
inheritance, 184, 188
`Automobile` class, 190
initialization operator (`=`), 20
initializing
arrays, 96–97
variables, 15
input, validating with loops, 89–90
input classes, 110
`inputFile.dat` file, 111
input files, 110–111
input operator (`>>`), 25
insertion operator (`<<`), 5, 113
instances, 3, 4
instantiation, 3, 4
`int` data type, 13–15, 17, 19
arrays, 94
integers, 5–6, 13
interactive input statements, 24
`int` keyword, 5, 6
invoking methods, 179
`iostream` header file, 4, 181

K

keyboard, 25
EOF character, 112
keywords, 5, 12

L

`lastName` attribute, 178–181
`lastName` variable, 111, 113, 143
LENGTH constant, 157
less than operator (`<`), 44–45, 98
`if` statements, 51
less than or equal to operator (`<=`), 44–45, 98
local constants, 35, 37–38
local variables, 35, 37–38, 142
logical operators, 45–46
precedence and associativity, 46–48

logic error, 51–52, 72
loop body, 72–73
accumulator, 87
counter, 87
loop control variable, 72, 78, 85
loopIndex variable, 97–98
loops
accessing array elements, 98
accumulating totals, 87–88
assigning values to array elements, 96–97
braces ({}), 32, 33, 81
counter-controlled, 74–75, 81
decrement operator (`--`), 70
definite, 80
`do until` loops, 82
`do while` loops, 82–83
Increment operator (`++`), 70
infinite, 72, 75
loop body, 72–73
loop control variable, 72, 85
`for` loops, 80–81
nesting, 84–85
priming read, 31, 77
reading data, 112
sentinel values, 72, 76–78
sequential statements, 32
validating input, 89–90
`while` loops, 72–74
LOW constant, 181
low-level programming languages, 2
lvalue, 70–71

M

`MailOrder2.cpp` file, 104
Mail Order program, 100–102
Mail Order 2 program, 104–106
`main()` function, 5–6, 31, 35, 37–38, 119–120, 131–132, 143, 147, 153, 157–159, 162, 166, 181, 187
`main()` module, 35
manipulators, 172–173
`maxSpeed` attribute, 186, 190
`MAX_STUDENTS` constant, 17
`medicalPlan` variable, 63–64
memory allocation
arrays, 94–95
`cities` array, 95–96
`cityPopulations` array, 95

- dynamic, **95–96**
string objects, **95**
- methods, **3**
 associating with class, **179**
 calling, **3**
 constructors, **181**
 derived classes, **190**
Employee class, **179–182**
 encapsulated, **176**
 invoking, **179**
 nonstatic, **179**
 overriding inherited methods, **190**
 programmer-defined classes, **179–182**
 public, **180**
 set methods, **179**
 modular programs, **35–40**
 modulus operator (%), **19, 21, 147**
Motorcycle class, **184**
 multidimensional arrays, **135–138**
 multiline block /* */ comments, **24**
 multipath decision statements, **57–58, 60–61**
 multiple comparisons, **63–64**
 multiple-parameter functions, **149–151**
 multiplication operator (*), **19, 21, 163**
myAge variable, **13**
MyAutomobileClassProgram.cpp file, **191–192**
MyEmployeeClassProgram, **180–182**
MyEmployeeClassProgram.cpp file, **182**
myGardener object, **181**
My Program folder, **7**
myRent variable, **136**
MyVehicleClassProgram.cpp file, **187–188**
- N**
- nameAndAddress()** function, **143–144, 152**
name argument, **179**
 named constants, **17**
 arrays, **98**
 namespaces, **4–5**
name variable, **36, 38, 78**
 naming conventions, **13**
- nesting
 if statements, **57–58**
 loops, **84–85**
net variable, **36, 78**
 newline, **5–6**
newNum variable, **159**
new_salary argument, **162**
newSalary variable, **113**
new_salary variable, **162**
 nonstatic methods, **179**
Notepad, **7**
 not equal to operator (!=), **44–45**
 NOT operator (!), **45–46, 112**
 null statement (;), **51–52, 72**
NumberDouble.cpp file, **31**
Number-Doubling program, **31–34**
 body of loop, **32–33**
 flowchart and pseudocode, **30**
 input and output, **33**
 main() function, **31**
 starting, **31**
numberOfElts variable, **132**
number parameter, **147**
numbers array, **98–99**
number variable, **70, 80–81, 146–147**
number1 variable, **44–47**
number2 variable, **44–47**
num data type, **14, 31**
 numeric constants, **17**
num loop control variable, **73–74**
NUM_LOOPS constant, **81**
numStudents variable, **88**
num variable, **163–164**
num1 variable, **19–20, 166**
num2 variable, **19–20, 166**
- O**
- object code, **7–8**
 object-oriented programming, **176**
 object-oriented programming languages, **2–3**
 object-oriented programs, **3**
 object-oriented terminology, **2–4**
 objects, **3**
 .obj extension, **7–8**
ofstream class, **110, 112–113**
oldState variable, **120**
 one-dimensional arrays, **135**
 opening files for reading, **110–111**
- open()** method, **111–113**
 operators, **18**
 arithmetic operators, **18–19**
 assignment operators, **19–20**
 assignment statement, **19–20**
 Boolean operators, **44–48**
 logical operators, **45–46**
 precedence and associativity, **20–21**
 relational operators, **44–45**
originalNumber variable, **31–32**
 OR logic, **64**
 OR operator (||), **45–47, 64**
 output classes, **110**
outputFile.dat file, **113**
 output files, **110**
 writing data to, **112–114**
 output operator (<>), **5**
Overloaded.cpp file, **171**
 overloading functions, **168–171**
 signature, **169**
OVERTIME_RATE OVE constant, **54–55**
overtime variable, **55**
- P**
- parallel arrays, **103–106**
 subscripts, **104**
 parentheses () precedence and associativity, **21**
partCounter variable, **85**
 pass by address, **161, 162–167**
PassByAddress.cpp file, **167**
 pass by reference, **161–162**
 Pass Entire Array program, **156**
 main() function, **157–159**
 quadrupleTheValues() function, **157–159**
 passing argument by value, **147**
 passing by value, **157**
PAY_RATE constant, **153**
PayrollReport.cpp file, **37**
 Payroll Report program
 detailLoop() module, **38–39**
 endOfJob() module, **39–40**
 flowchart, **36–40**
 housekeeping() module, **37–38**
 main() function, **37**
 sentinel values to control while loop, **76–78**

- pCnt** parameter, 162
PointerPractice.cpp file, 165
pointers, **162–163**
 - address operator (`&`), 163
 - dereferencing, **165****pointer variables**, 163–164
polymorphic code, **169**
polymorphic functions, 169
postfix form, **70–71**, 73
precedence, **20–21**
Precision.cpp file, 173
prefix form, **70–71**, 73
preprocessor, **4**
preprocessor directives, **4–5**
price variable, 104
priming loops, 31
priming read, 31, 77, 120
primitive data types, **13**
 - arrays, 95**printBill()** function, 169–171
private keyword, **178–179**
private section, **178–179**
procedural programming, **2–3**, 176
procedural programs, **2–3**
 - declaring variables and constants, 35
 - detail loop tasks, **35**
 - end-of-job tasks, **35**
 - housekeeping tasks, **35**
 - users, **3****produceReport()** function, 120
programmer-defined classes, **176**
 - adding attributes, 178
 - creation, 176–178
 - get methods, **180**
 - methods, 179–182
 - private section, **178**
 - public section, **179****programming languages**, **2–3**
programs
 - comments, **23–24**
 - compiling, 7–8
 - derived classes, 191–192
 - executing, 8–9
 - flowcharts, 30–34
 - flow of control, **50**
 - logic error, **51–52**
 - main()** module, 35
 - modular, 35–40
 - object-oriented, **3****procedural**, **2–3**
pseudocode, 30–34
single-level control break, **116**
structure, 4–6
prompt, **24**, 25
pseudocode, 30–34
ptr_age variable, 163
ptr_num variable, 164
public derivation, **189–190**
public keyword, 189
public methods, 180
public section, **179**
- Q**
- quadrupleTheValues()** function, 157–159
questionCounter variable, 85
QUIT constant, 36, 39, 78, 132
- R**
- raise** argument, 162
raise variable, 162
RATE constant, 36
rate formal parameter, 150–151
rate variable, 150–151
reading data
 - from input files, 111
 - with loops and EOF (end of file), 112**reading files**, 110–111
records, **110**
 - sorting, 126**relational operators**, **44–45**
 - comparing string objects, 48–50
 - precedence and associativity, 46–48**rent** array, 136–138
REPORT_HEADING constant, 36–37
return 0; statement, 33
return statement, 154
return type, **152**
reusing classes, 184–192
reusing code, 4
Ritchie, Dennis, 2
- S**
- salary** argument, 162
salary variable, 111, 113, 162
salePrice variable, 13
- scope resolution operator (::)**, **179**
scores array, 131–134
Score Sorting program, 128–130
 - displayArray()** function, 132–134
 - fillArray()** function, 132
 - main()** function, 131–132
 - sortArray()** function, 132–133**score2** variable, 126
search algorithms, 126
searching arrays for exact match, 100–102
sentinel values, **72**, 76–78, 146
sequences, **23**
 - interactive input statements, **24****sequential files**, **116–121**
sequential statements, **23–24**
 - loops, 32**setConvertibleStatus()**
 - method, 190–192**setFuelCapacity()** method, 186, 190
setHourlyWage() method, 179, 182
setLastName() method, 179
setMaxSpeed() method, 186–187, 190, 192
set methods, **179**, 186
setprecision() function, 172–173
setSpeed() method, 186–187, 190–192
signature, **169**, 190
simple statements, 51
single alternative statements, 50
single-dimensional arrays, **135**
single-level control break
 - programs, **116****single-line comments (//)**, 24
single-parameter functions, 145–147
single-path decision statements, 50–52
single-sided statements, 50
SIZE constant, 131–132
someNums array, 157, 159
sortArray() function, 132–133
sorting
 - bubble sort, 128–134
 - records, 126

- s**ource array, 97
source code, 7
 compiling, 7
 documenting, 23–24
 writing, 6–7
source code files, 7
 syntax errors, 8
Standard C++ library, 5
standard input device, 25
Start button, 7
Start screen, 7
state control break variable, 117, 120
statements
 block, 51
 ending with semicolon (;), 5
 executing, 2
 interactive input statements, 24
 sequential statements, 23
 simple, 51
std namespace, 4–5
string constants, 5, 17
 comparing to string objects, 48–49
string objects, 14, 37
 arrays, 94, 95
 ASCII values of characters in, 49
 comparing, 48–50
 comparing to string constants, 48–49
 dynamic memory allocation, 95–96
 memory allocation, 95
 testing for equality, 52
strings, 14
 comparing, 48–50
 empty (""), 96
string variables, 15, 17
Stroustrup, Bjarne, 2
stuCount variable, 88
stuID array, 103–104
subscripts, 94, 97–98, 133
 parallel arrays, 104
subtraction operator (-), 19, 21
supervisorName variable, 57–58
swap() function, 133, 166
swapping values, 126–127
switch keyword, 60
switch statements, 60–61
syntax, 50
syntax errors, 8
- T**
target array, 97
tasks, objects encapsulating, 3
TaxiCab class, 184
Temperature.cpp file, 24–25
temp variable, 132, 166
TestAverage.cpp file, 88
Te~~st~~ing folder, 7
testScore variable, 88
testTotal accumulator, 88
testTotal variable, 87
test1 variable, 21
test2 variable, 21
text editors, 7
tripleTheNumber() function, 159
true or false variables, 13
two-dimensional arrays, 135–138
- U**
unary minus operator (-), 19, 21
unary plus operator (+), 19, 21
unamed constants, 17
users, 3
 data input, 24–25, 31
using statement, 178, 187, 189
- V**
VALID_ITEMS array, 104
VALID_PRICES array, 104
va11 parameter, 166
va12 parameter, 166
value-returning functions, 5, 152–154
values
 relationship between, 44–45
 swapping, 126–127
value variable, 15
variable declaration, 14–15
variables, 12–14
 aliases, 161
 assigning initial value to, 15, 20
 camel case, 13
 characters, 13, 14
 data types, 13–14
 declaring, 12, 14–15
 global, 35, 37
 initializing, 15
- i**nvalid names, 12
local, 35, 37–38, 142
memory address, 15
naming, 12–13
numeric values, 13–14
pass by address, 161, 162–167
pass by reference, 161
swapping values, 126–127
true or false, 13
values stored in, 12
Vehicle class, 184–187, 189–190
 as base class, 188
Vehicle.cpp file, 189
vehicleOne object, 187–188
void keyword, 143–144, 147, 152, 179
- W**
wage argument, 179
weeklyPay attribute, 178, 180
while loops, 32, 38–39, 72–74, 90, 112, 120, 132–134, 146–147, 157–158
 counter-controlled, 74–75
 loop body, 32–33
 sentinel values, 76–78
whole numbers, 13
Windows 7
 opening Command Prompt window, 7
 starting Notepad, 7
Windows 8
 opening Command Prompt window, 7
 starting Notepad, 7
workHours variable, 154
WORK_WEEK_HOURS constant, 179
writing
 to files, 112–113
 source code, 6–7
- X**
x variable, 132–133, 157
- Y**
Y string constant, 52
- Z**
zero (0) indicating end of input, 32

